

# Machine Learning and Deep Learning

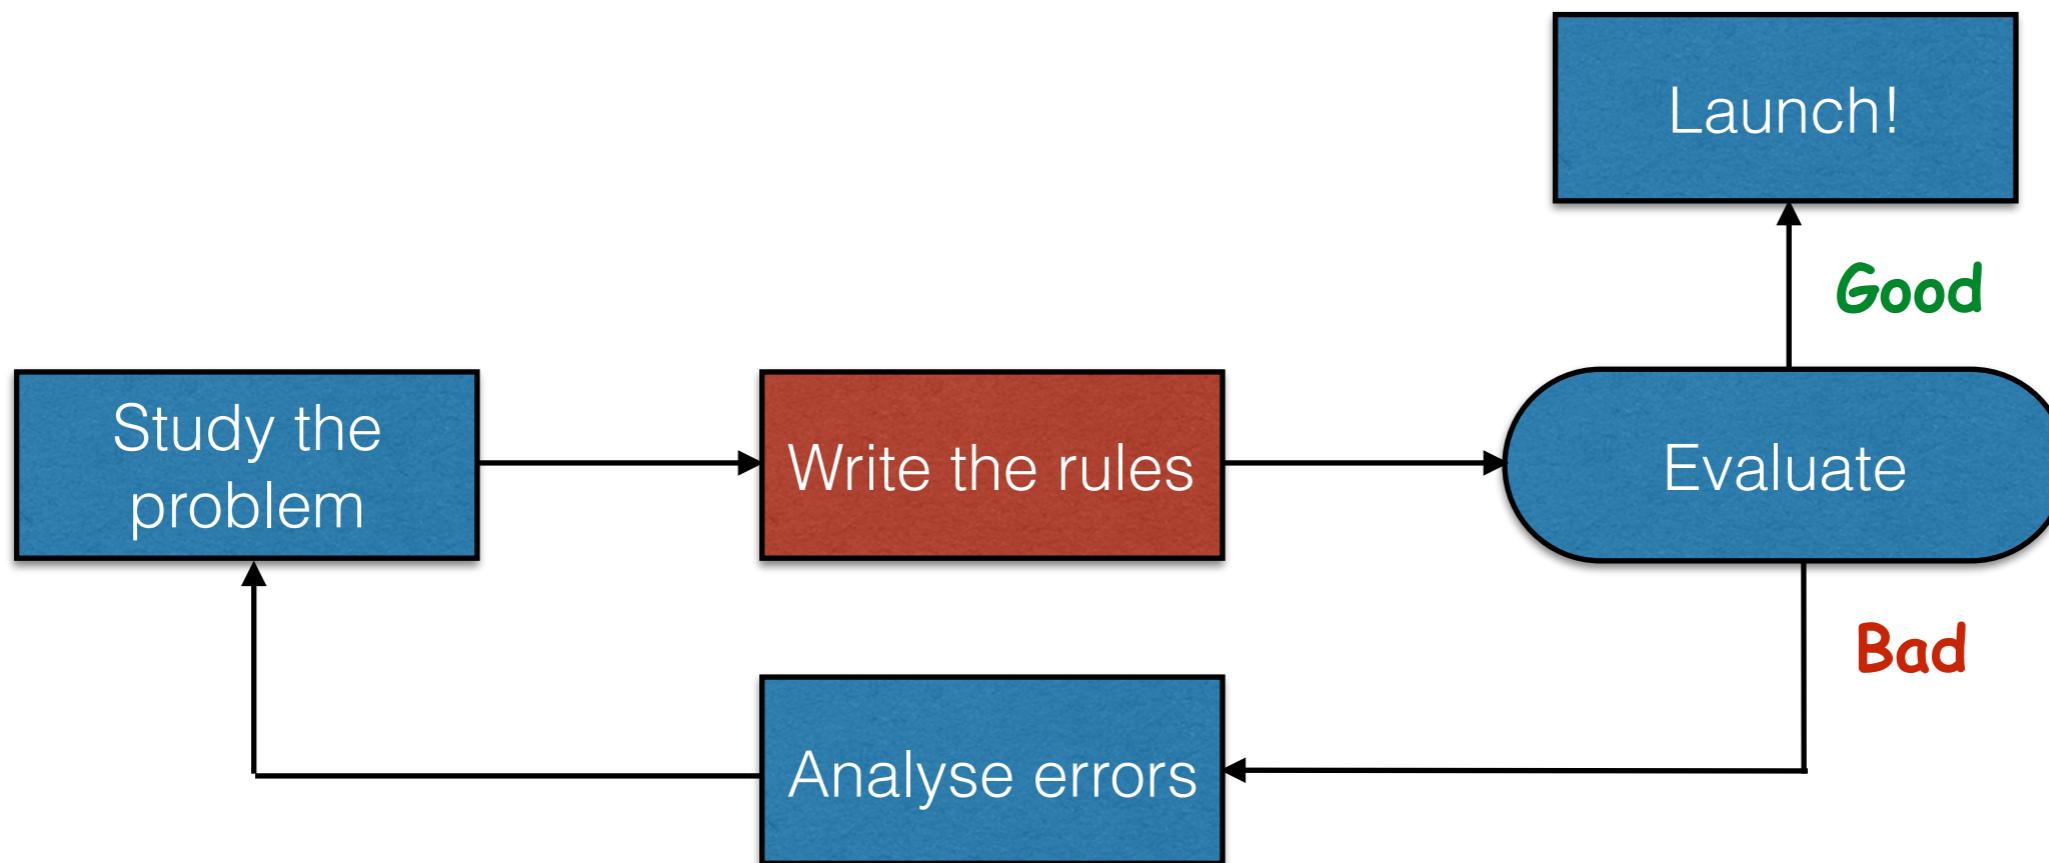
Simon Jenni  
(slides by Paolo Favaro)

# Deep Learning

- **Objective:** Build a machine that can learn from experience and understand the world as a hierarchy of concepts

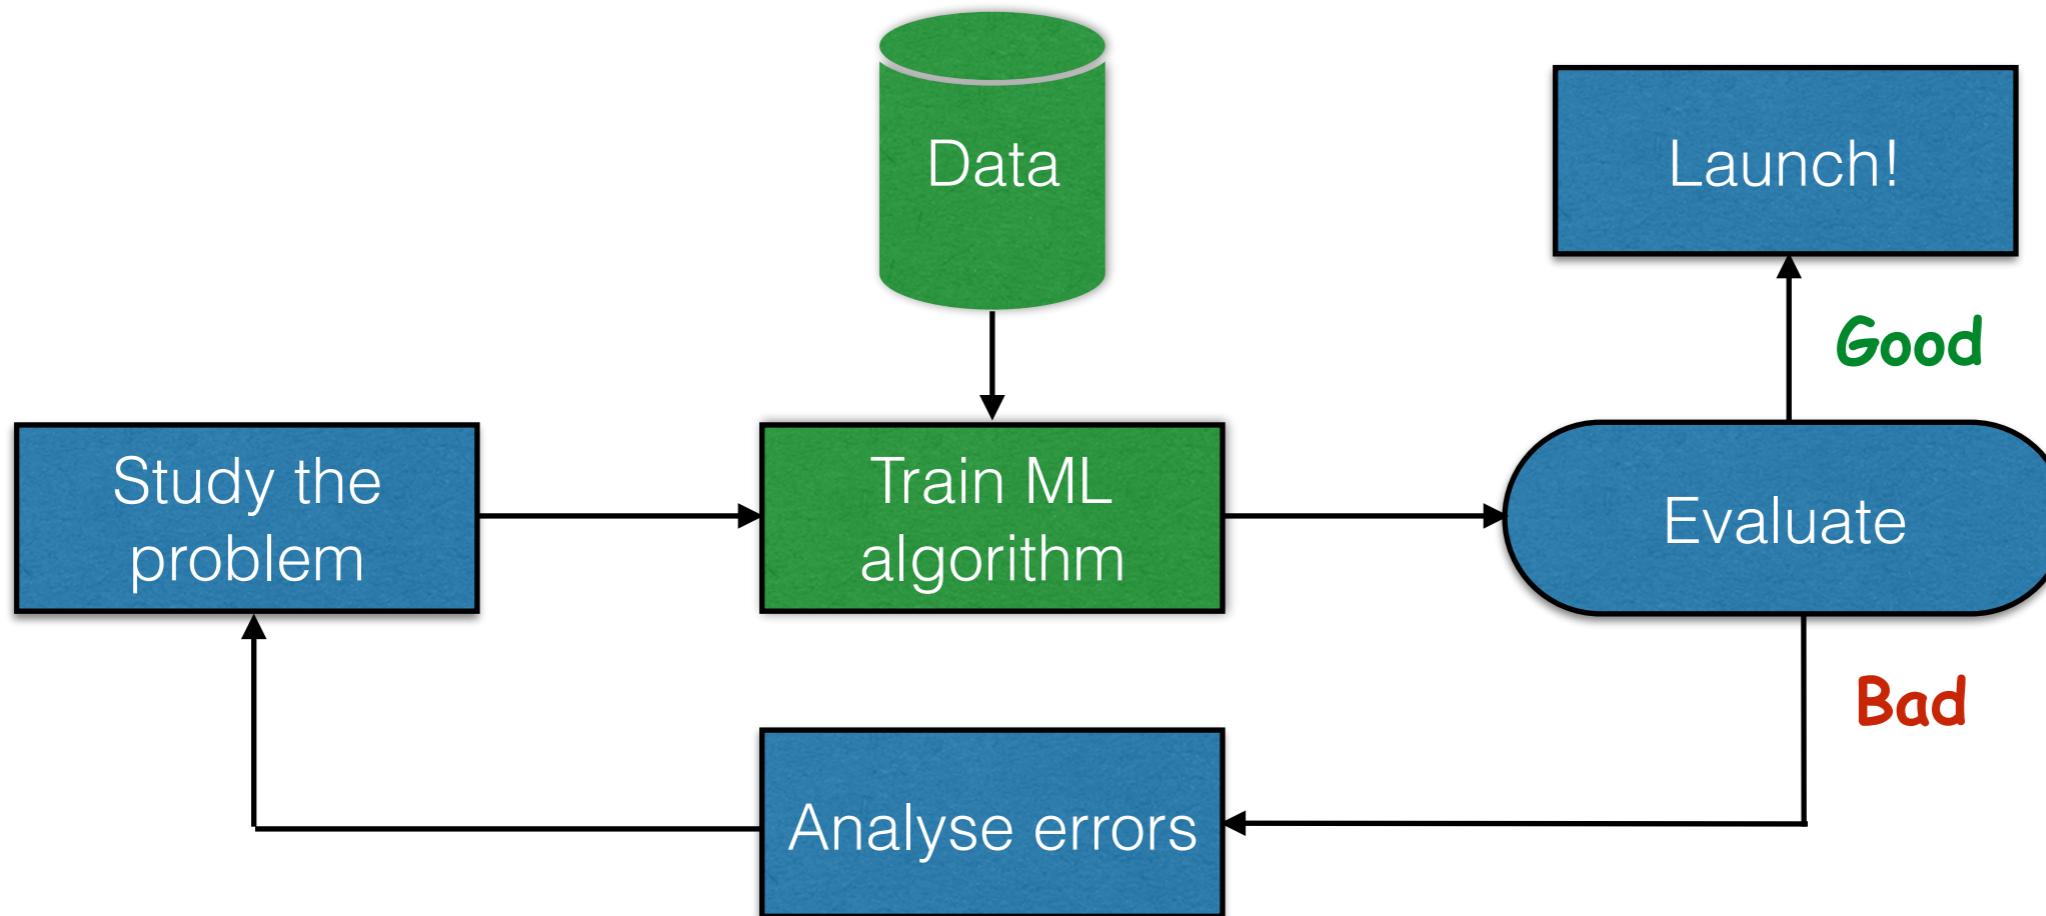
# Traditional Approach

- List of all the knowledge and formal rules
  - works for games and simple systems
  - leads to a combinatorial problem
  - **not general (often we do not know the rules)**



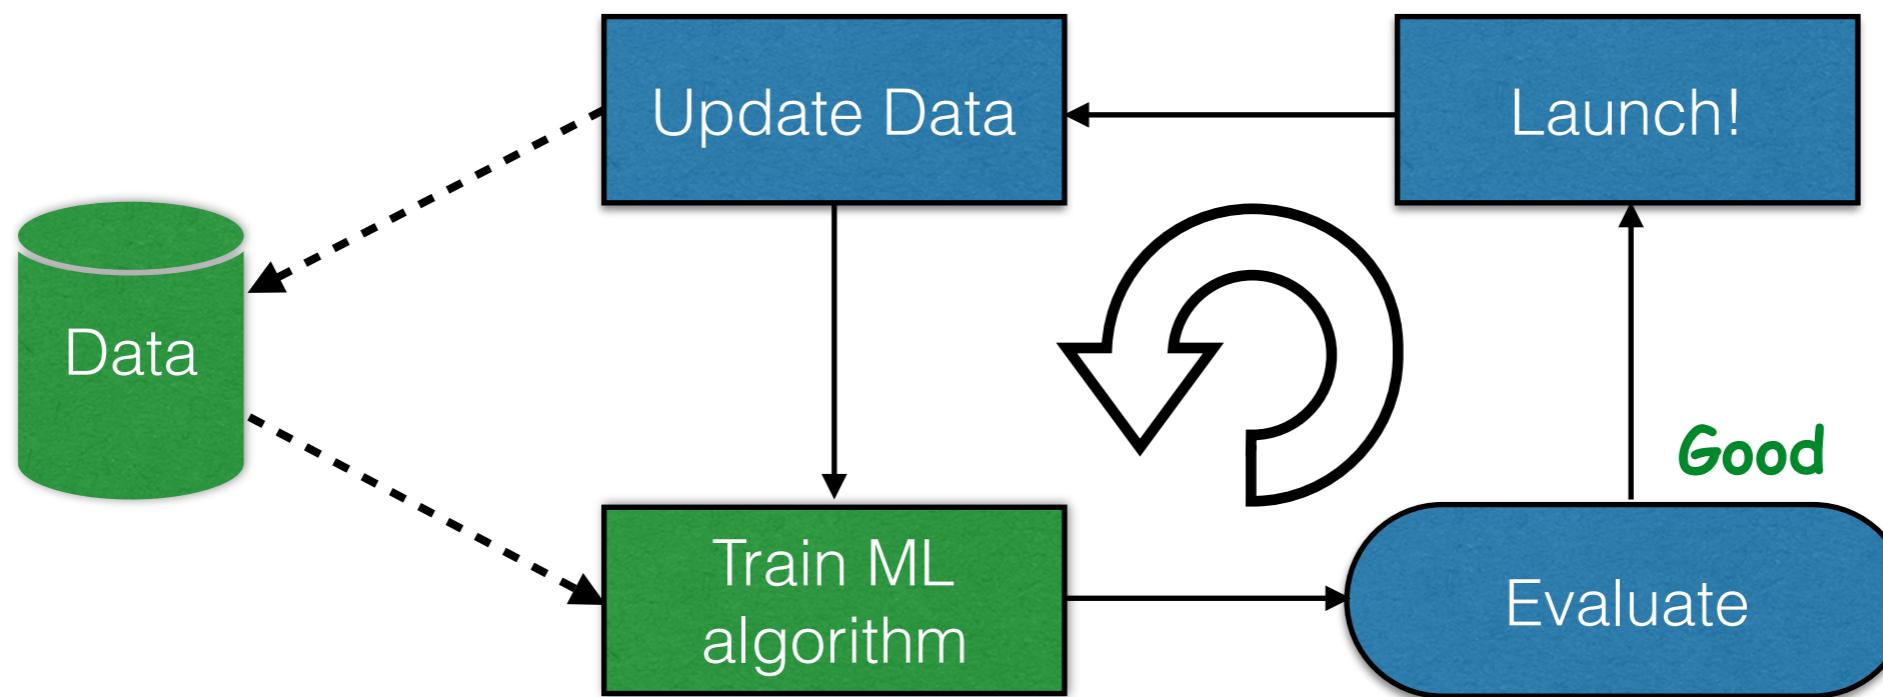
# Learning from Examples

- The machine automatically learns from examples
  - machine learning
  - no need to identify and explain rules
  - **general and flexible**



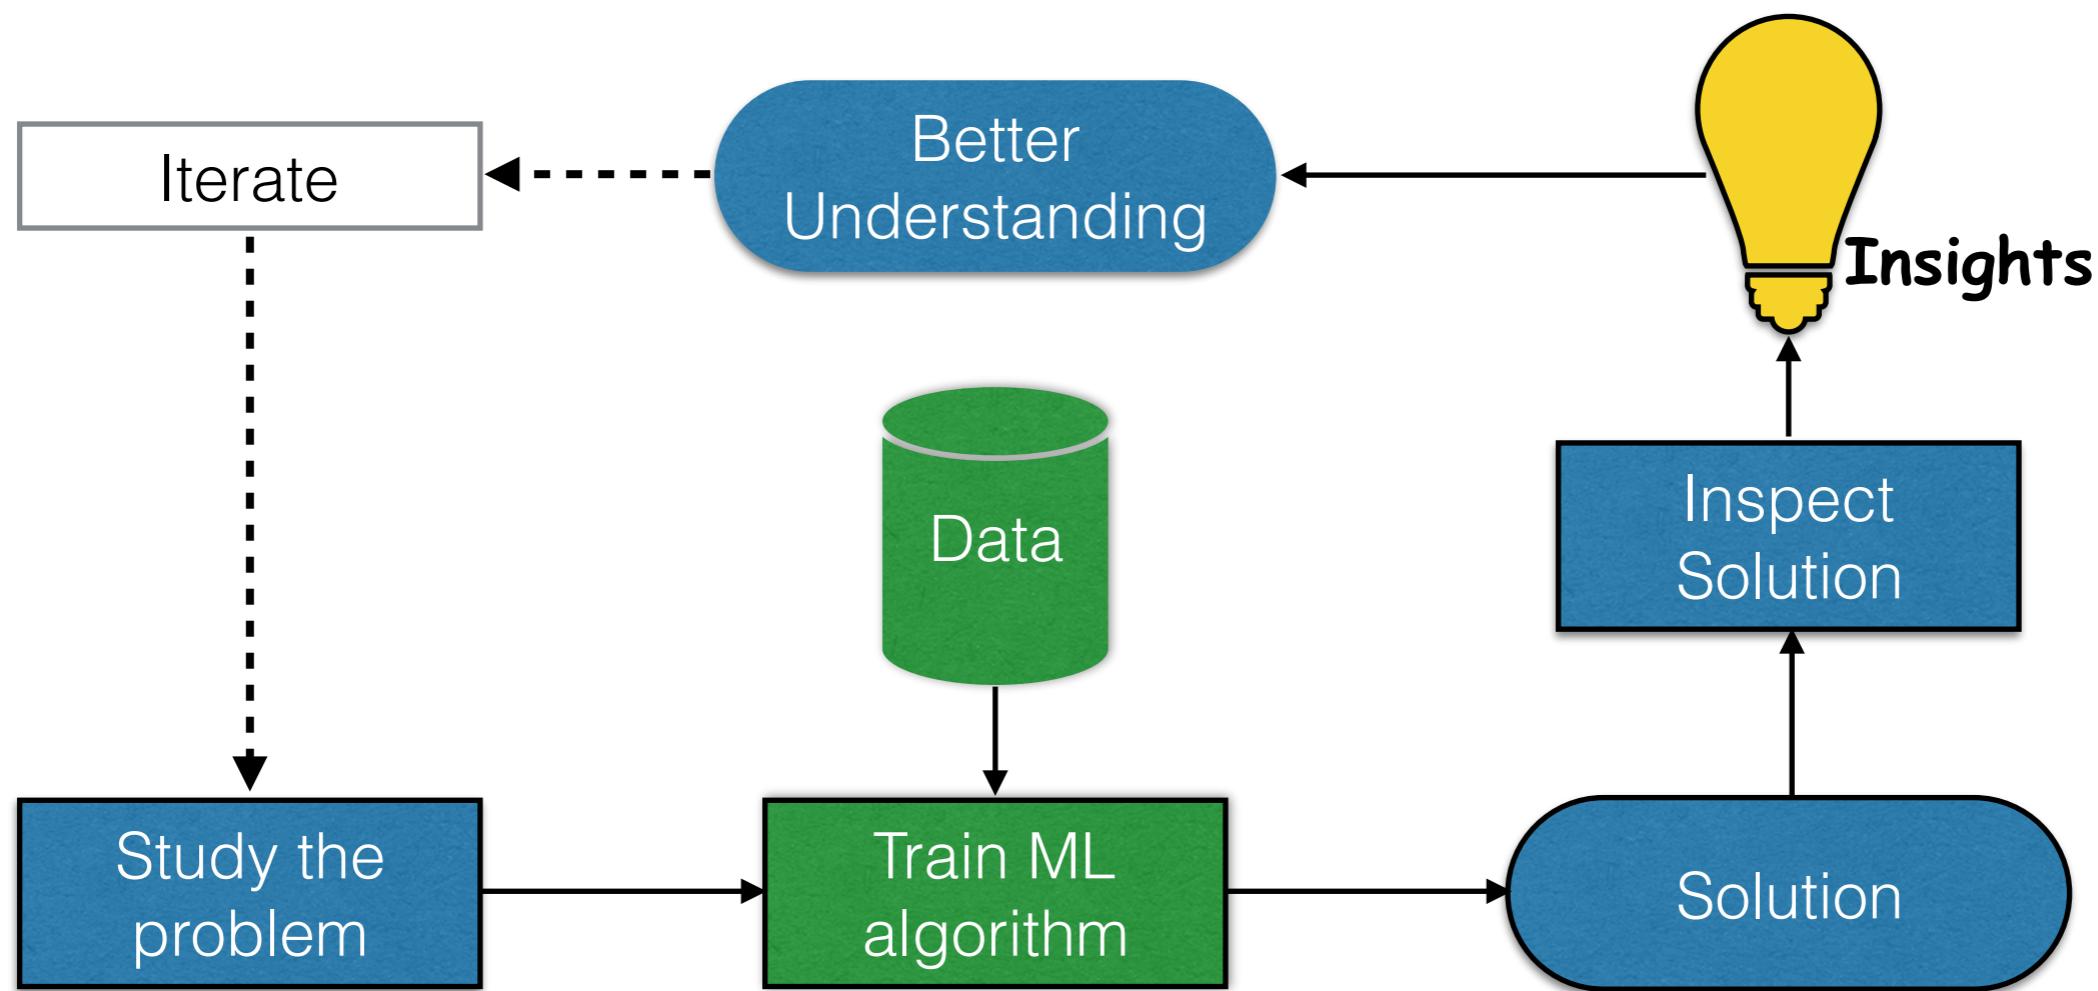
# Adapting to Change

- Machine Learning can automatically adapt to change
  - Simply update the data and train again
  - No need to change the underlying algorithm



# Help Humans Learn

- Machine Learning algorithms can be inspected
  - Might lead to new insights
  - Can uncover patterns in the data



# Features

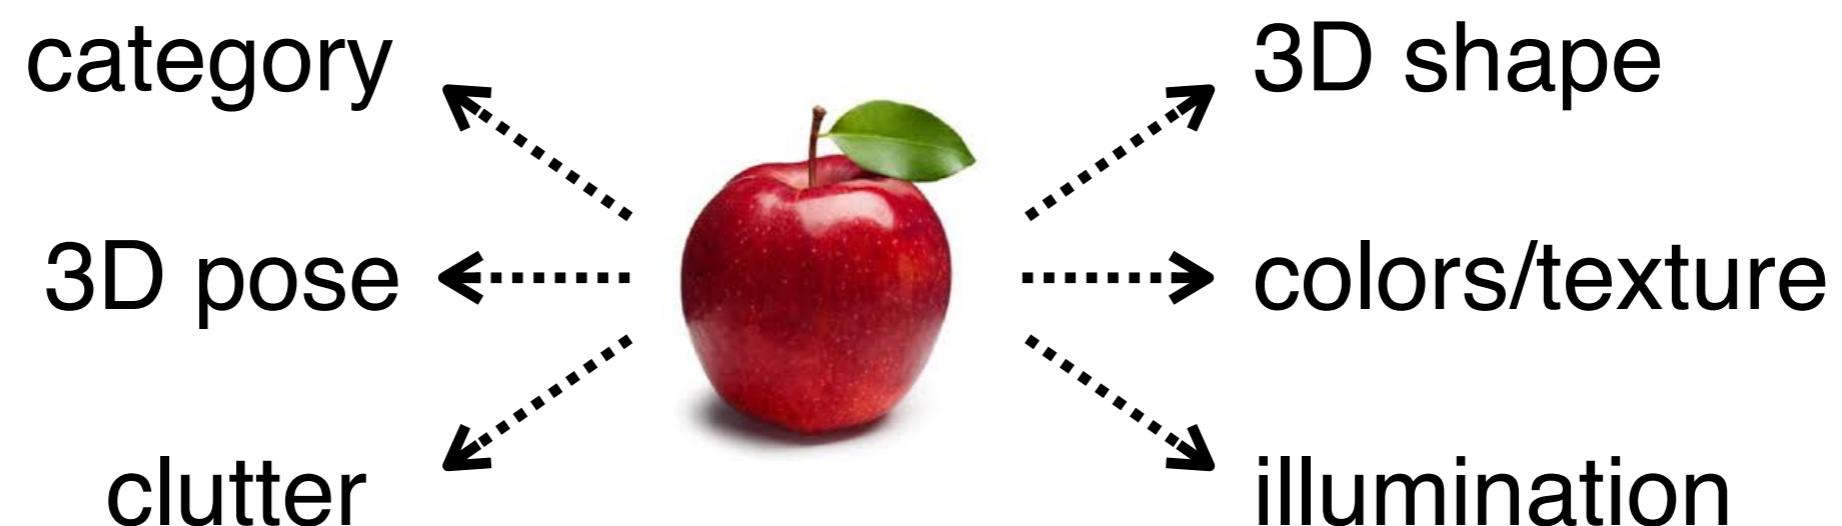
- Machines solve tasks/decisions by using the provided information (data)
- Data is often encoded into more focused relevant information (features) to simplify the decision

$$\text{data} \xrightarrow{x \rightarrow \phi(x)} \text{feature}$$

- Features can be hand-made/encoded
  - Operators often do not know the optimal features

# Representation Learning

- Features or, more in general, an **internal representation** or a hierarchy of concepts should be learned automatically
- The internal representation should separate all **factors of variation** (i.e., concepts that summarize important variation of the data)

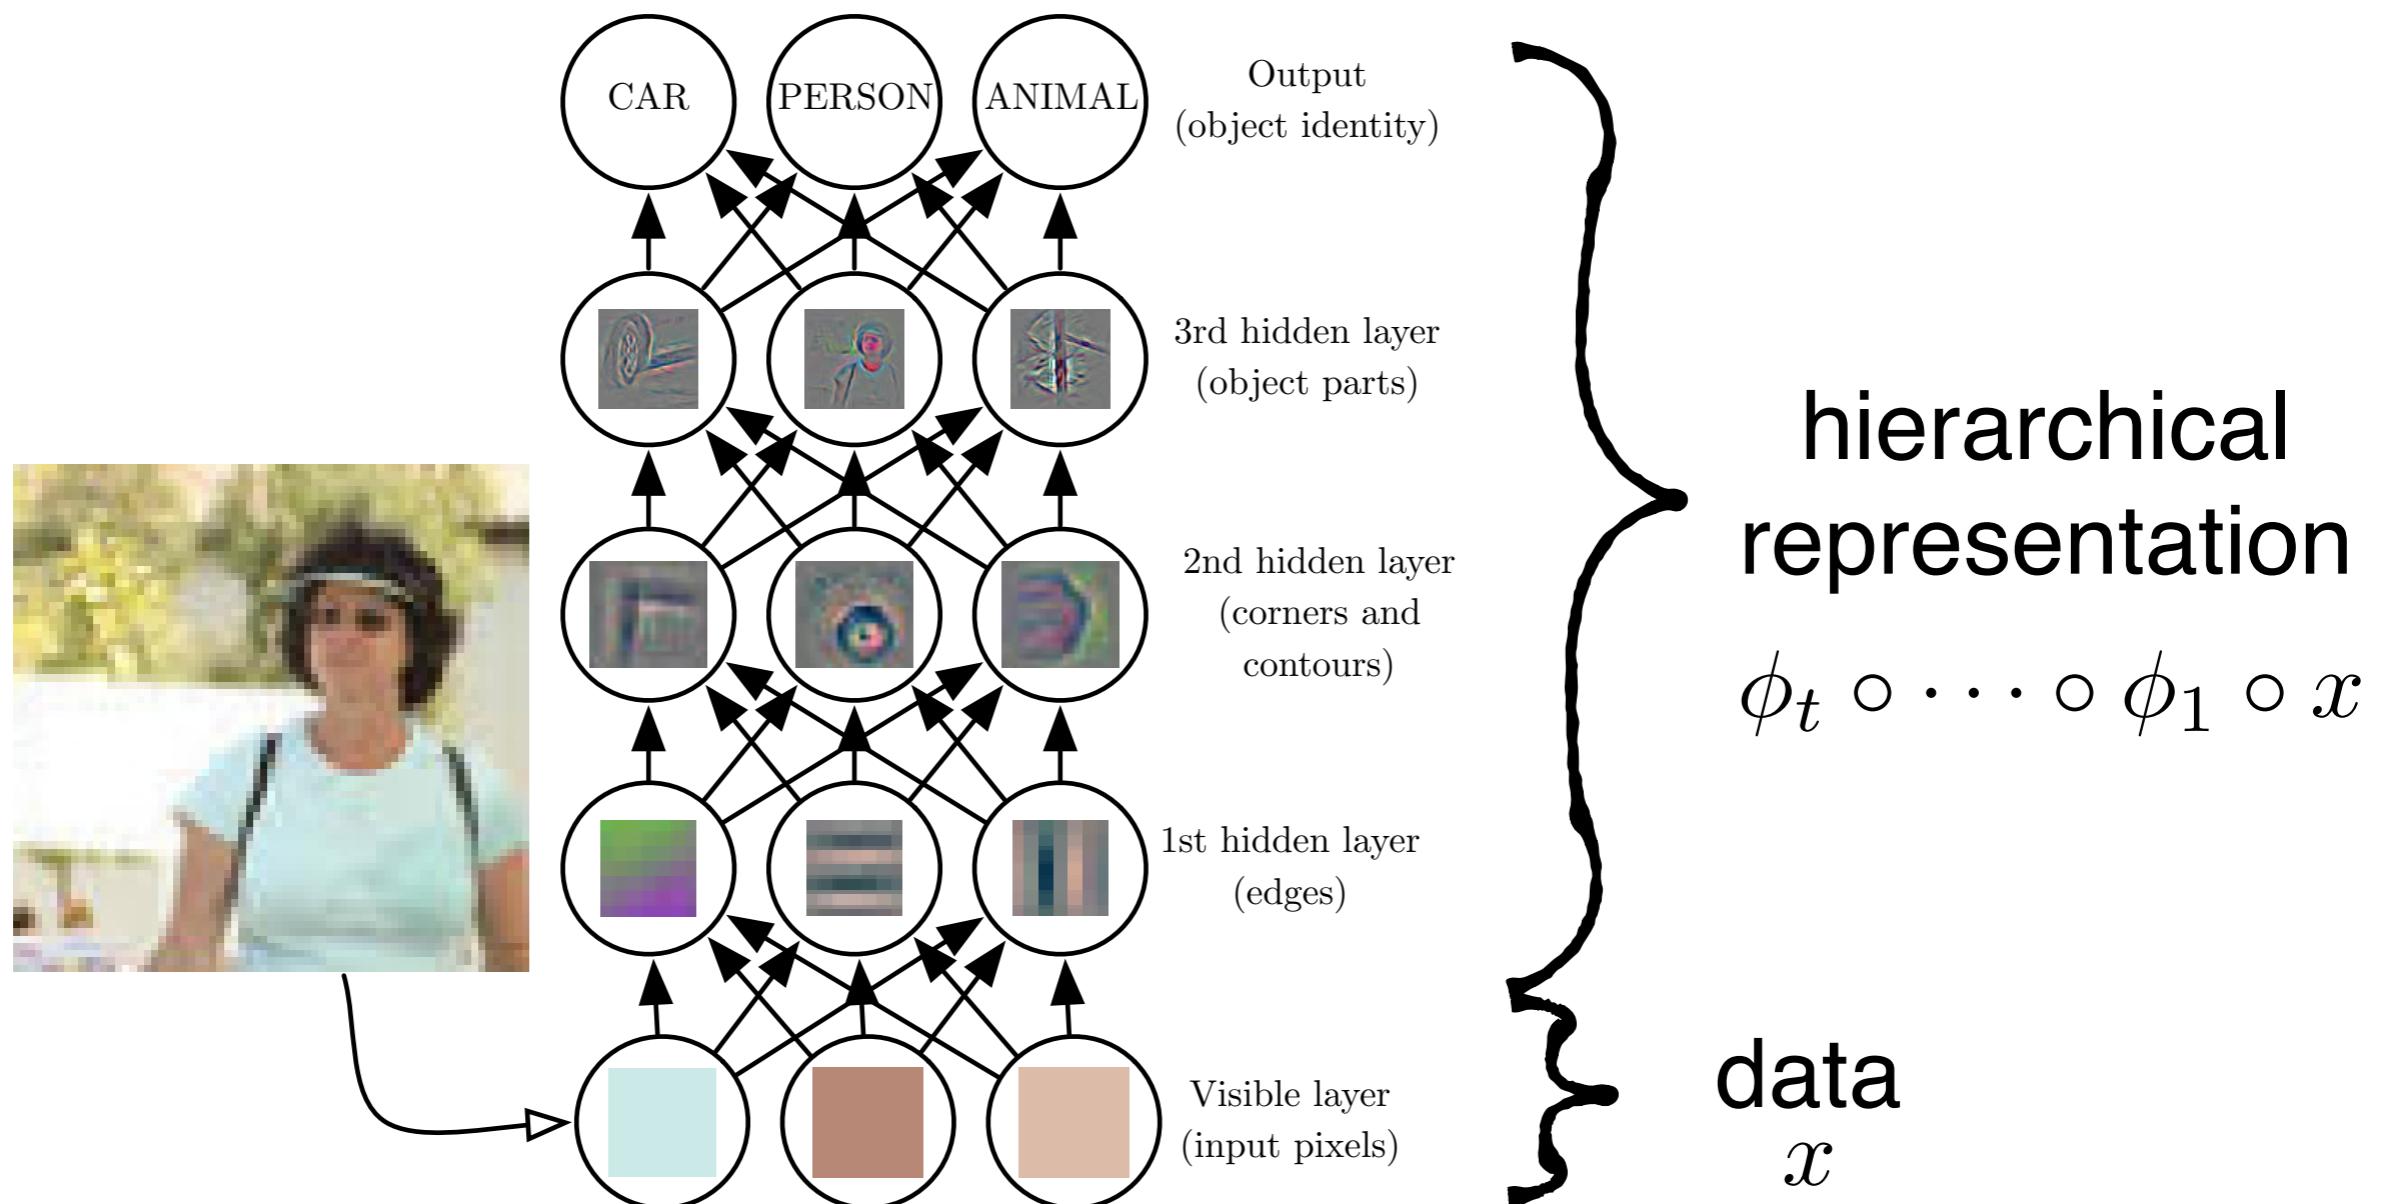


# Distributed Representation

- Use many features to represent data and each feature should handle multiple data samples
- Example: Recognition of cars, trucks and birds and each can be red, green or blue
- Case #1: 1 feature for each case  
( $3 \times 3 = 9$  features)
- Case #2: 3 features for identity and 3 for color  
( $3+3 = 6$  features)

# Deep Learning

Introduces hierarchical representations (from simple to complex, from low-level features to high-level features)



# Machine Learning Review

Simon Jenni  
(slides by Paolo Favaro)

# Contents

- Revision of basic concepts of Machine Learning
- Based on **Chapter 5** of Deep Learning by Goodfellow, Bengio, Courville

# Context

- A more complete introduction to Machine Learning through the following courses
  - Machine Learning @ UniBe
  - Machine Learning and Data Mining @ UniNe
  - Pattern Recognition @ UniFr
  - Statistical Learning Methods @ UniNe

# Resources

- Books and online material for further studies
  - Machine Learning @ Stanford (Andrew Ng)
  - **Pattern Recognition and Machine Learning**  
by Christopher M. Bishop
  - **Machine Learning: a Probabilistic Perspective**  
by Kevin P. Murphy

# Learning Pillars

- Supervised learning
- Semi-supervised learning
- Self-taught learning (unsupervised feature learning)
- Unsupervised learning (+self-supervised learning)
- Reinforcement learning

# Definition

- Mitchell (1997)

A computer program is said to learn from **experience  $E$**  with respect to some class of **tasks  $T$**  and **performance measure  $P$** , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .

# The Task T

- Example: if we want a robot to be able to walk, then **walking** is the task
- Approaches
  1. We could directly input **directives** for how we think a robot should walk, or
  2. We could provide **examples** of successful and unsuccessful walking (this is machine learning)

# The Task $T$

- Given an input  $x$  (e.g., a vector) produce a function  $f$ , such that  $f(x) = y$  (e.g., an integer, a probability vector)
- Examples
  - Classification
  - Regression
  - Machine translation
  - Denoising
  - Probability density estimation

# The Performance Measure P

- To evaluate a ML algorithm we need a way to measure how well it performs on the task
- It is measured on a separate set (**the test set**) from what we use to build the function  $f$  (**the training set**)
- Examples
  - Classification accuracy (portion of correct answers) or error rate (portion of incorrect answers)
  - Regression accuracy (e.g., least squares errors)

# The Experience E

- Specifies what data can be used to solve the task
- We can distinguish it based on the learning pillars
  - **Supervised**: data is composed of both the input  $x$  (e.g., features) and output  $y$  (e.g., labels/targets)
  - **Unsupervised**: data is composed of just  $x$ ; here we typically aim for  $p(x)$  or a method to sample  $p(x)$
  - **Reinforcement**: data is dynamically gathered based on previous experience

# Data

- We assume that all collected data samples in all datasets:
  1. come from the same distribution  $\rightarrow p_{x^{(i)}}(x) = p_{x^{(j)}}(x)$
  2. are independent  $\rightarrow p(x^{(1)}, \dots, x^{(m)}) = \prod_{i=1}^m p(x^{(i)})$
- This assumption is denoted **IID** (independent and identically distributed)

# Example: Linear Regression

- Given IID data inputs  $x \in \mathbb{R}^n$  and outputs  $y \in \mathbb{R}$
- **Task T:** predict  $y$  with the linear regressor  $\hat{y} = w^\top x$   
need to find the weights  $w$
- **Experience E:** training set  $X^{\text{train}} \in \mathbb{R}^{m \times n}$ ,  $Y^{\text{train}} \in \mathbb{R}^m$   
and test set  $X^{\text{test}} \in \mathbb{R}^{m \times n}$ ,  $Y^{\text{test}} \in \mathbb{R}^m$
- **Performance P:** Mean squared error

$$\text{MSE}^{\text{test}}(w) = \frac{1}{m} |X^{\text{test}}w - Y^{\text{test}}|^2$$

# Linear Regression

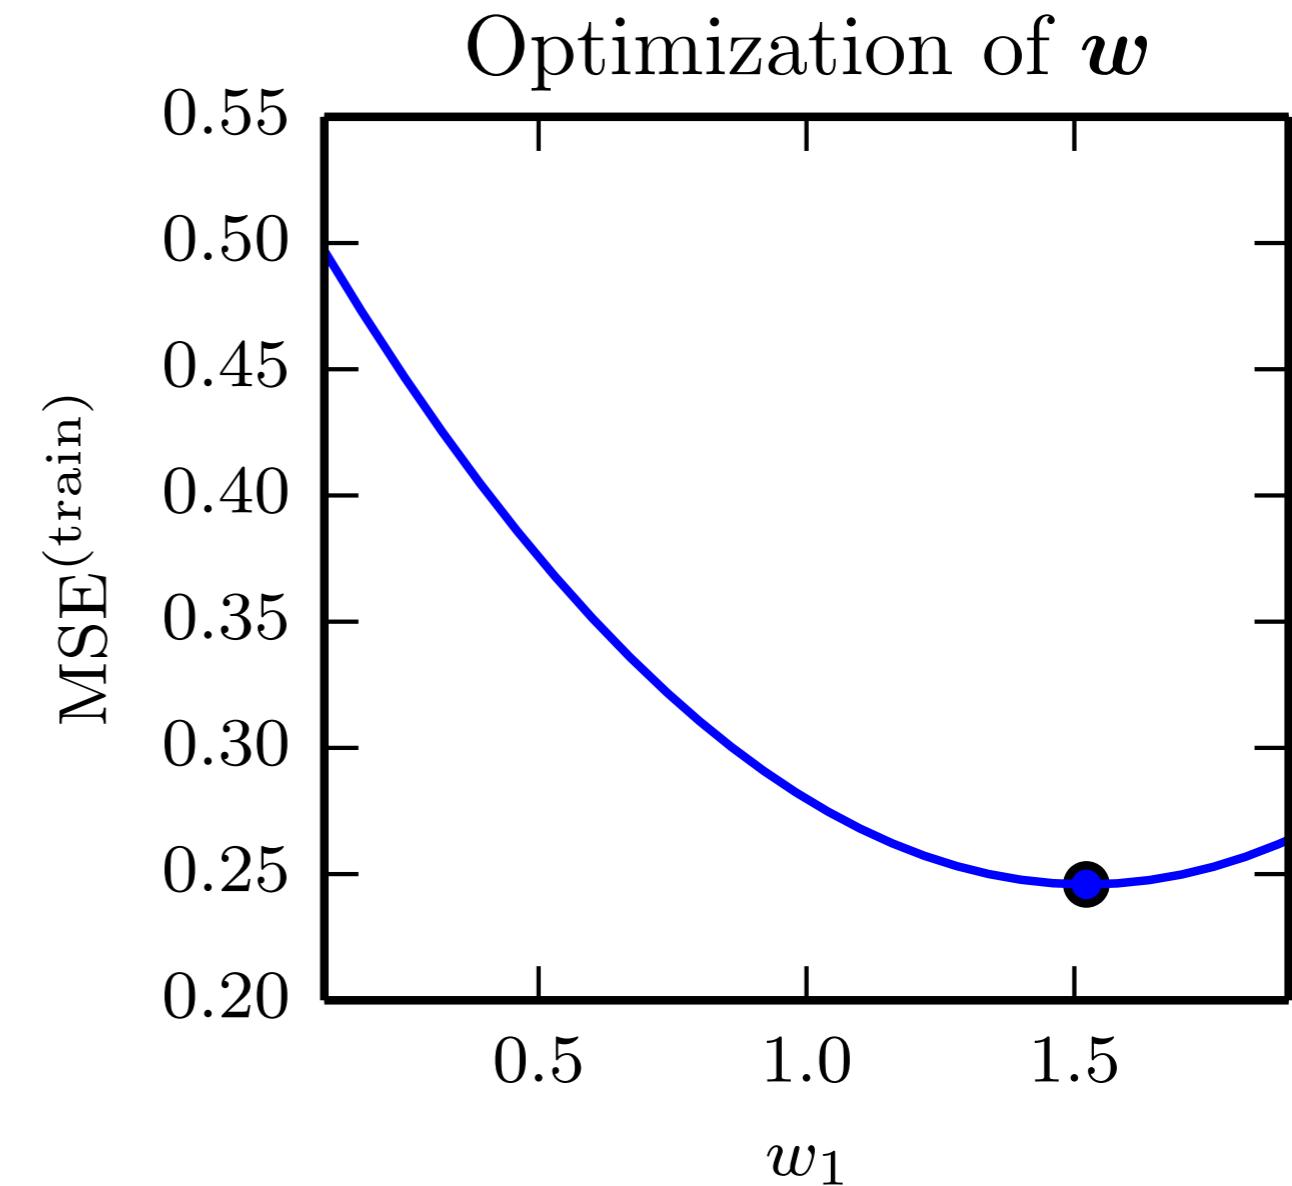
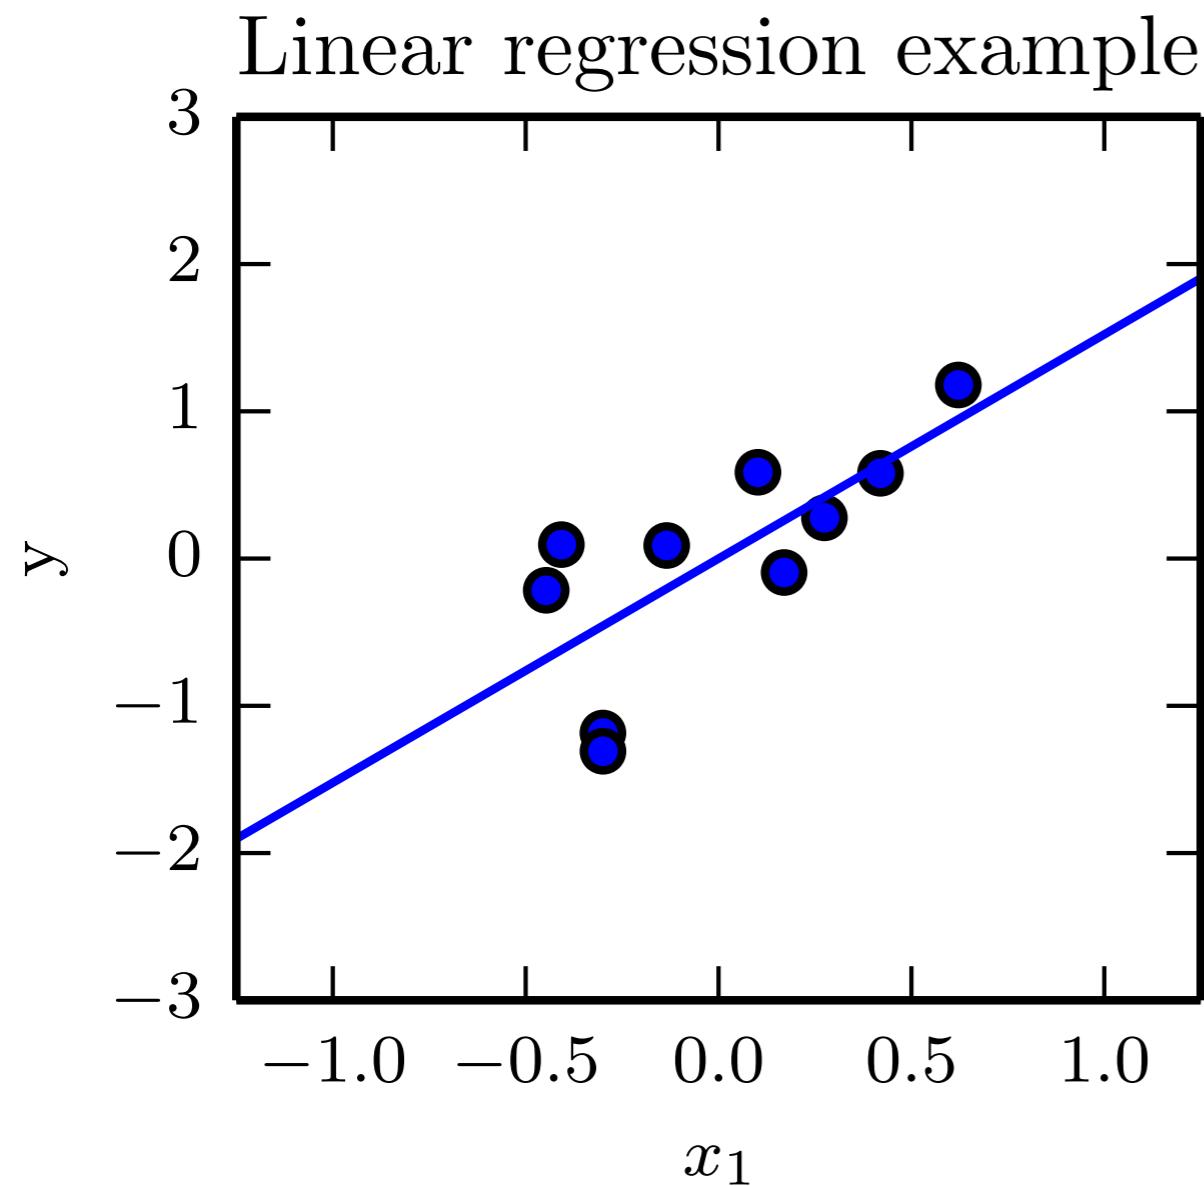
- Solve task T by minimizing the  $\text{MSE}^{\text{train}}$

$$\text{MSE}^{\text{train}}(w) = \frac{1}{m} |X^{\text{train}}w - Y^{\text{train}}|^2$$

- Compute the gradient of  $\text{MSE}^{\text{train}}(w)$  with respect to  $w$  and set to 0 (normal equations)
- The solution is (pseudo-inverse)

$$w = (X^{\text{train}}{}^\top X^{\text{train}})^{-1} X^{\text{train}}{}^\top Y^{\text{train}}$$

# Linear Regression



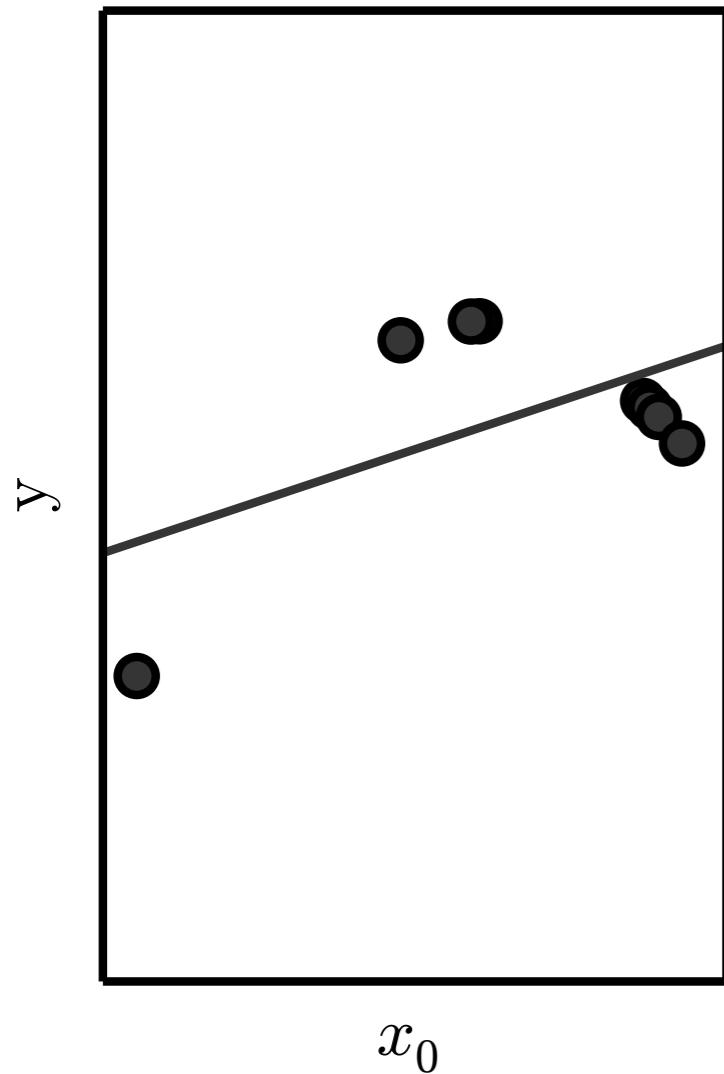
# Overfitting and Underfitting

- Performance  $P$  captures how well the learned model predicts new unseen data
- Ideally we want to select the predictor with the best performance
- What happens when we use predictors of different complexity/capacity?

# Overfitting and Underfitting

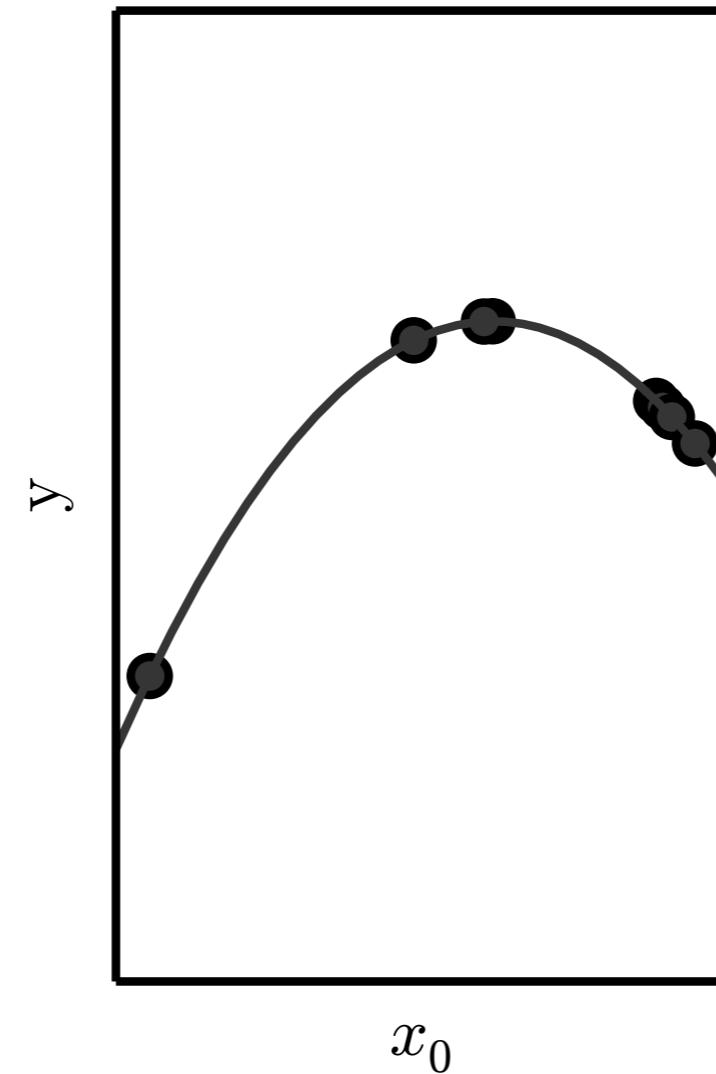
shown data is the training set

Underfitting



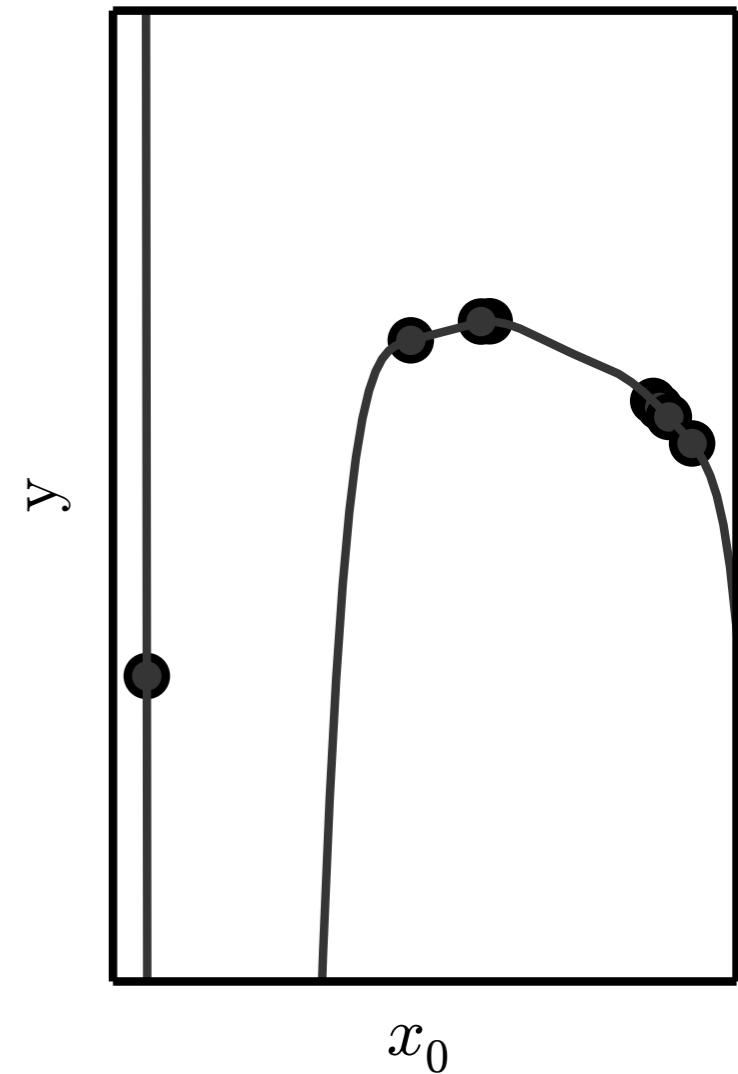
simple predictor

Appropriate capacity



optimal predictor

Overfitting



complex predictor

# Loss function

- Define a **predictor** function  $f : \mathcal{X} \mapsto \mathcal{Y}$
- Define a **loss** function  $l : \mathcal{Y} \times \mathcal{Y} \mapsto \mathbb{R}$  which measures how different the two inputs are

- Examples

- 0-1 loss

$$l(y, f(x)) = \begin{cases} 0 & \text{if } y = f(x) \\ 1 & \text{if } y \neq f(x) \end{cases}$$

- Quadratic loss

$$l(y, f(x)) = (y - f(x))^2$$

# Bayes Risk

- **Bayes risk** is defined as (average loss)

$$R(f) = E_{x,y}[l(f(x), y)] = \int l(f(x), y)p(x, y)dxdy$$

- The optimal predictor function is

$$f^* = \arg \min_f R(f)$$

# Empirical Risk

- Given  $(x_i, y_i)$  with  $i = 1, \dots, m$  the **empirical risk** is

$$\hat{R}(f) = \frac{1}{m} \sum_{i=1}^m l(f(x_i), y_i)$$

- The empirical predictor is

$$\hat{f} = \arg \min_{f \in \mathcal{F}} \hat{R}(f)$$

# Risks

- Bayes risk

$$R(f^*) = E_{x,y}[l(f^*(x), y)]$$

- Empirical risk

$$\hat{R}(\hat{f}) = \frac{1}{m} \sum_{i=1}^m l(\hat{f}(x_i), y_i)$$

- Bayes risk restricted to function family

$$\min_{f \in \mathcal{F}} R(f)$$

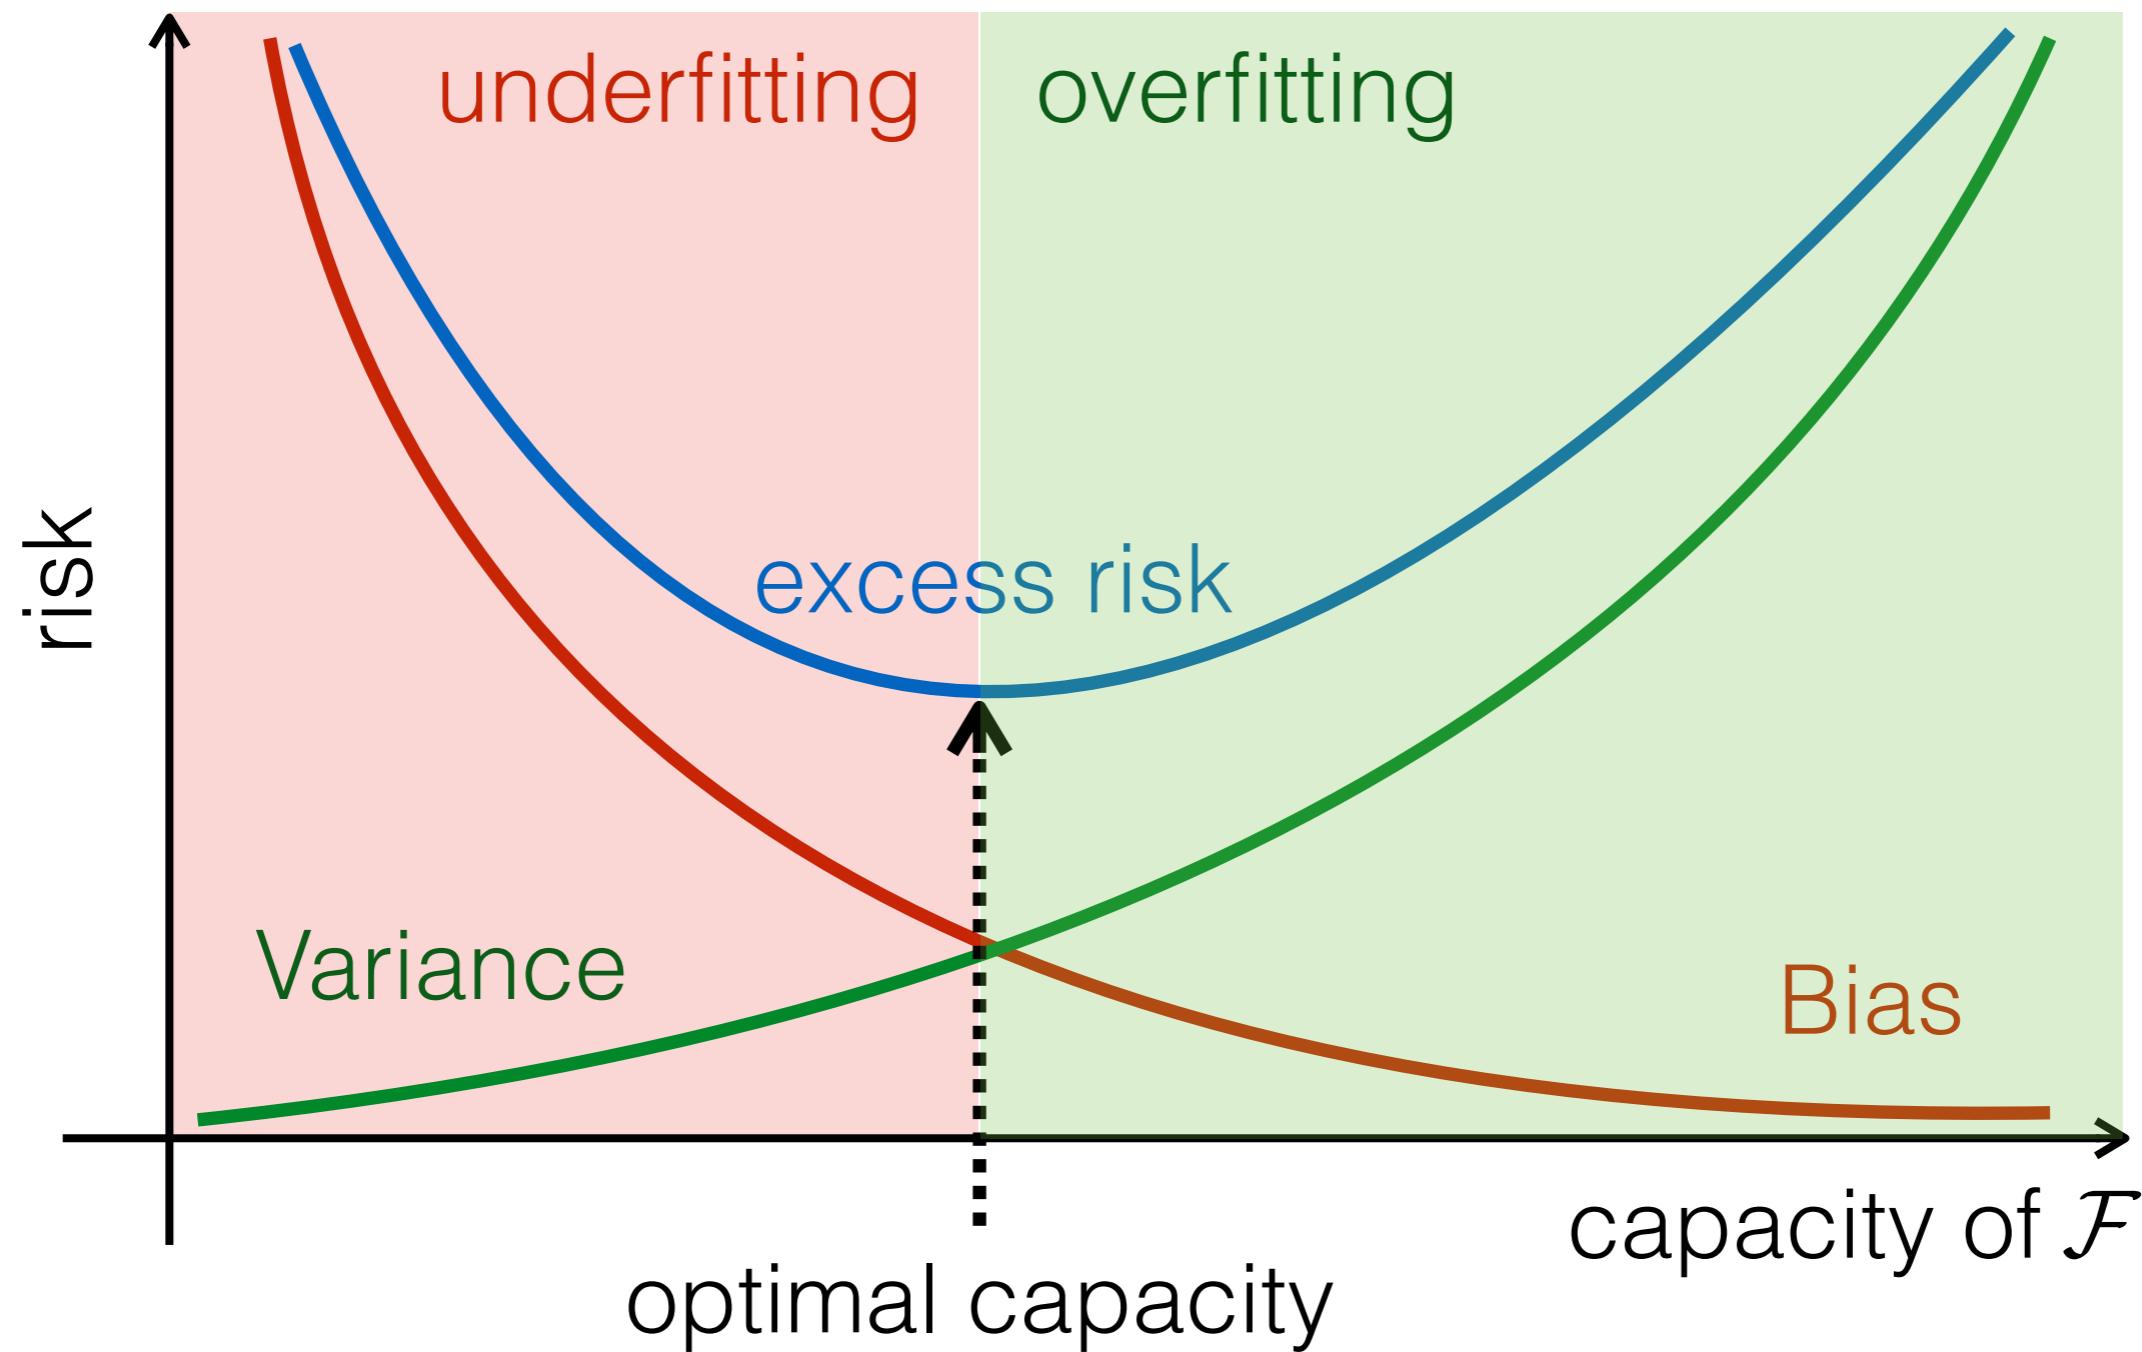
# Estimation vs Approximation

- The **excess risk** is the gap between the empirical risk and the optimal Bayes risk

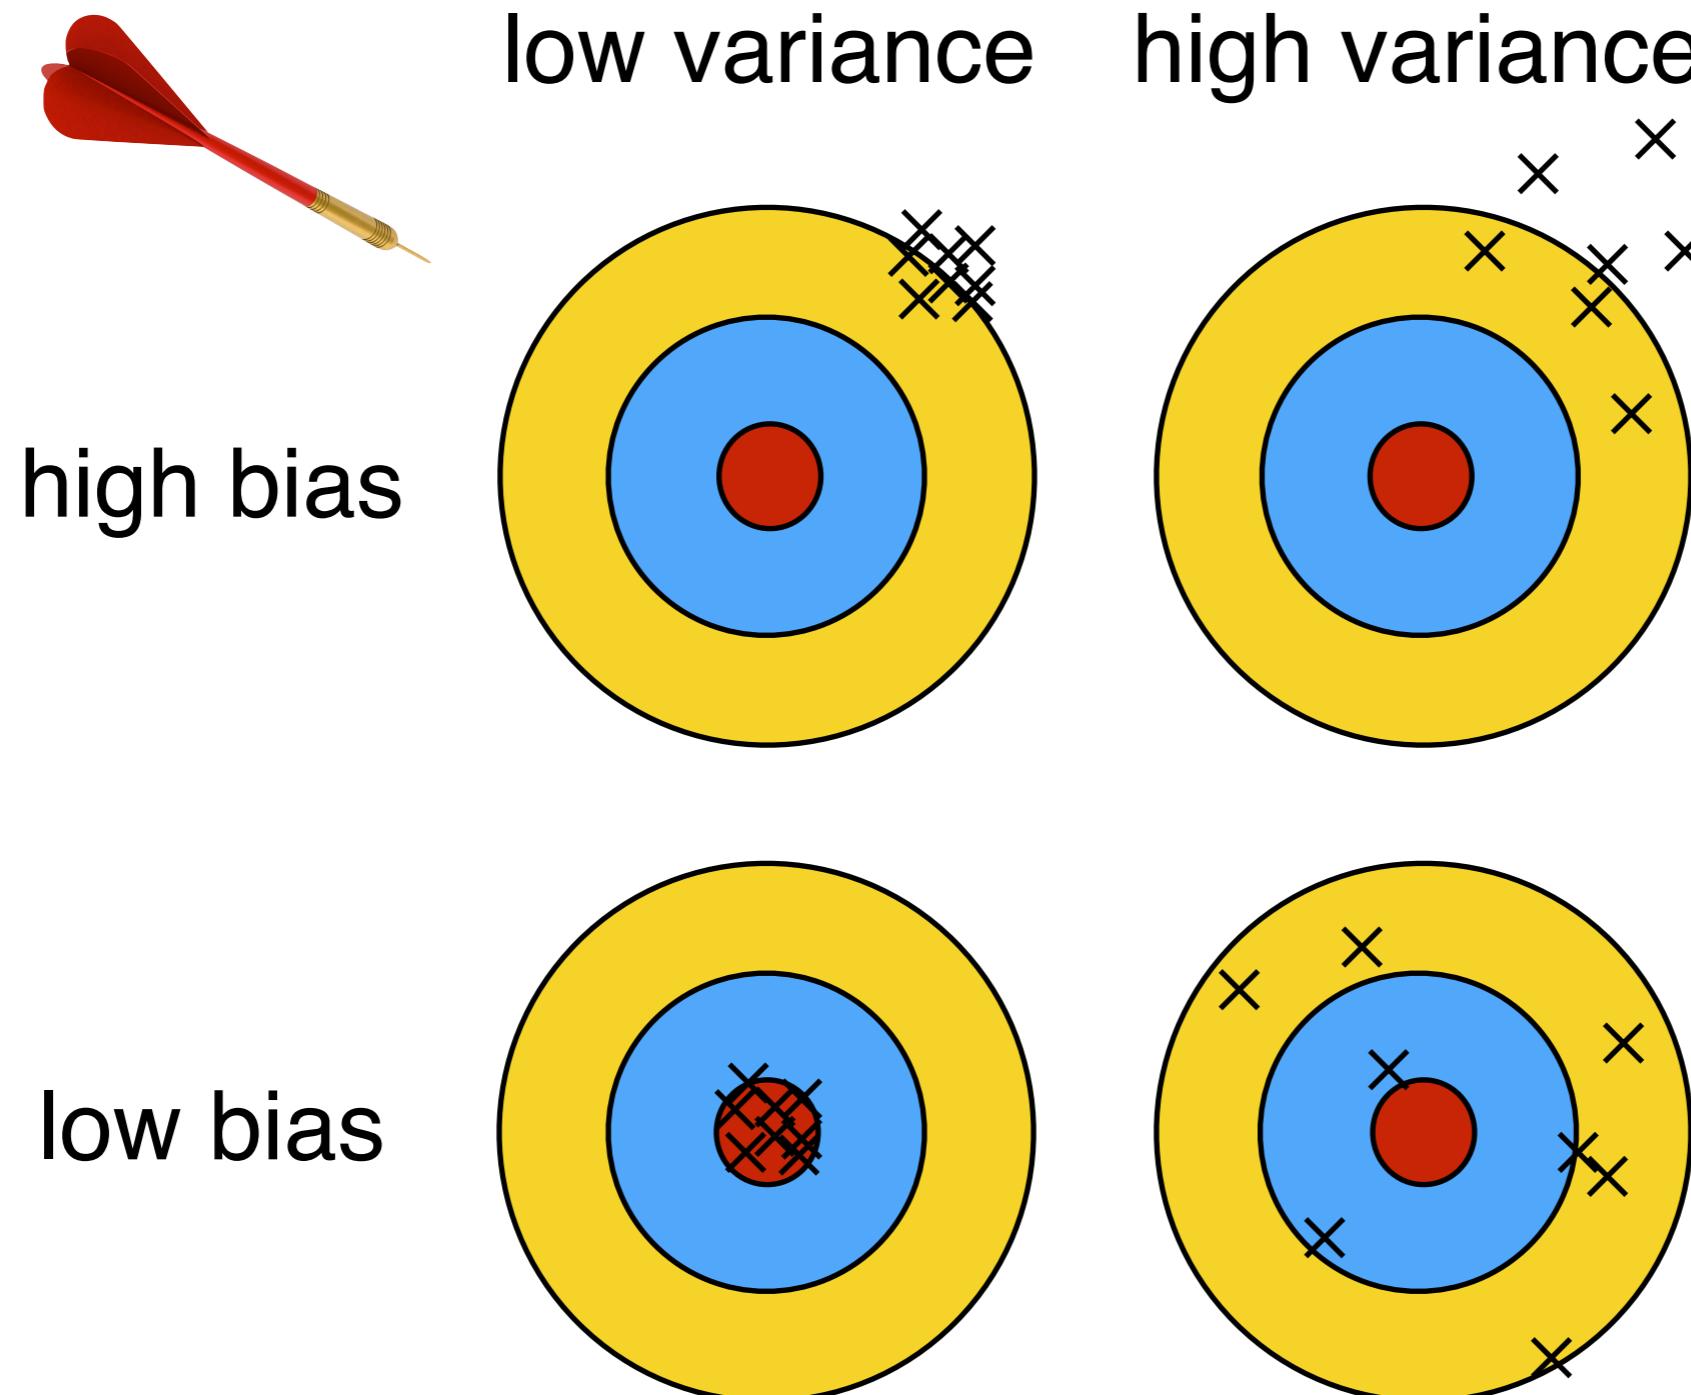
$$\hat{R}(\hat{f}) - R(f^*) = \underbrace{\hat{R}(\hat{f}) - \min_{f \in \mathcal{F}} R(f)}_{\text{estimation error}} + \underbrace{\min_{f \in \mathcal{F}} R(f) - R(f^*)}_{\text{approximation error}}$$

- Estimation (variance)**: due to training set
- Approximation (bias)**: due to function family  $\mathcal{F}$

# Estimation vs Approximation



# Bias and Variance





**THE BEST WAY TO  
EXPLAIN OVERFITTING**

# Regularization

- Define a parametric family  $\mathcal{F}_\lambda$  of functions, where  $\lambda$  regulates the complexity/capacity of the predictors
- Given the optimal predictor from the empirical risk

$$\hat{f}_\lambda = \arg \min_{f \in \mathcal{F}_\lambda} \hat{R}(f)$$

we would like to choose the capacity based on Bayes risk

$$R(\hat{f}_\lambda)$$

# Training, Validation and Test

- In alternative, collect samples into training set  $D_{\text{train}}$ , validation set  $D_{\text{val}}$  and test set  $D_{\text{test}}$
- Use the **training set** to define the optimal predictor

$$\hat{f}_\lambda = \arg \min_{f \in \mathcal{F}} \hat{R}_{D_{\text{train}}}(f)$$

- Use the **validation set** to choose the capacity

$$\hat{\lambda} = \arg \min_{\lambda} \hat{R}_{D_{\text{val}}}(\hat{f}_\lambda)$$

- Use the **test set** to evaluate the performance

$$\text{performance } P = R_{D_{\text{test}}}(\hat{f}_{\hat{\lambda}})$$

# Supervised Learning

- Make a prediction of an output  $y$  given an input  $x$
- Boils down to determining the conditional probability

$$p(y|x)$$

- Formulate problem as that of finding  $\theta$  for a parametric family (Maximum Likelihood)

$$p(y|x; \theta)$$

# Maximum Likelihood

- Given IID input/output samples  $(x^i, y^i) \sim p_{\text{data}}(x, y)$

the **conditional maximum likelihood** estimate is

$$\begin{aligned}\theta_{\text{ML}} &= \arg \max_{\theta} \prod_{i=1}^m p_{\text{data}}(y^i | x^i; \theta) \\ &= \arg \max_{\theta} \sum_{i=1}^m \log p_{\text{data}}(y^i | x^i; \theta)\end{aligned}$$

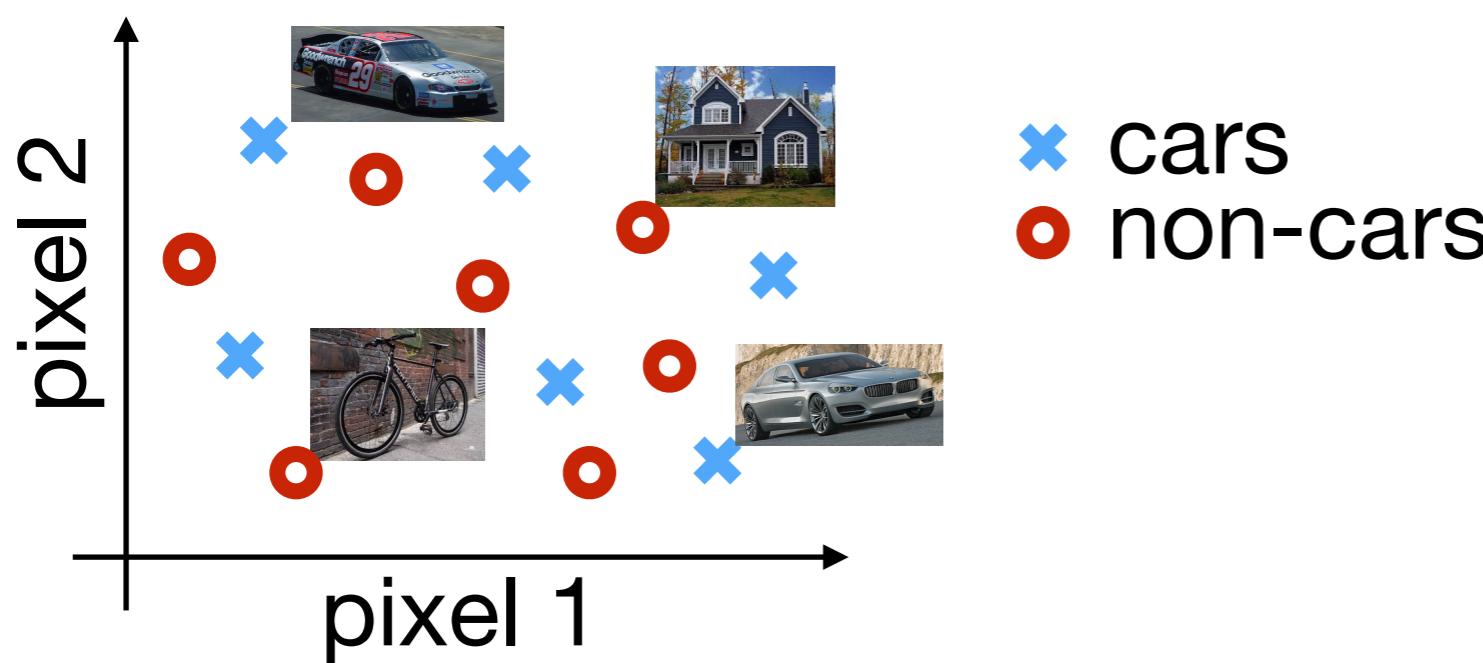
# Logistic Regression

- **Example:** Binary classification  $y \in \{0, 1\}$
- We aim at determining  $p(y = 1|x; \theta) = \sigma(\theta^\top x)$   
where  $\sigma(z) = \frac{1}{1 + e^{-z}}$  is the sigmoid function
- Class  $y=1$  can be picked when

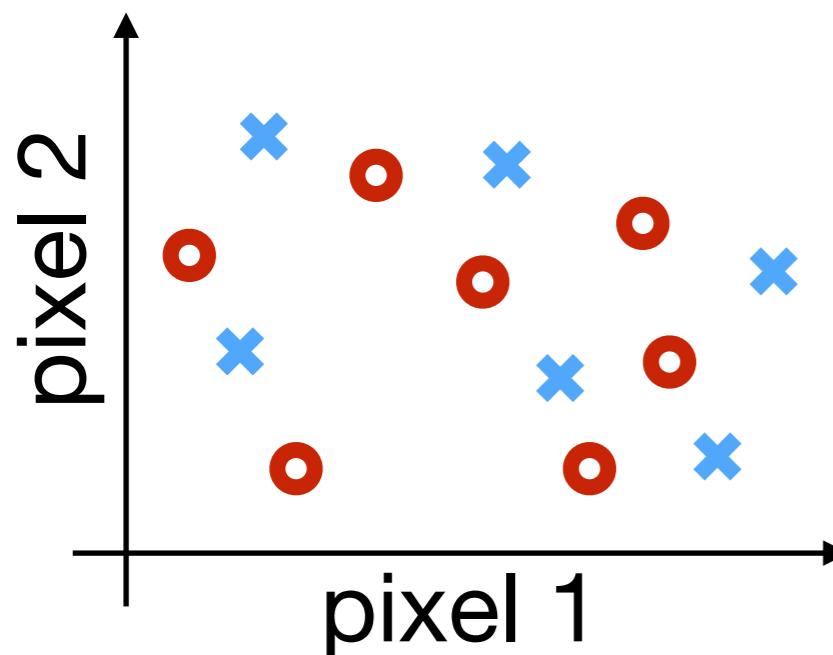
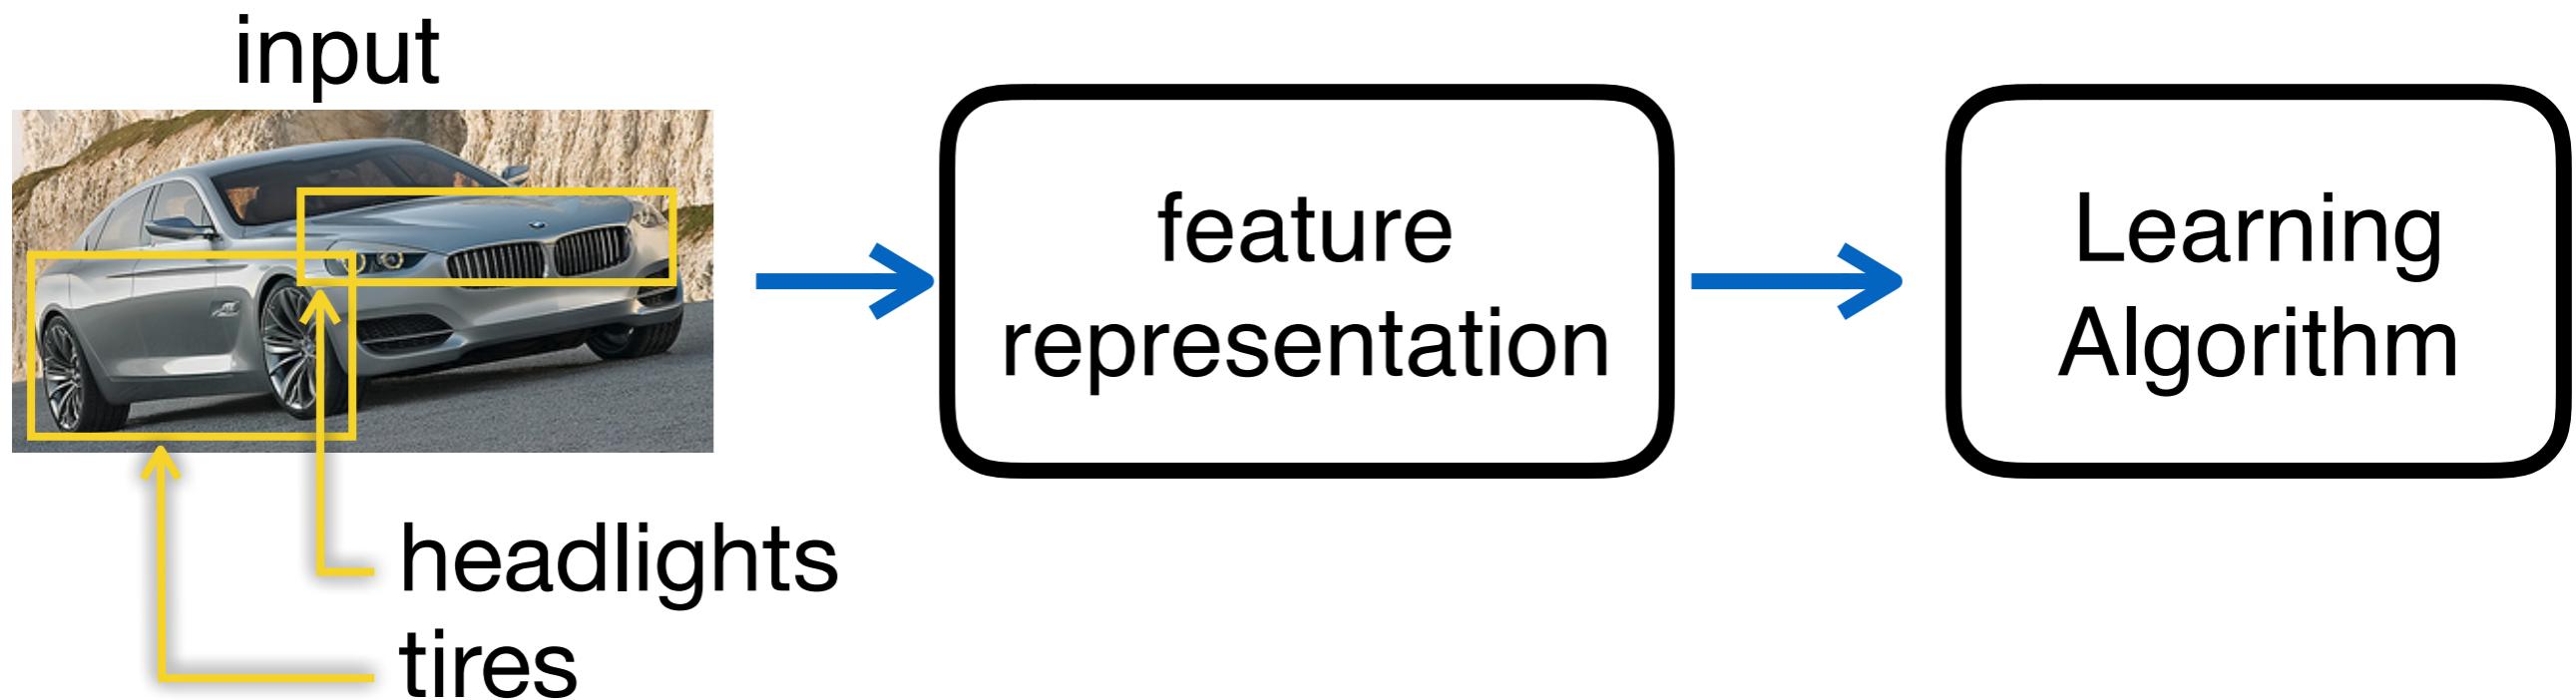
$$p(y = 1|x; \theta) > p(y = 0|x; \theta)$$

which is equivalent to  $\theta^\top x > 0$

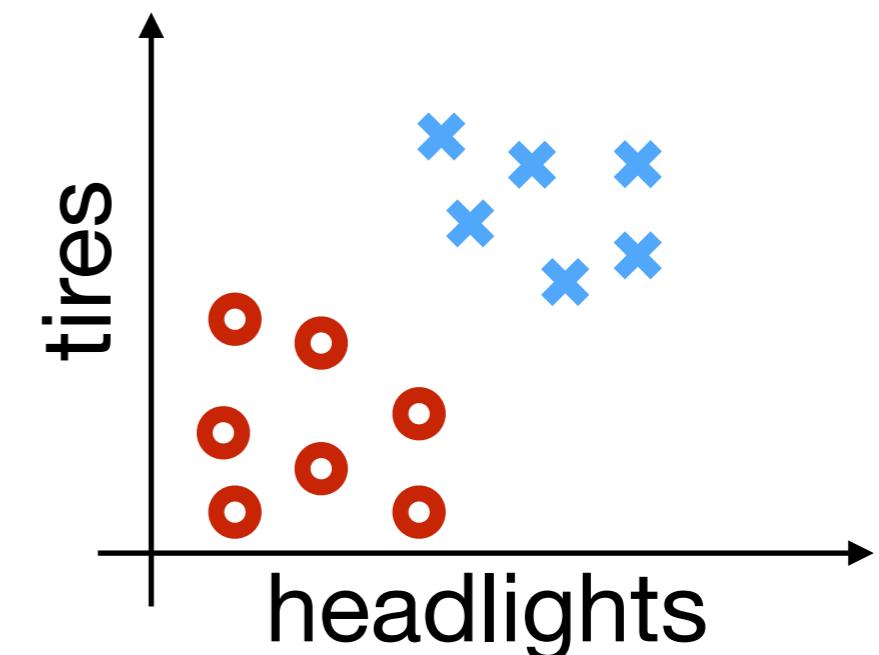
# Features



# Features



cars  
non-cars



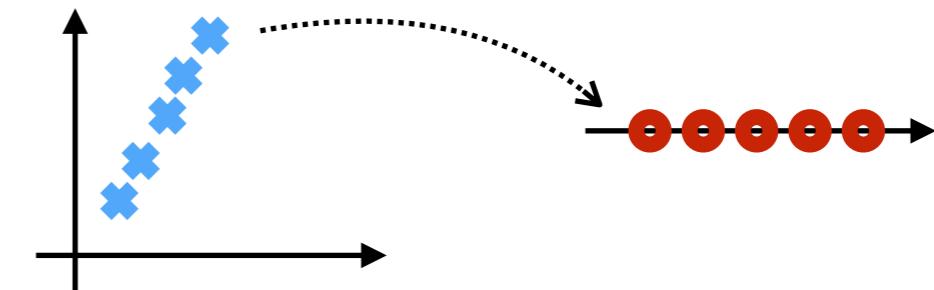
# Unsupervised Learning

- Aim is to find a suitable **data representation**
  - Probability density estimator
  - Sampling procedure
  - Data denoising
  - Manifold learning
  - Clustering

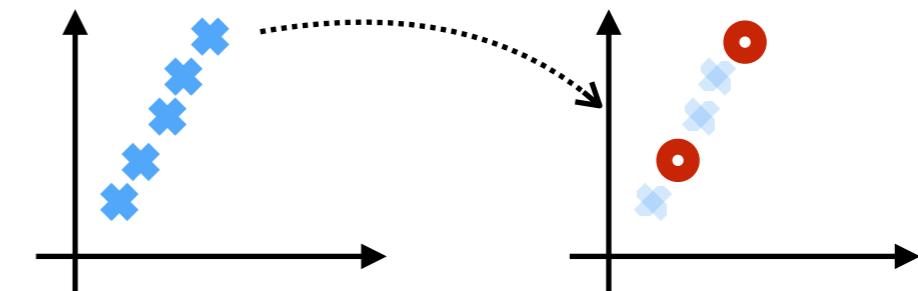
# Data Representation

- The ideal data representation should:
  1. **Preserve** all task-relevant information
  2. Be **simpler** than the original data and **easier** to use

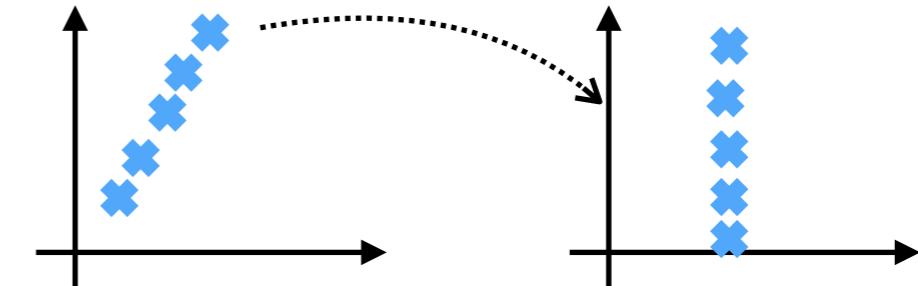
(i) low-dimensional



(ii) sparse

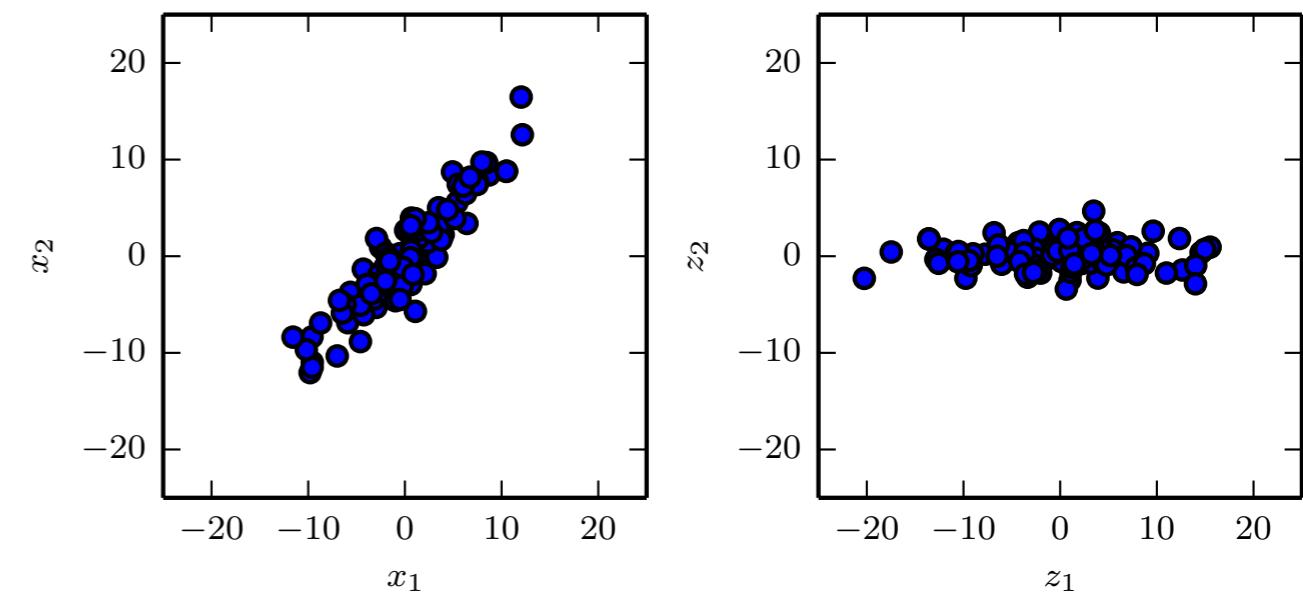


(iii) independent



# Principal Components Analysis

- **Definition:** Project data  $X$  so that the largest variation of the projected data  $Z = U^\top X$  is axis-aligned



$$X = U\Sigma V^\top$$

$U^\top U \doteq I$        $\Sigma \doteq \begin{bmatrix} \sigma_0 & 0 & \cdots & 0 \\ 0 & \sigma_1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_n \end{bmatrix}$        $V^\top V \doteq I$

**singular values**  $\longrightarrow \sigma_0 \geq \sigma_1 \geq \cdots \geq \sigma_n \geq 0$

# Principal Components Analysis

- Unsupervised learning method for **linearly** transformed data
- A low-dimensional representation (by thresholding the singular values)
- Yields independent (uncorrelated) components

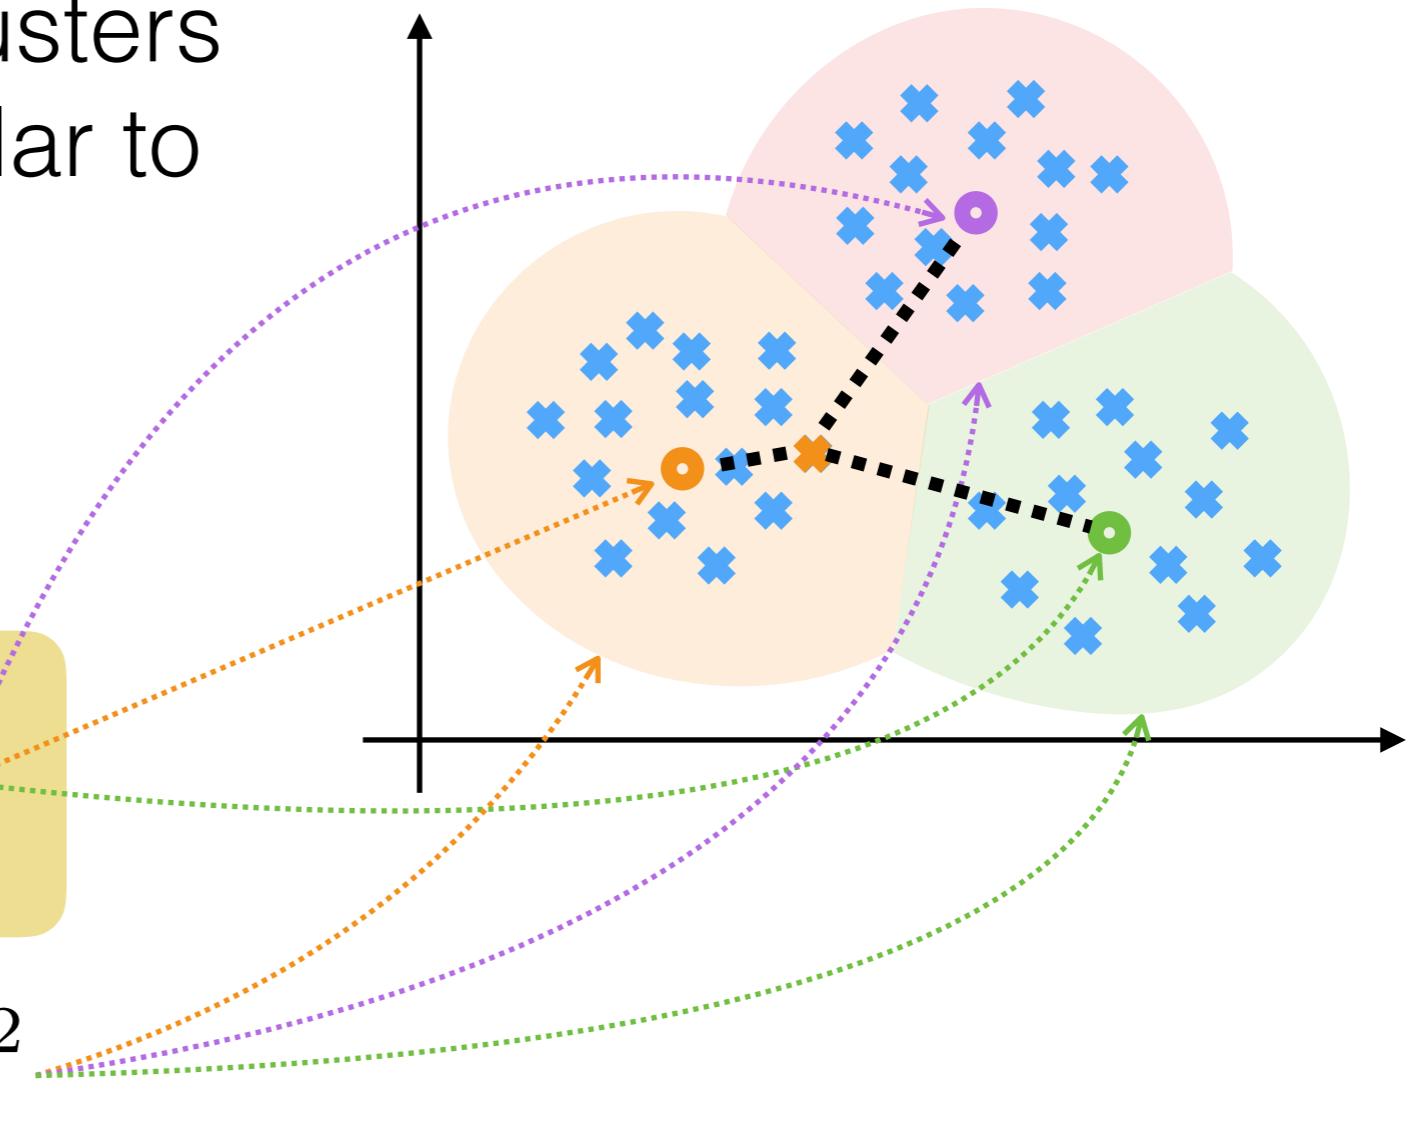
# K-Means Clustering

- **Definition:** Find  $k$  clusters of data samples similar to each other

Alternate between:

$$c_j = \frac{\sum_i \delta[w_i = j]x_i}{\sum_i \delta[w_i = j]}$$

$$w_i = \arg \min_j |x_i - c_j|^2$$



# K-Means Clustering

- Unsupervised learning method (handles nonlinearly transformed data)
- A sparse representation (assignments  $w_i$  encode one sample with one of the cluster centers  $c_j$ )
- Depends on initialization
- Ill-posed (multiple solutions can be valid)
- Number of clusters is usually unknown

# Conclusion

- Machine Learning is about making computers better at some task by learning from data
- Many different ML systems:
  - Supervised (regression, classification, ...)
  - Unsupervised (clustering, dim. reduction, ...)
- We maximize the model likelihood over the training set and hope it will generalise to unseen data
- Data is important (garbage in, garbage out)!  
Model complexity should fit the data.

Thank you for your attention!

Questions?