

# Fundamentals in Digital Image processing

Lectures on Advanced Microscopy

Guillaume Witz, PhD  
Science IT Support - MIC

# Why do we need image processing?



We are easily fooled and tend to see patterns everywhere.

Reproducibility requires a standardized data analysis

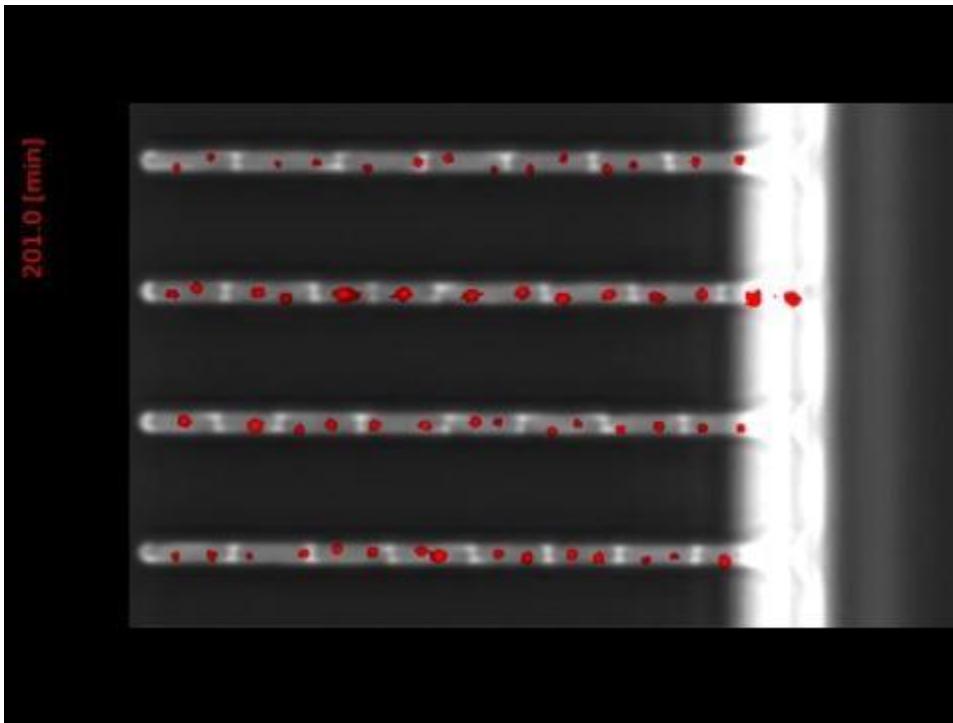


<https://www.flickr.com/photos/eworm/3401596917>

# Why do we need image processing?



# Why do we need image processing?



10 positions  
24 hours  
3 min intervals  
3 wavelengths

==

14'400 images

Modern microscopy methods and the trend towards quantitative biology make computational approaches necessary.

E. coli growing in microfluidics channels.  
Tagged origin of replication.

# Tools

One can do 99% of image processing tasks using open source tools. Two ecosystems are currently dominating:



Fiji (Java)

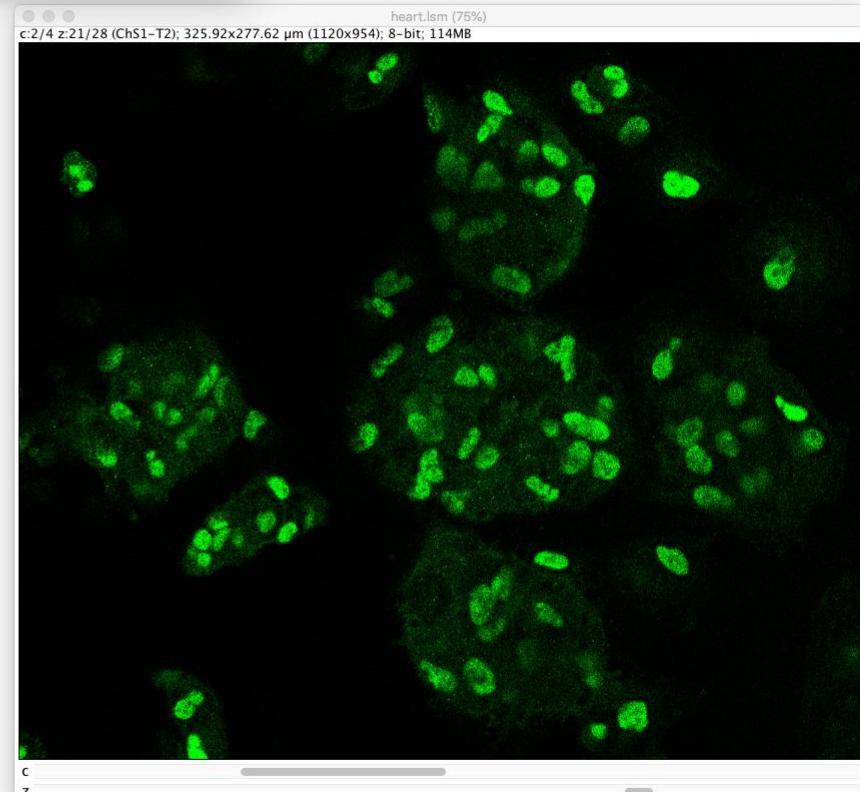
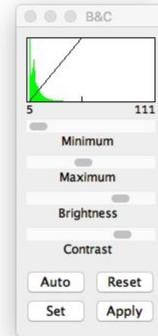


Fiji

- Macros
- Shortcuts
- Utilities
- New
- Compile and Run...
- Install...
- Install Plugin...
- 3D
- 3D Viewer
- 3D script
- Analyze
- BigDataTools
- BigDataViewer
- BigStitcher
- Bio-Formats
- Cidre Plugin
- Cluster
- Color Inspector 3D
- Color Segmentation
- Create Template
- DeconvolutionLab2
- Ellipse Split
- Examples
- Feature Extraction
- FeatureJ
- HDF5
- Image5D
- Integral Image Filters
- Janelia H265 Reader
- LOCI
- LSM Toolbox
- Landmarks
- LoG3D
- MTrackJ
- MoMA
- MorphoLibJ
- Mosaic
- Multiview Reconstruction
- NTA
- NeuronJ
- Non-local Means Denoising
- Optic Flow
- PSF Generator
- Process
- RandomJ
- Registration
- Ridge Detection
- Segmentation
- Shape Filter
- Shape Smoothing
- Skeleton
- Spatial Statistics
- SpotCaliper
- Stacks



\*Rectangle\*, rounded rect or rotated rect (right click to switch)



# Python/Jupyter

jupyter jupyter\_example Last Checkpoint: a few seconds ago (autosaved) ✓

Logout Trusted Python 3

## Thresholding

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import skimage

In [9]: image = skimage.io.imread('Samples/heart.lsm', plugin='tifffile')

In [18]: projection = np.mean(image[0,:,:,:],axis = 0)
thresholded = projection > skimage.filters.threshold_otsu(projection)

In [35]: fig, ax = plt.subplots(1,2, figsize = (15,7))
ax[0].imshow(np.max(image[0,:,:,:],axis = 0),cmap = 'Reds');
ax[0].set_title('Nuclei',fontdict = {'fontsize':20})
ax[1].imshow(thresholded,cmap = 'gray');
ax[1].set_title('Thresholded',fontdict = {'fontsize':20});
```

Nuclei

Thresholded

# Interactive examples

You will find interactive examples illustrating some of the concepts presented in this lecture here :

<https://fl-7-45.zhdk.cloud.switch.ch>

# If the link on the previous page does not work 1/2

You can play with little interactive web-apps by clicking following this link:

[https://github.com/quiwitz/Fundamentals\\_image\\_processing](https://github.com/quiwitz/Fundamentals_image_processing)

And pressing this button:



Or by following this link:

[https://mybinder.org/v2/gh/quiwitz/Fundamentals\\_image\\_processing/master?urlpath=voila](https://mybinder.org/v2/gh/quiwitz/Fundamentals_image_processing/master?urlpath=voila)

For those who are curious, the code for demonstration is available on GitHub (looking at the code IS **NOT A REQUIREMENT OF THE COURSE**):

README.md

launch binder

Fundamentals in digital image\_processing

Course material for fundamentals in digital image processing

# If the link on the previous page does not work 2/2

This should appear:

Thanks to Google Cloud and OVH for sponsoring our computers !

binder

Starting repository:  
guiwitz/Fundamentals\_image\_processing/master

We use the [repo2docker](#) tool to automatically build the environment in which to run your code.

Show

Build logs

jupyter nbviewer

JUPYTER FAQ

Here's a non-interactive preview on nbviewer while we start a server for you. Your binder will open automatically when it is ready.

Just wait until this appears:

Select items to open with voila.

- [04-Fundamentals\\_template\\_matching\\_interactive.ipynb](#)
- [02-Fundamentals\\_contrast\\_interactive.ipynb](#)
- [05-Fundamentals\\_filtering\\_interactive.ipynb](#)
- [06-Fundamentals\\_thresholding\\_ineractive.ipynb](#)
- [07-Fundamentals\\_binary\\_interactive.ipynb](#)
- [03-Fundamentals\\_browsing\\_ineractive.ipynb](#)
- [01-Fundamentals\\_digitization\\_interactive.ipynb](#)

To run the web-app, just click on the appropriate link. You can come back to this view using the ← button of the browser.

If things stop working, just click again on the binder button:



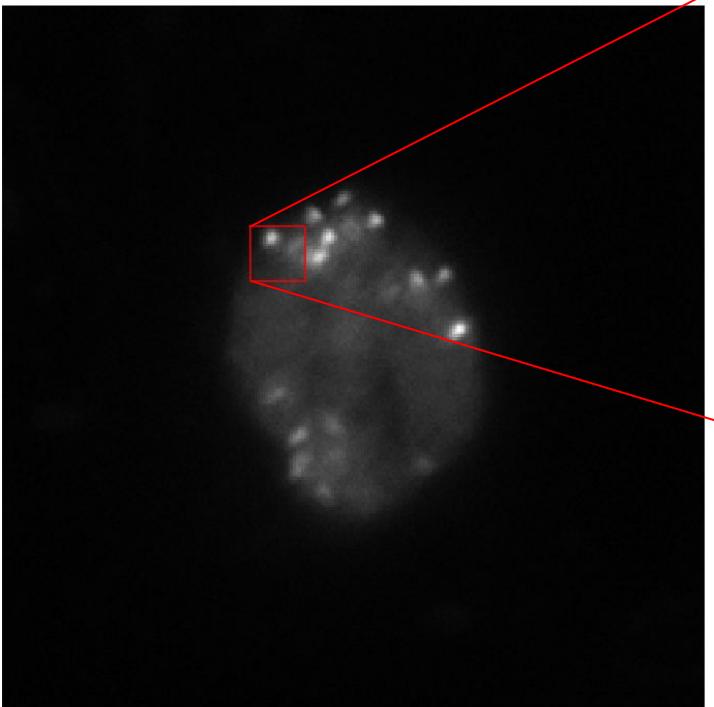
# Table of contents

- I. What is an image?
- II. How is an image stored?
- III. Representing an image: non-destructive operations
- IV. Processing an image: filtering, resampling etc (destructive operations)
- V. Segmentation
- VI. Advanced image processing

I.

What is an image ?

# What is an image?



0	131	136	134	132	132	136	133	133	141	141	143	147	149	152	151	154	160	157	170	
1	136	135	135	135	141	139	134	137	144	148	151	154	158	152	153	156	159	163	181	184
2	141	137	141	139	144	142	147	144	151	156	160	158	157	159	154	158	166	171	171	
3	141	145	143	151	157	158	162	159	157	160	160	159	157	157	158	159	165	167	170	
4	144	144	165	169	180	184	193	186	173	170	166	159	164	161	169	164	161	171	172	160
5	149	164	190	206	232	241	244	230	208	172	168	152	160	155	164	170	164	173	172	177
6	158	179	219	268	315	317	292	267	222	194	180	175	170	170	173	169	171	173	180	175
7	164	202	249	326	379	386	351	271	218	193	192	180	177	177	177	173	172	179	181	176
8	168	201	265	364	394	403	362	279	237	199	196	181	175	180	185	172	180	178	183	183
9	159	190	244	318	376	355	299	254	222	308	308	169	165	191	182	185	176	162	179	185
10	155	184	213	261	295	288	256	240	223	226	196	189	185	190	186	176	179	183	180	185
11	160	170	187	220	235	239	230	225	215	211	193	190	180	189	182	188	187	189	194	194
12	163	173	181	200	207	212	217	219	221	222	217	215	209	206	199	207	197	192	192	205
13	161	177	177	190	199	206	212	225	229	229	227	216	213	202	194	199	197	196	193	195
14	173	173	187	194	203	207	228	242	249	247	245	239	236	206	211	209	204	201	201	
15	173	178	185	195	198	223	245	255	274	268	254	237	224	225	206	217	204	206	210	204
16	178	184	194	210	216	230	258	264	274	260	247	247	242	232	218	218	202	209	206	203
17	194	192	191	205	221	250	264	284	267	259	255	242	237	232	224	239	233	208	206	203
18	188	200	205	222	230	259	273	271	261	240	239	243	240	246	231	227	213	215	209	203
19	190	202	202	213	241	260	266	263	250	251	239	254	287	267	246	250	222	229	215	203

An image is simply a collection of pixels arranged as a **2D matrix**

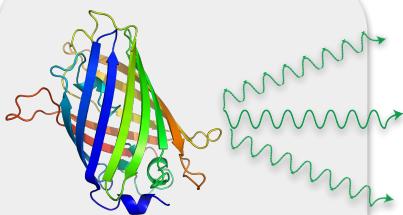
In raw image data, the pixel value has a **physical meaning** (e.g. number of photons emitted by a fluorescent protein).

# Keeping raw data raw

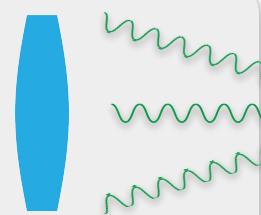
0	131	136	134	132	132	136	133	133	141	141	143	147	149	152	152	151	154	160	157	170
1	136	135	135	135	141	139	134	137	144	148	151	154	158	152	153	156	159	163	162	165
2	141	137	141	139	144	142	147	144	151	156	156	158	157	159	154	158	166	171	173	171
3	141	145	143	151	157	148	162	159	157	160	160	158	157	157	156	159	165	167	174	170
4	144	144	165	169	180	168	193	186	173	170	166	159	184	201	189	184	185	172	172	180
5	149	164	190	206	232	243	244	230	208	172	168	162	168	151	164	155	164	173	172	177
6	158	179	219	268	315	317	292	267	222	194	180	173	170	160	170	169	171	173	180	175
7	164	202	249	326	378	396	351	271	218	193	192	180	177	177	177	173	172	179	181	176
8	168	201	265	364	394	401	362	275	227	193	180	161	175	180	185	172	180	178	183	183
9	159	190	244	318	376	365	299	254	220	308	186	186	181	181	181	182	185	176	182	185
10	155	184	213	261	295	288	256	280	223	226	196	189	186	180	186	176	179	183	180	185
11	160	170	187	220	235	239	230	225	215	211	193	200	202	199	182	188	187	189	194	194
12	163	173	181	200	207	212	217	219	221	222	217	215	209	206	199	197	197	192	192	205
13	161	177	177	190	199	206	212	225	229	229	227	216	213	202	194	194	197	196	193	195
14	173	173	187	194	203	207	228	242	249	247	245	239	216	208	211	202	206	201	202	201
15	173	178	185	195	198	223	245	255	274	268	254	237	224	225	206	217	204	206	210	204
16	178	184	194	210	216	230	258	264	274	280	247	247	242	232	218	218	202	208	208	203
17	194	192	191	205	221	250	264	284	267	259	255	242	237	232	224	215	215	209	205	205
18	188	200	205	222	230	259	273	271	261	240	239	245	240	246	234	237	211	215	209	200
19	190	202	202	213	241	260	266	263	250	251	239	254	287	287	248	230	223	220	212	201

Any time you change the pixel values e.g. by drawing or by saving in another format e.g. jpg the physical information is **LOST**

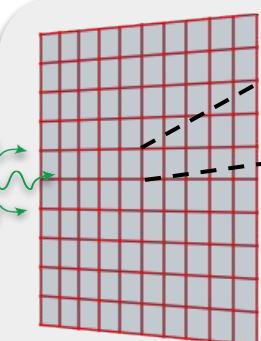
# From physics to numbers



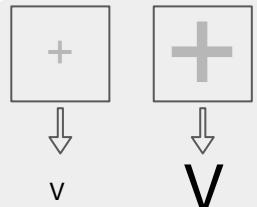
1. Fluorescent protein emits light



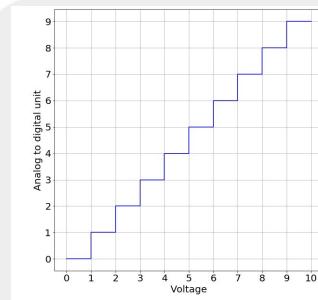
2. Emitted light goes through microscope optics



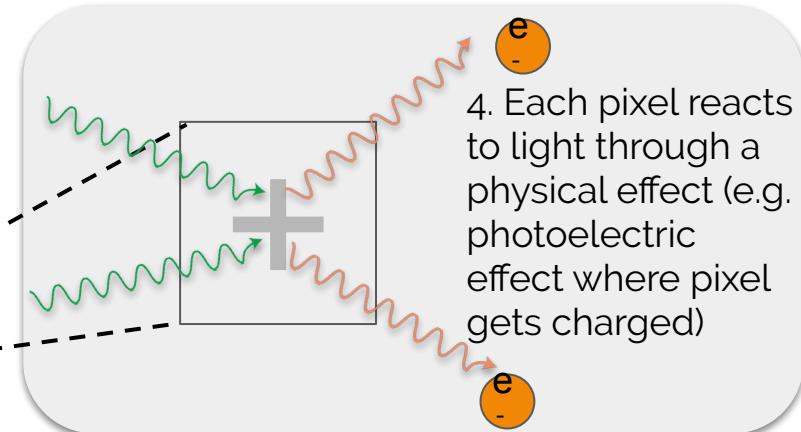
3. Light lands on camera detector composed of pixels



5. Pixel charge translated into voltage



6. Voltage of each pixel is digitized, i.e. assigned a number depending on the voltage range to which it belongs.



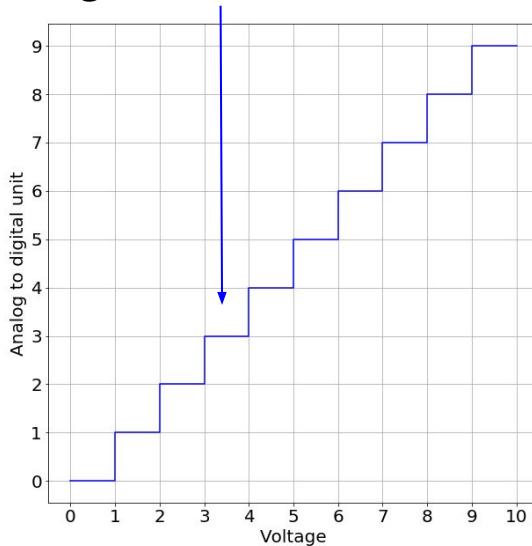
4. Each pixel reacts to light through a physical effect (e.g. photoelectric effect where pixel gets charged)

1	0	1	3	2	0	0	0	0	0
1	2	2	2	2	1	2	0	0	0
0	2	3	3	5	0	1	0	1	1
2	4	4	4	4	4	3	2	0	0
1	3	5	9	6	4	1	2	0	0
3	6	7	7	7	6	4	2	0	0
2	4	6	6	7	4	1	2	0	0
2	4	6	4	5	3	1	0	0	0
0	2	3	2	4	3	2	0	0	1
0	2	2	2	1	1	1	0	0	0

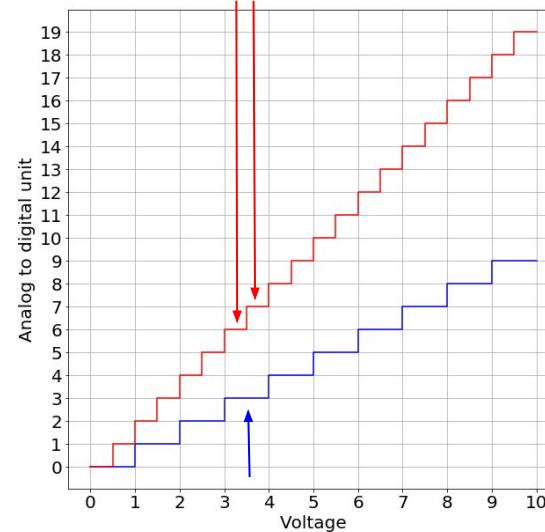
7. **For plotting purposes**, a color shade is assigned to each digitization level

# Digitization

Voltage from 3V to 4V is assigned value 3



Voltage from 3V to 4V is split into two values: more “detail” visible



**By convention,**  
images are most of  
the time encoded  
into 8, 16 or 32 bits, i.e

8 bits : 256 values  
16 bits : 65536 values  
32 bits:  $\sim 4 \cdot 10^9$  values

You can play with the *digitization* web-app

II.

How is an image stored ?

# Metadata: reconstructing the physical meaning

## Metadata:

- What microscope, what camera etc
- Exposure time
- Position of acquisition (e.g. 96 well plate)

31	30	31	32	35	37	36	36	31	28
31	30	31	33	32	33	32	30	27	25
29	30	30	32	32	31	31	28	25	23
30	29	30	30	30	29	30	27	25	25
28	27	29	29	29	29	28	26	25	25
27	27	27	27	29	27	27	25	25	25
27	26	26	28	27	27	28	27	25	25
27	26	28	26	27	29	29	27	25	24
28	28	27	26	27	27	29	27	25	19
27	27	26	26	28	28	29	26	21	

Information about the image is stored in a “**header**”

Pixel information is stored in binary form in the file body

# Metadata: reconstructing the physical meaning

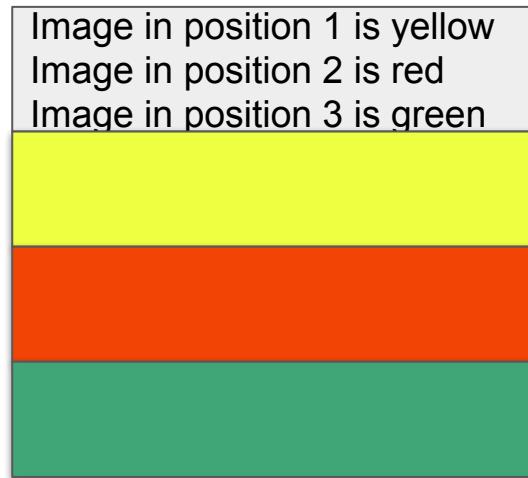
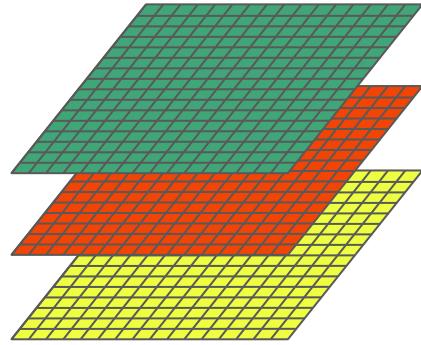


The screenshot shows a code editor displaying an XML file titled "OME Metadata - 190322\_MacrophagesSlaug\_DAPI\_Phall633\_Iba1555\_GFPLm\_2h\_60x\_0005.oir". The XML structure is as follows:

- Instrument**:
  - Laser**:
    - ID="LightSource:0:0" Model="LD405\_65792" Wavelength="405.0" WavelengthUnit="nm"/>
    - ID="LightSource:0:1" Model="LD445\_65793"/>
    - ID="LightSource:0:2" Model="LD488\_65794" Wavelength="488.0" WavelengthUnit="nm"/>
    - ID="LightSource:0:3" Model="LD514\_65795"/>
    - ID="LightSource:0:4" Model="LD561\_65796" Wavelength="561.0" WavelengthUnit="nm"/>
    - ID="LightSource:0:5" Model="LD594\_65797"/>
    - ID="LightSource:0:6" Model="LD640\_65798" Wavelength="640.0" WavelengthUnit="nm"/>
  - Detector**:
    - Gain**:
      - ID="Detector:0:0" Offset="0.0" Voltage="419.0" VoltageUnit="V"/>
      - ID="Detector:0:1" Offset="0.0" Voltage="350.0" VoltageUnit="V"/>
      - ID="Detector:0:2" Offset="2.0" Voltage="350.0" VoltageUnit="V"/>
      - ID="Detector:0:3" Offset="1.0" Voltage="397.0" VoltageUnit="V"/>
      - ID="Detector:0:4" Offset="0.0" Voltage="408.0" VoltageUnit="V"/>
  - Objective**:
    - ID="Objective:0:0" Immersion="Oil" LensNA="1.3" Model="UPLSAPO 60" NominalMagnification="60.0" WorkingDistance="0.3" WorkingDistanceUnit="mm"/>
  - Image**:
    - ID="Image:0" Name="190322\_MacrophagesSlaug\_DAPI\_Phall633\_Iba1555\_GFPLm\_2h\_60x\_0005.oir">
      - AcquisitionDate**:>2019-05-29T12:41:34.450
      - InstrumentRef**: ID="Instrument:0"/>
      - ObjectiveSettings**: ID="Objective:0:0" RefractiveIndex="1.406"/>
  - Pixels**:
    - BigEndian="false" DimensionOrder="XYCZT" ID="Pixels:0" Interleaved="false" PhysicalSizeX="0.096775563118596" PhysicalSizeXUnit="μm" PhysicalSizeY="0.096775563118596" PhysicalSizeYUnit="μm">
      - Channel**:
        - Color="16777215" EmissionWavelength="481.0" EmissionWavelengthUnit="nm" ExcitationWavelength="441.0" ExcitationWavelengthUnit="nm" ID="Channel:0:0" Name="CH1" PinholeSize="238.0">
          - LightSourceSettings**: ID="LightSource:0:0"/>
          - DetectorSettings**: ID="Detector:0:1"/>
          - LightPath**:
        - Color="16711935" EmissionWavelength="540.0" EmissionWavelengthUnit="nm" ExcitationWavelength="500.0" ExcitationWavelengthUnit="nm" ID="Channel:0:1" Name="CH2" PinholeSize="238.0">
          - LightSourceSettings**: ID="LightSource:0:2"/>
          - DetectorSettings**: ID="Detector:0:2"/>
          - LightPath**:
        - Color="-8453889" EmissionWavelength="614.0" EmissionWavelengthUnit="nm" ExcitationWavelength="564.0" ExcitationWavelengthUnit="nm" ID="Channel:0:2" Name="CH3" PinholeSize="238.0" I
          - LightSourceSettings**: ID="LightSource:0:4"/>
          - DetectorSettings**: ID="Detector:0:3"/>
          - LightPath**:
        - Color="167776961" EmissionWavelength="744.0" EmissionWavelengthUnit="nm" ExcitationWavelength="644.0" ExcitationWavelengthUnit="nm" ID="Channel:0:3" Name="CH4" PinholeSize="238.0" I
          - LightSourceSettings**: ID="LightSource:0:6"/>
          - DetectorSettings**: ID="Detector:0:4"/>
          - LightPath**:
      - MetadataOnly**:>

# More than one image in a file

Often images of the same “object” are acquired in **multiple dimensions**: different **modes** (wave length), **positions** (x, y and depth z) or over **time**. Such ensembles of pictures can be saved within a single file. The header can be used to specify “what is where”.



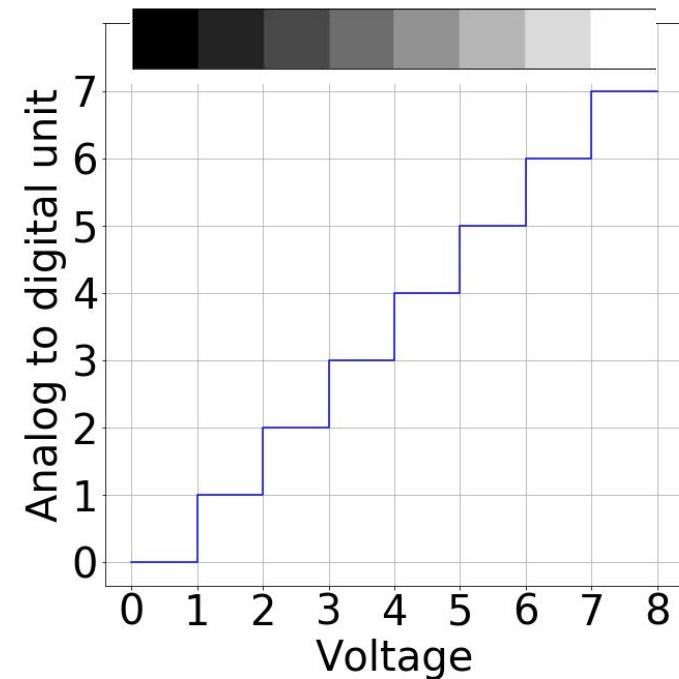
my\_image.tif

You can play with the *browsing* web-app

# III

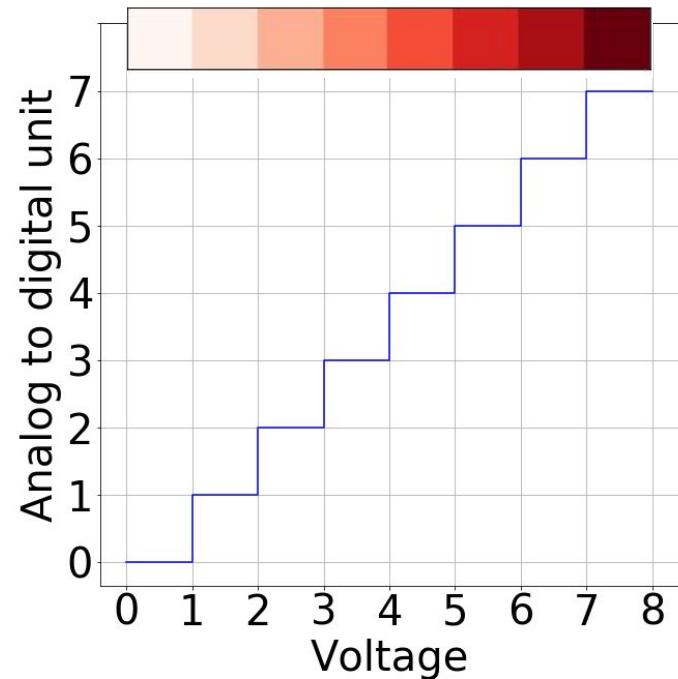
## Representing an image: Non-destructive operations

# Color maps



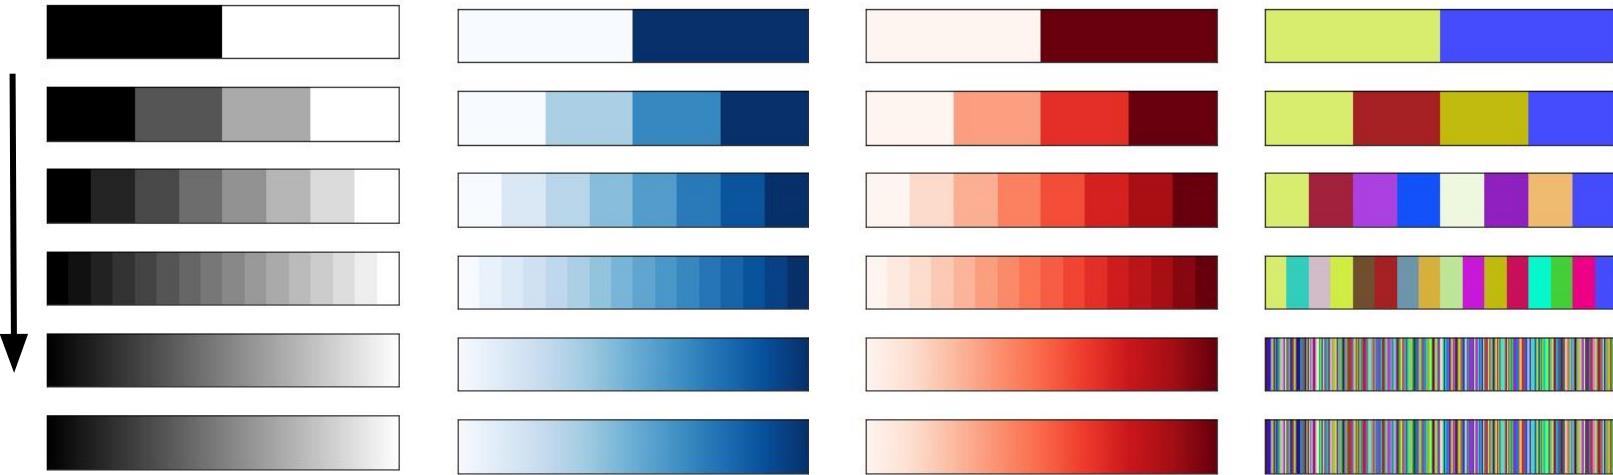
The color we see on a screen is only a mapping of ADU to a given color map.

The choice of color map does not have a physical meaning.

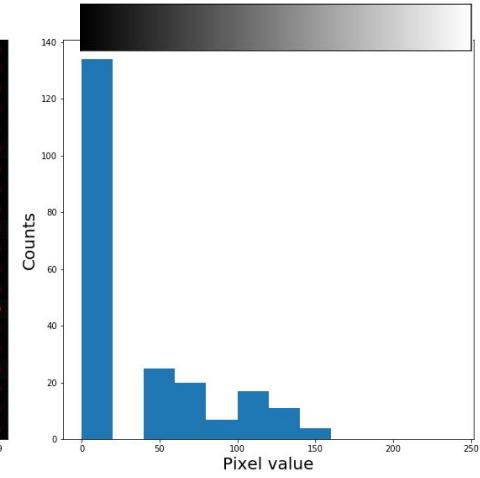
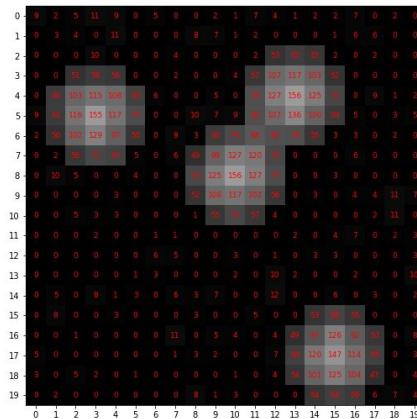


# Color maps

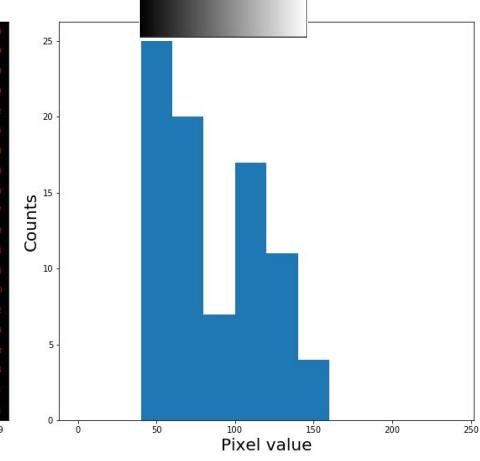
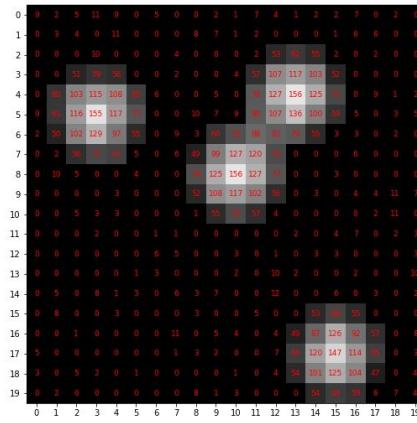
Digitization detail



# Low contrast



# High contrast



# Contrast

If a range of values is not actually used in the image, the associated color shades are “lost”. One can then use a single color shade for all those values, and use the other shades to see better see intensity changes in the image, i.e. increase the contrast.

You can play with the **contrast** web-app

# Volumetric representation

3D data can also be represented in 3D. Different strategies exist to show the depth e.g. projection of the brightest pixel along the viewing axis, surface representation etc.

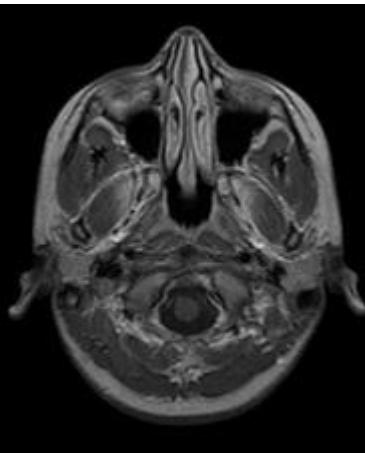
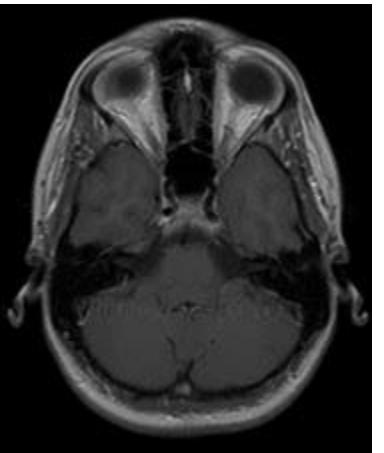
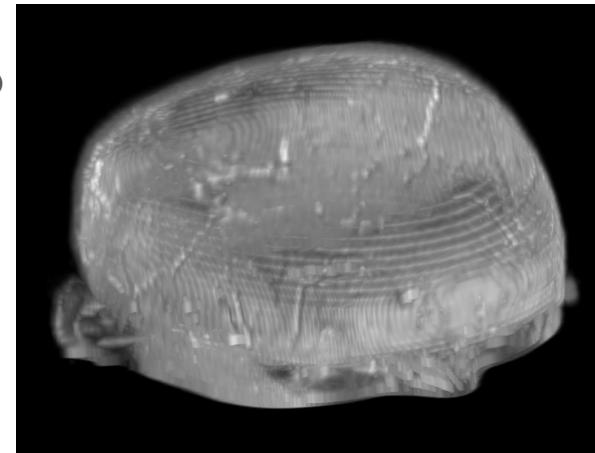
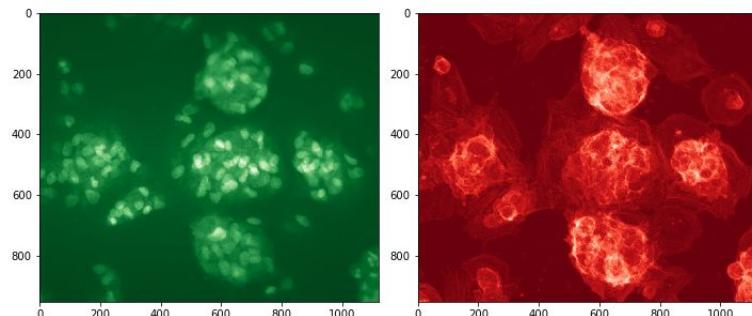


Image stack turned into  
transparent volume

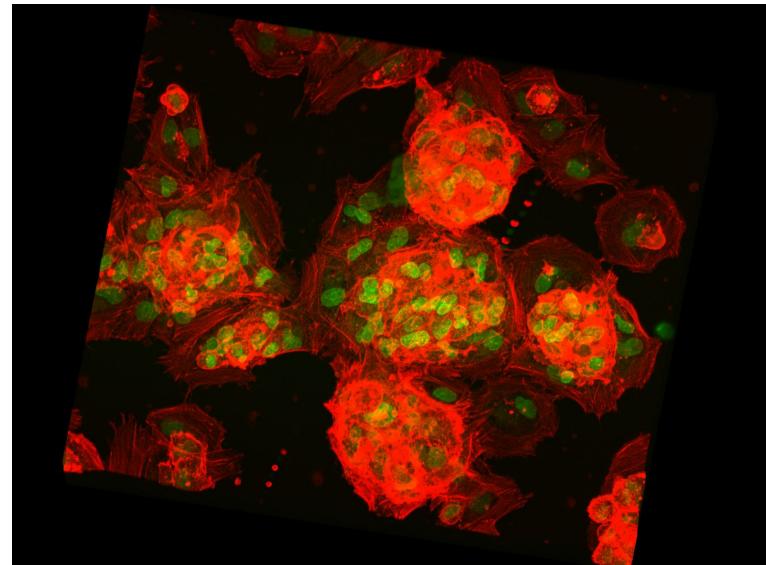


# Volumetric representation

3D data can also be represented in 3D. Different strategies exist to show the depth e.g. projection of the brightest pixel along the viewing axis, surface representation etc.



Blending two  
channels in 3D



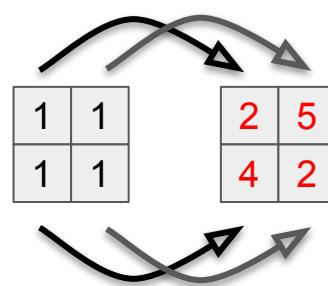
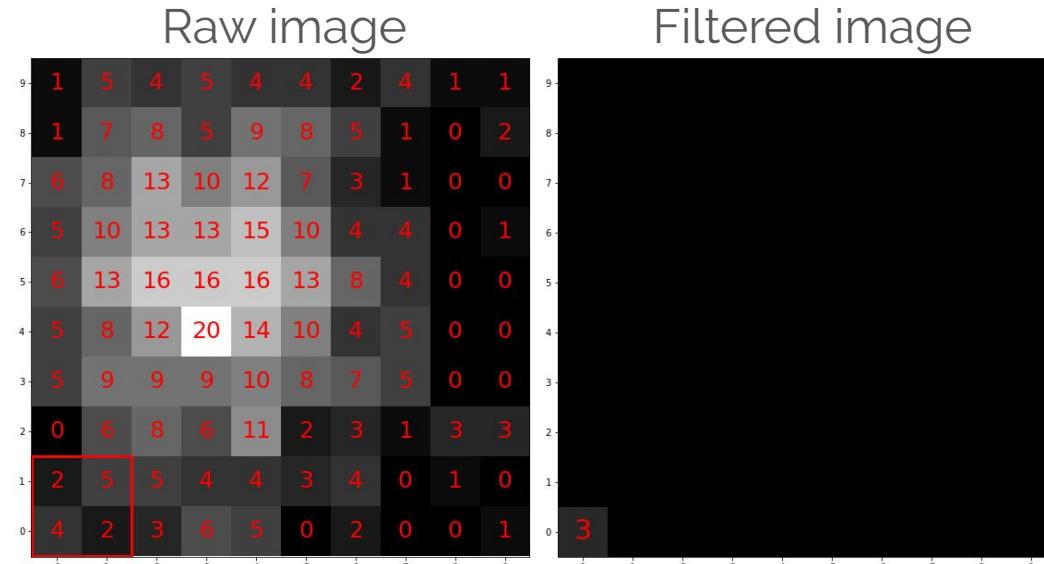
# IV.

## Resampling, filtering etc. Destructive operations

# Linear Filtering

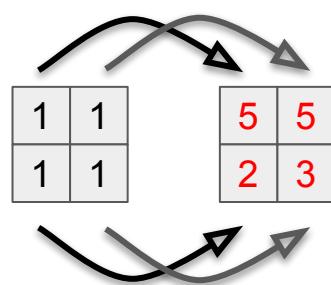
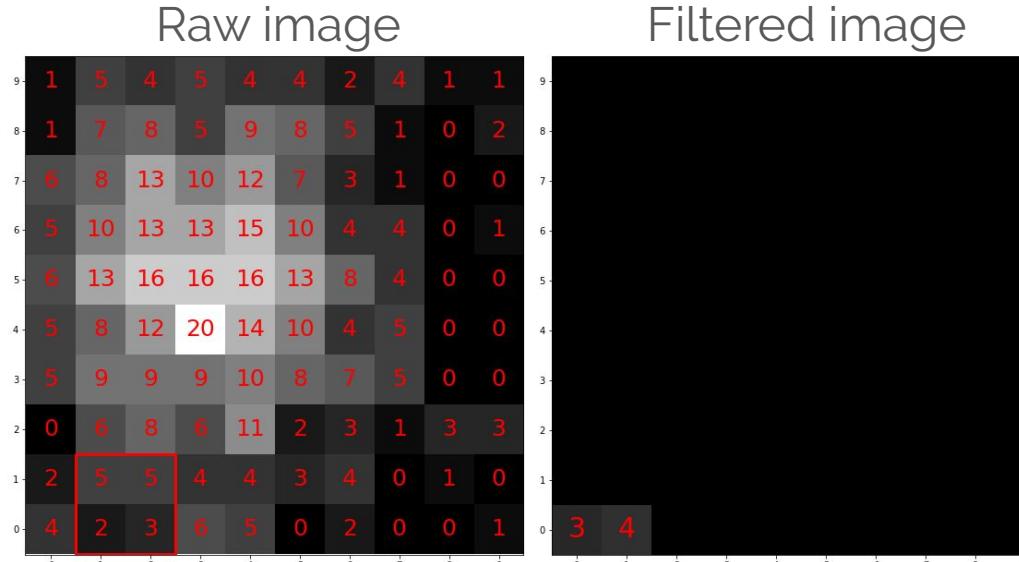
A small filter window travels across the image. At each position, each pixel of the travel window is multiplied with the corresponding pixel of the image window. The average (or sum) of these multiplications is the pixel value of the filtered image.

This operation is also called a **convolution** (can be done in Fourier space for large images).



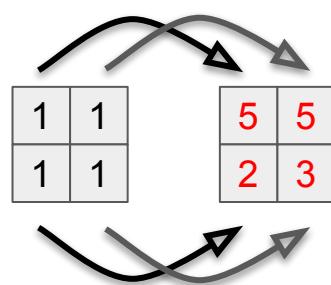
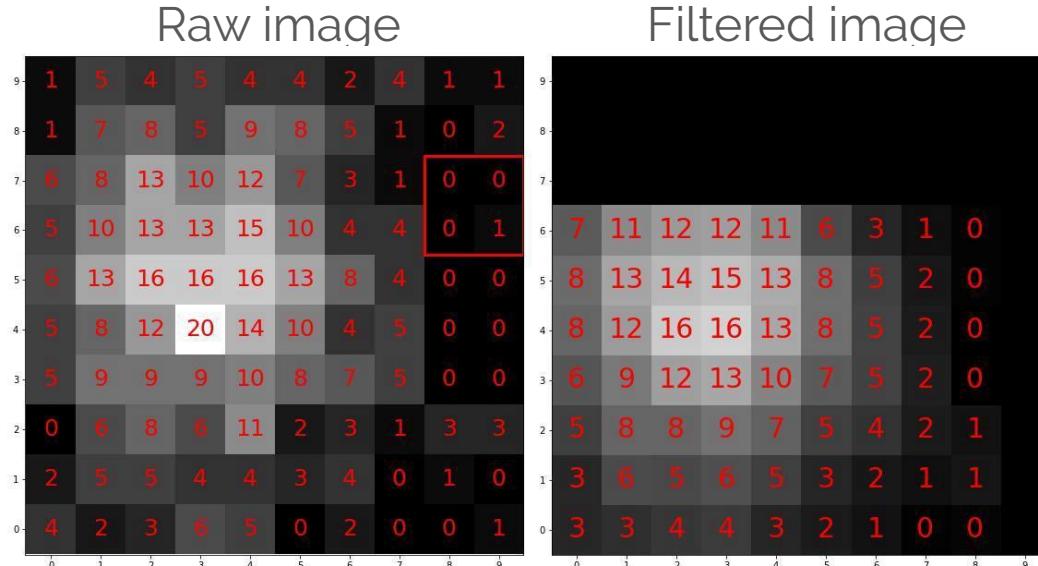
# Linear Filtering

A small filter window travels across the image. At each position, each pixel of the travel window is multiplied with the corresponding pixel of the image window. The average (or sum) of these multiplications is the pixel value of the filtered image.



# Linear Filtering

A small filter window travels across the image. At each position, each pixel of the travel window is multiplied with the corresponding pixel of the image window. The average (or sum) of these multiplications is the pixel value of the filtered image.

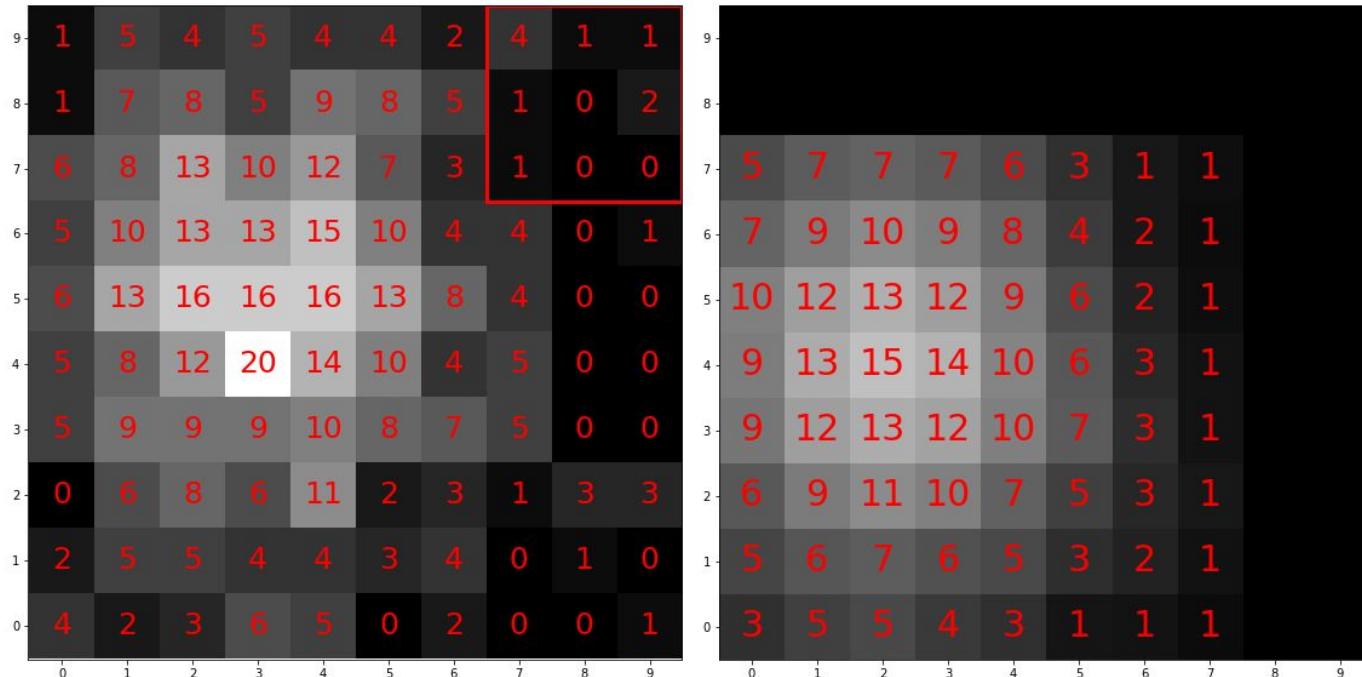


$$(5*1+5*1+2*1+3*1)/4 = 15/4 = 3.75$$

$$\text{round}(3.75) = 4$$

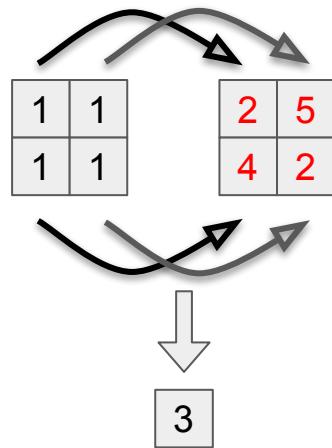
# Linear Filtering

The filter can have **any size**. Here it covers a 3x3 region



# Linear filtering: the filter can be any image

The filter is generally not filled with 1's



Filtering gives maximal response when filter and image region are identical. Therefore convolution can be used to detect specific features.

0	0	1	1	1	2	1	1	1	0	0
0	1	2	3	4	4	4	3	2	1	0
1	2	3	5	7	8	7	5	3	2	1
1	3	5	9	12	13	12	9	5	3	1
1	4	7	12	16	18	16	12	7	4	1
2	4	8	13	18	20	18	13	8	4	2
1	4	7	12	16	18	16	12	7	4	1
1	3	5	9	12	13	12	9	5	3	1
1	2	3	5	7	8	7	5	3	2	1
0	1	2	3	4	4	4	3	2	1	0
0	0	1	1	1	2	1	1	1	0	0

Gaussian filter  
(spot detection)

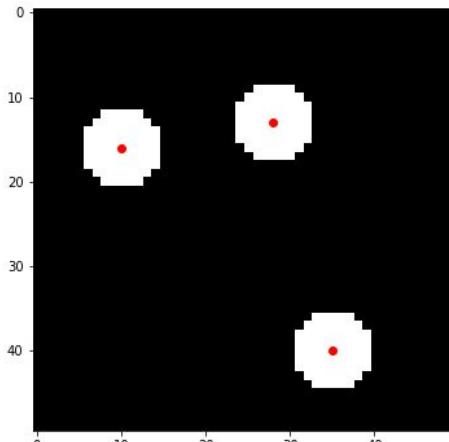
-1	-1	-1
0	0	0

1      1      1

Prewitt filter  
(ridge detection)

# Linear filtering for detection

If the shape of the object to be detected is known, one can use a special type of filtering (template matching) to detect them: high response (white) in the filtered image corresponds to objects.



Clean image

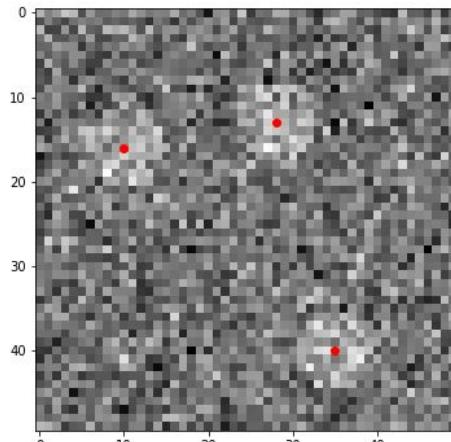
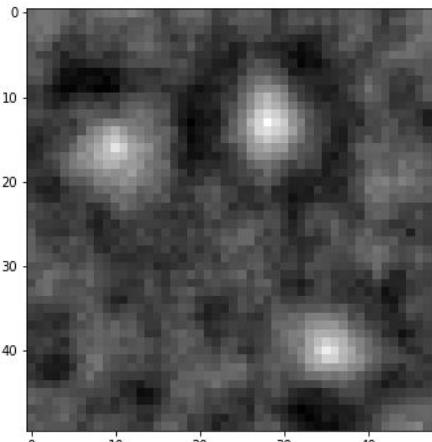
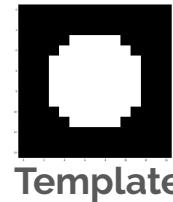


Image corrupted by noise



Filtered image

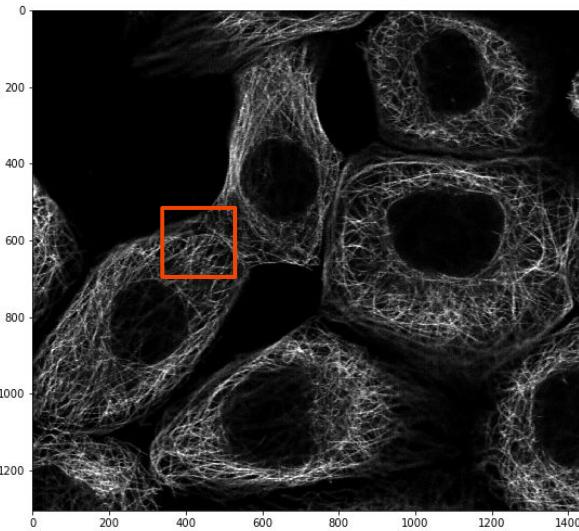


Template

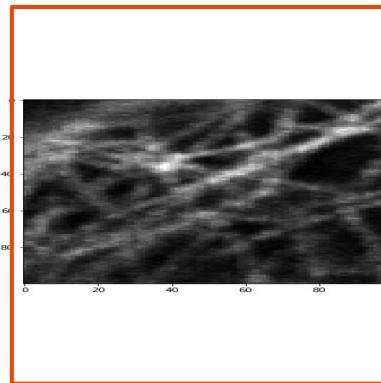
You can play with the [\*\*\*template-matching\*\*\*](#) web-app

# Linear filtering for detection

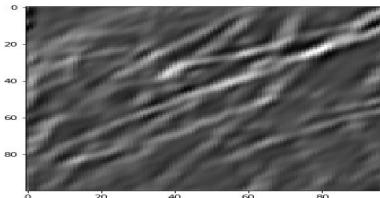
Other filters can for example detect the orientation of elongated structures (e.g. Gabor filters).



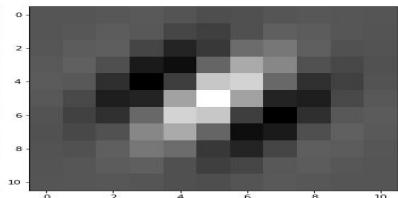
Actin labeling



Zoom



Filtered image

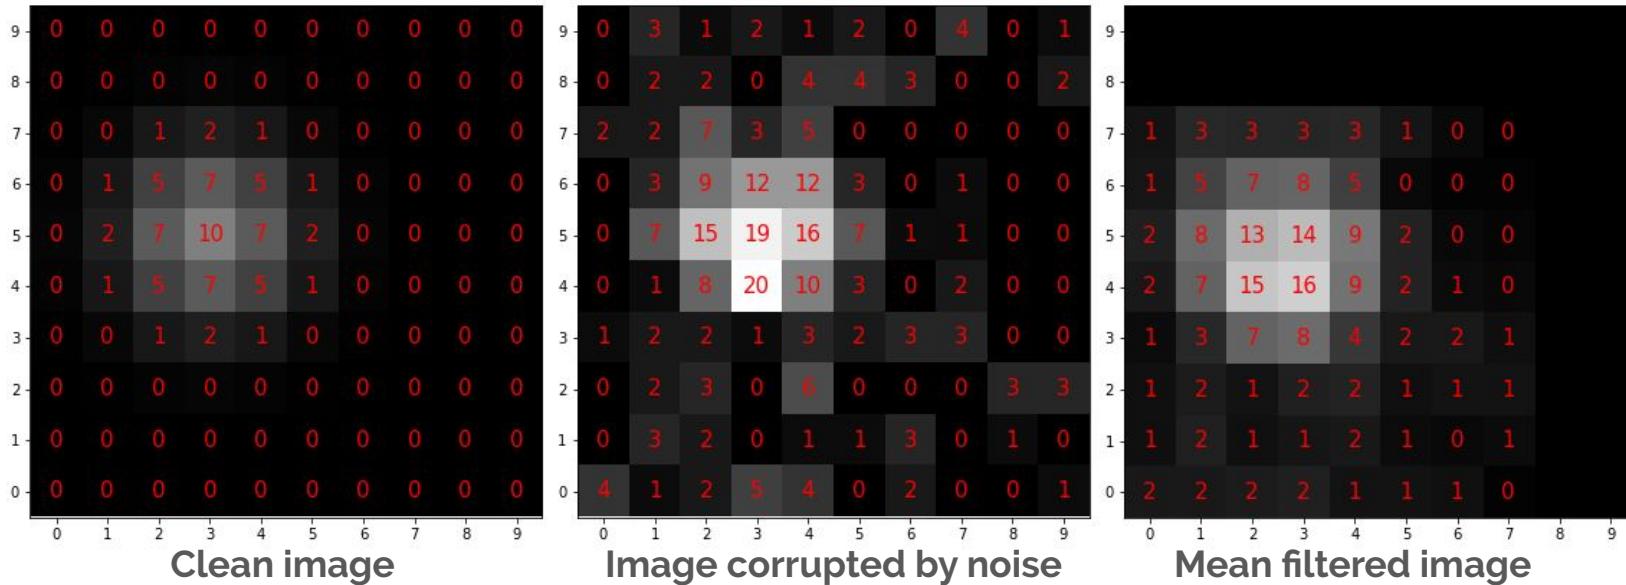


Filter

[You can play with the \*orientation\* web-app](#)

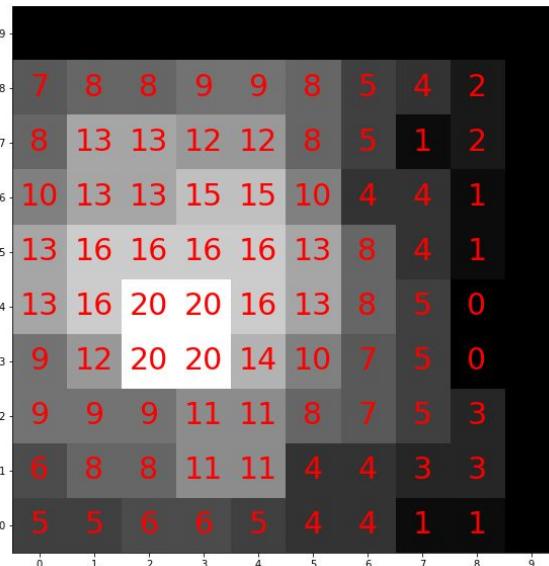
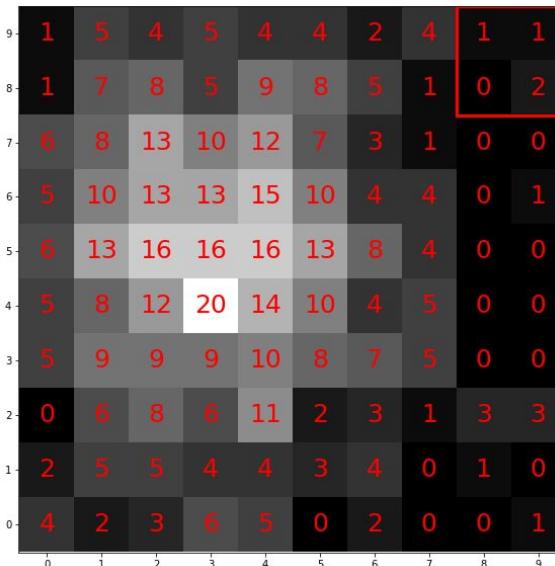
# Linear filtering for noise removal

As seen before, if the filter is composed of 1's, it simply calculates a local mean for each pixel (mean filter). This can be used to remove noise in an image: pixel with deviating values are averaged out in this process



# Non-linear filters

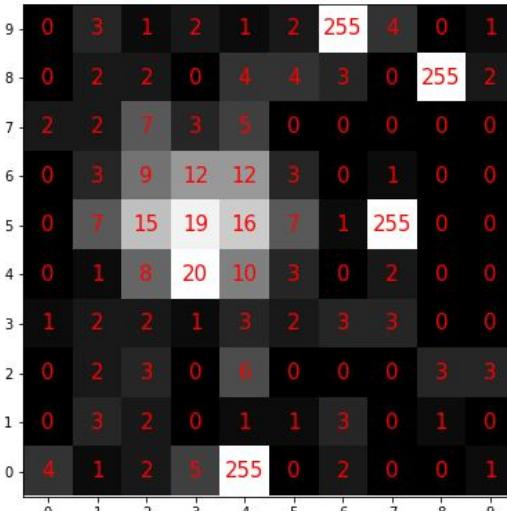
Instead of simply multiplying and summing pixels, non-linear filters can apply an operation on the sub-regions they travel through. For example, one can calculate the local maximum at each pixel position.



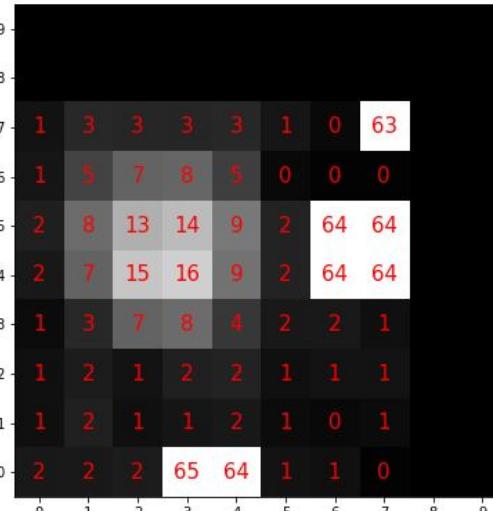
$$\max \left( \begin{array}{|c|c|} \hline 1 & 1 \\ \hline 0 & 2 \\ \hline \end{array} \right) = 2$$

# Different filters for different usages

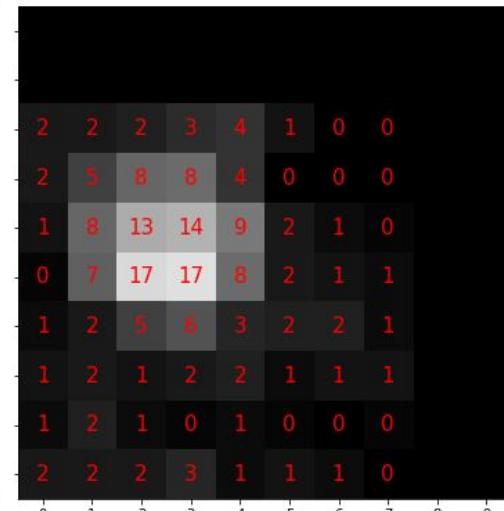
In presence of a few pixels with extreme values, the mean filter is a bad choice.  
The median filter in contrast is very efficient at this task.



Salt & pepper corrupted image



Mean filtered

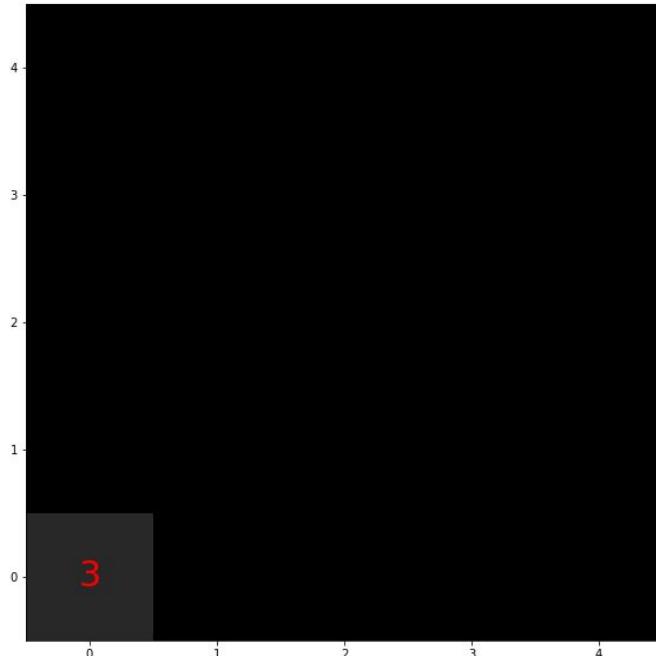
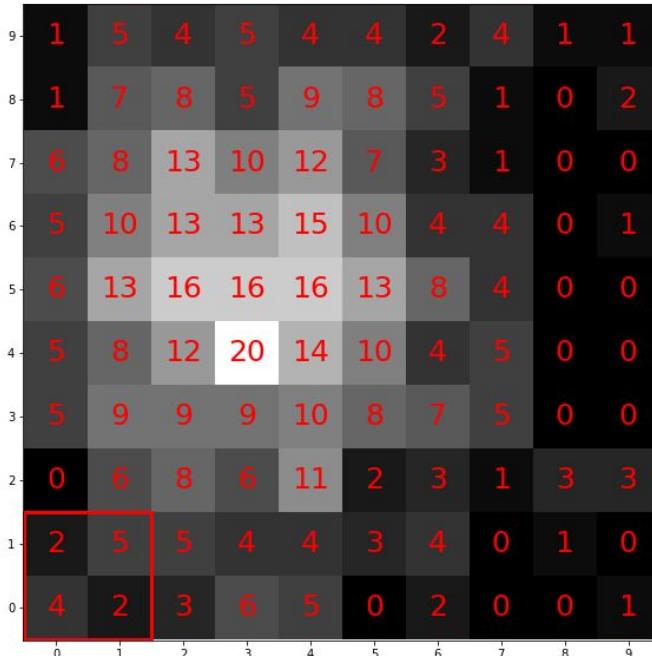


Median filtered

You can play with the *filtering* web-app

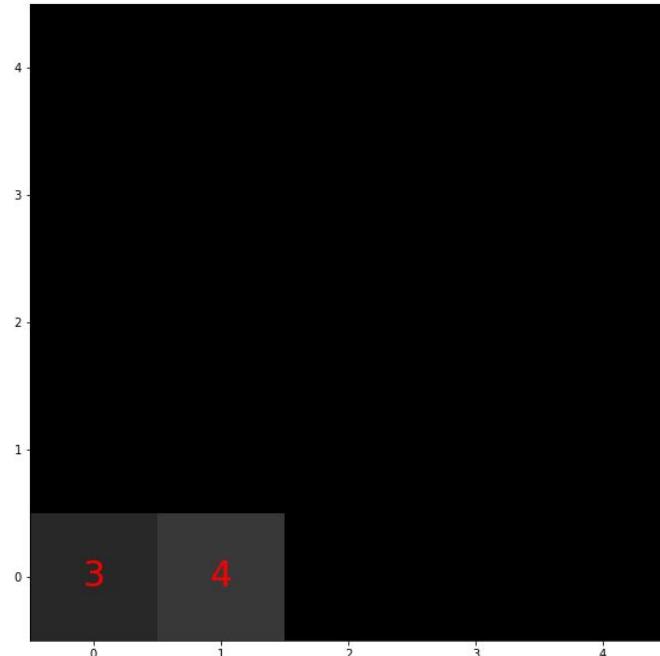
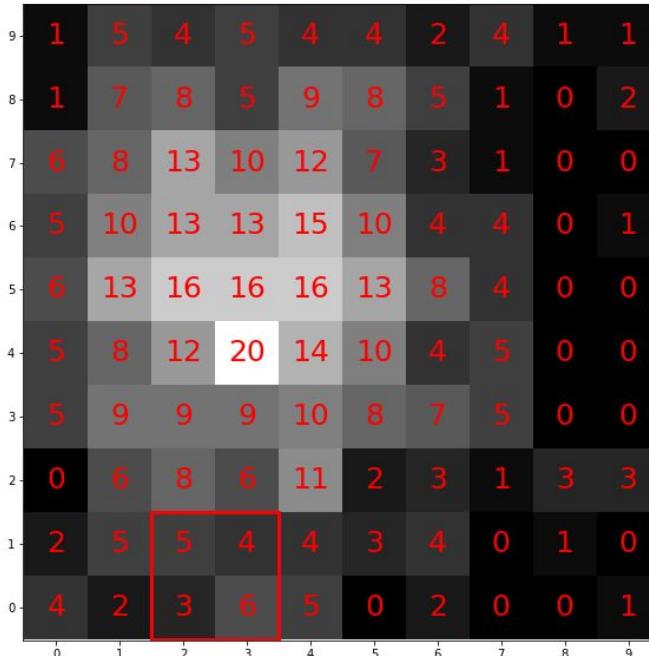
# Resampling (downsampling)

When applying a mean filter on non-overlapping sub-regions, one can down-sample an image. This is often used to decrease computational time.



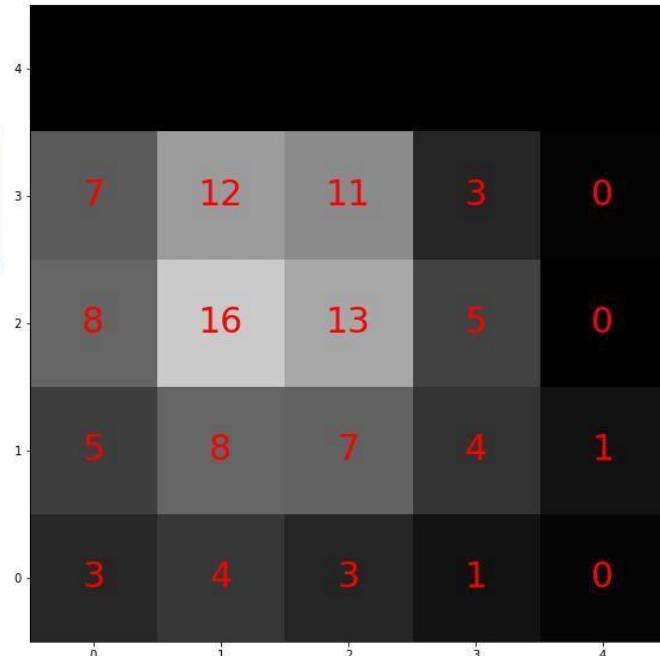
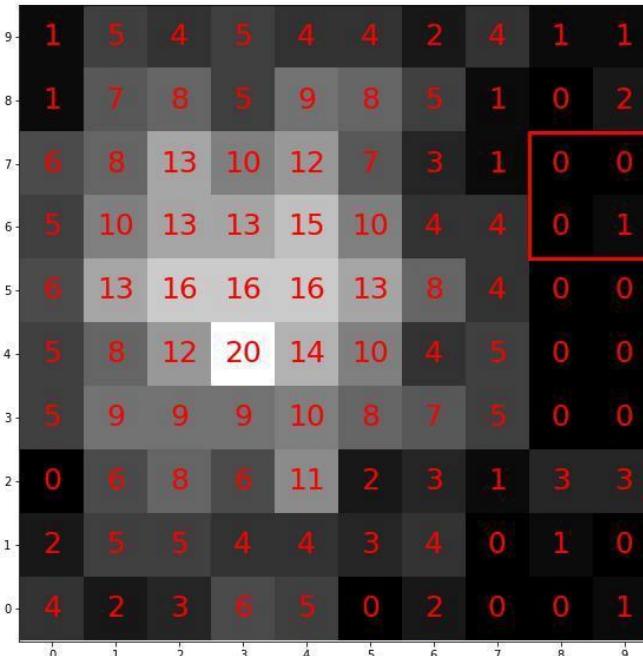
# Resampling (downsampling)

When applying a mean filter on non-overlapping sub-regions, one can down-sample an image. This is often used to decrease computational time.



# Resampling (downsampling)

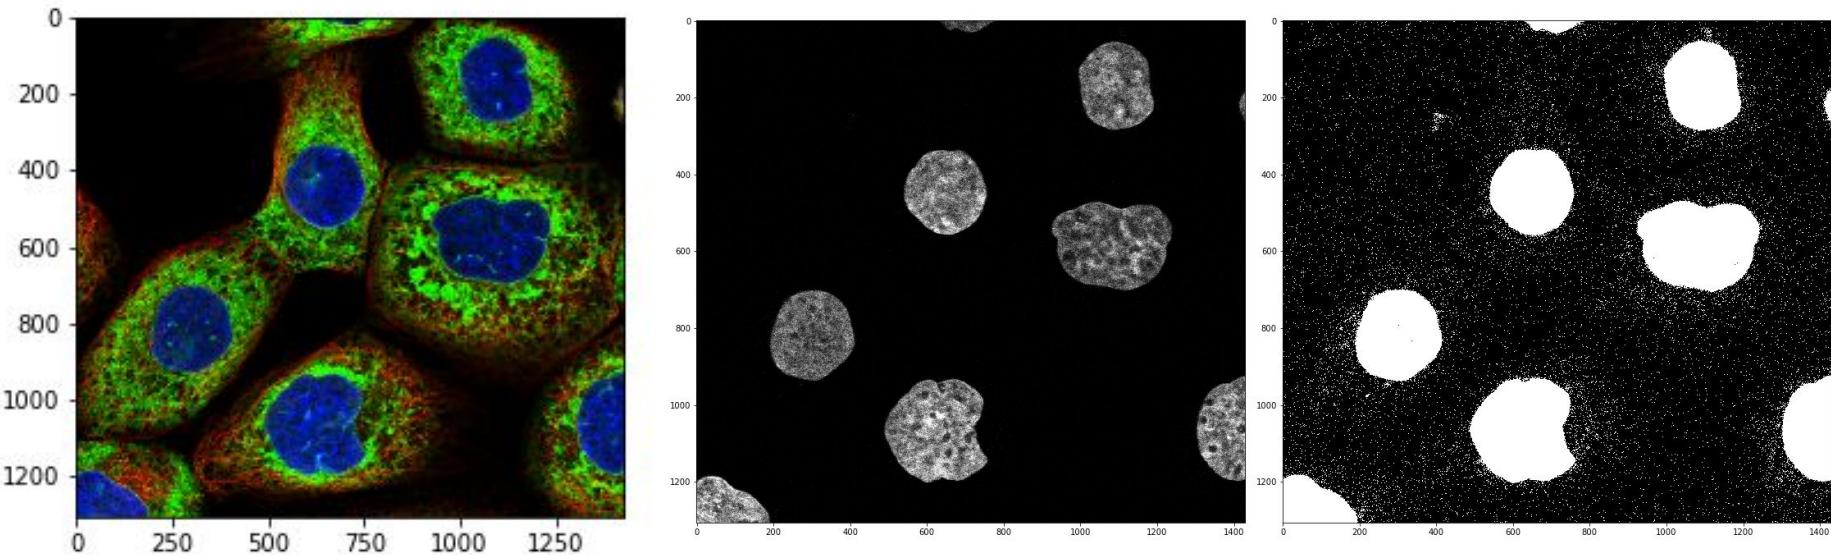
When applying a mean filter on non-overlapping sub-regions, one can down-sample an image. This is often used to decrease computational time.



V.  
Segmentation

# Thresholding (binarization)

Creation of a binary image where any pixel with a value within a certain range becomes 1 and the others 0. A typical example is the segmentation of nuclei in fluorescence microscopy.



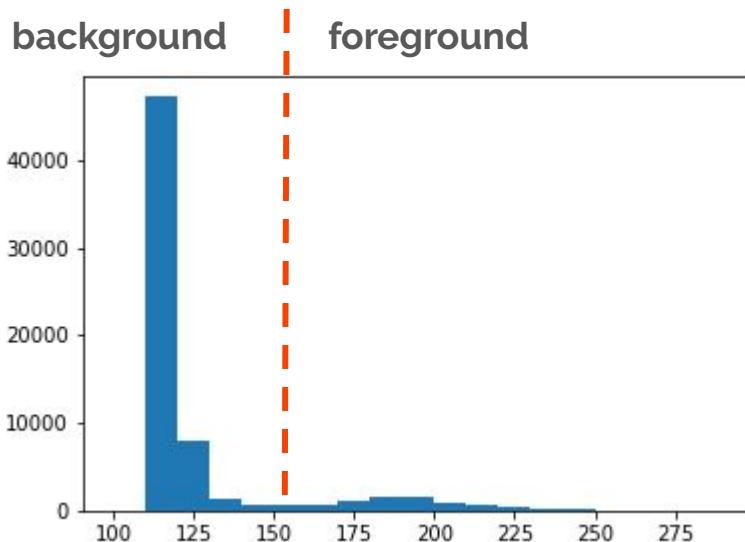
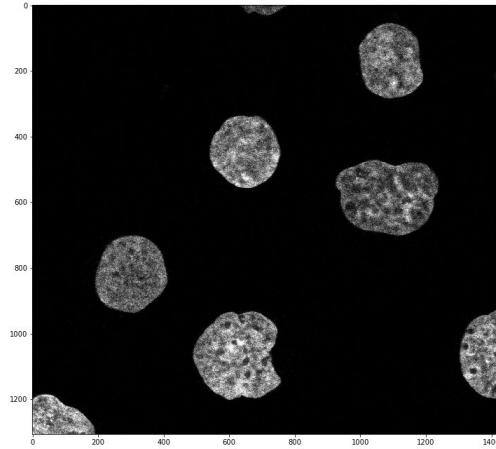
Epiderm human cells, image from the protein atlas (blue nuclei, red tubulin, green atlas antibody)

<http://cellimagelibrary.org/images/41678> <https://www.proteinatlas.org/> doi:10.7295/W9CIL41678

# Thresholding: automatization

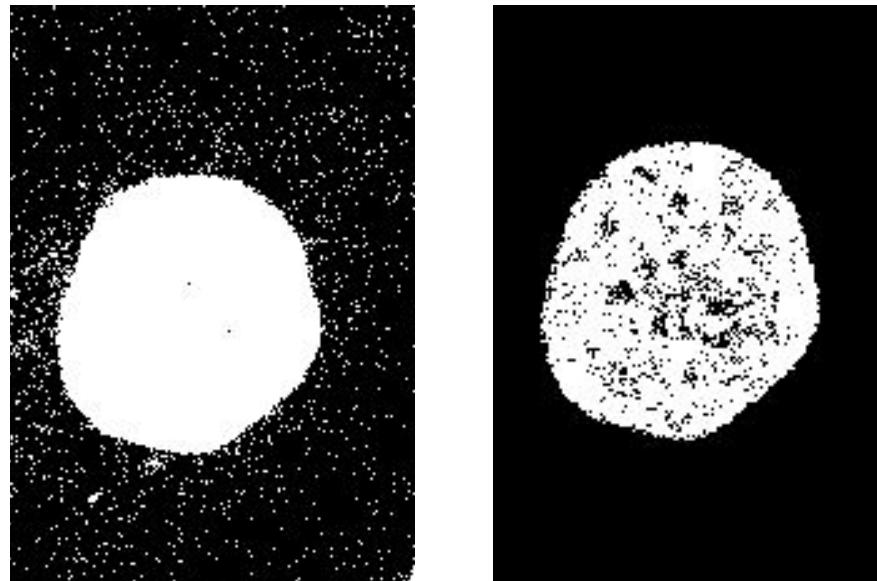
Typically, an image contains foreground and background pixels. This can be seen in the histogram of pixel values composed of two distributions.

Automated thresholding methods attempt to find the “best” separation of the two distributions (e.g. minimize the variance of both distributions)



# Binary operations

The binary image obtained from thresholding often needs to be cleaned-up. In our example, we might have noisy pixels ON outside the nucleus, or OFF inside



You can play with the ***thresholding*** web-app

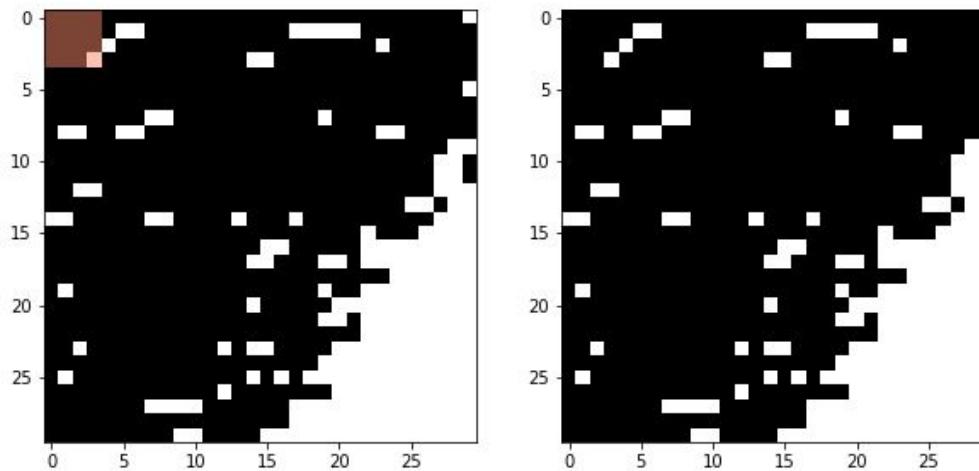
# Binary operations

To handle these cases, we use binary operators. These are filters that turn pixels ON/OFF depending on the local neighborhood. The simplest ones are the eroding and dilating filters.

## Erosion/dilation:

Whenever the subregion contains at least an OFF/ON pixel, turn that subregion location (here, upper left corner) OFF/ON

Erosion



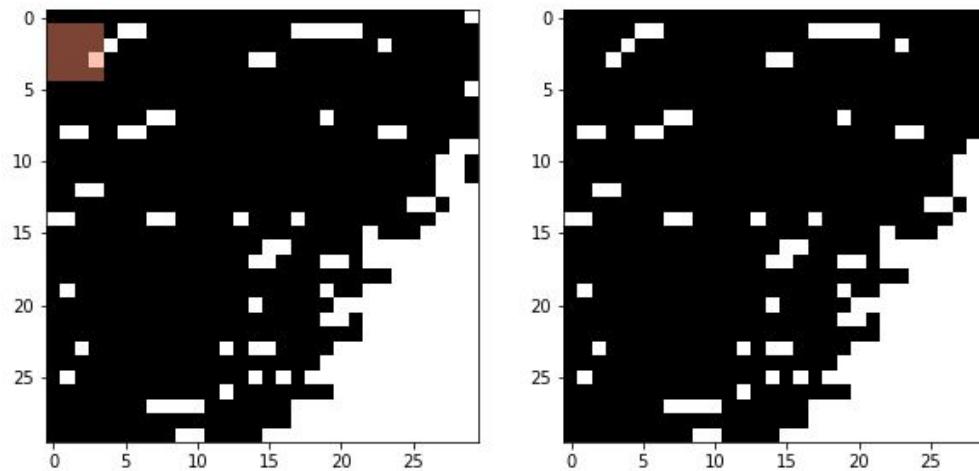
# Binary operations

To handle these cases, we use binary operators. These are filters that turn pixels ON/OFF depending on the local neighborhood. The simplest ones are the eroding and dilating filters.

## Erosion/dilation:

Whenever the subregion contains at least an OFF/ON pixel, turn that subregion location (here, upper left corner) OFF/ON

Erosion



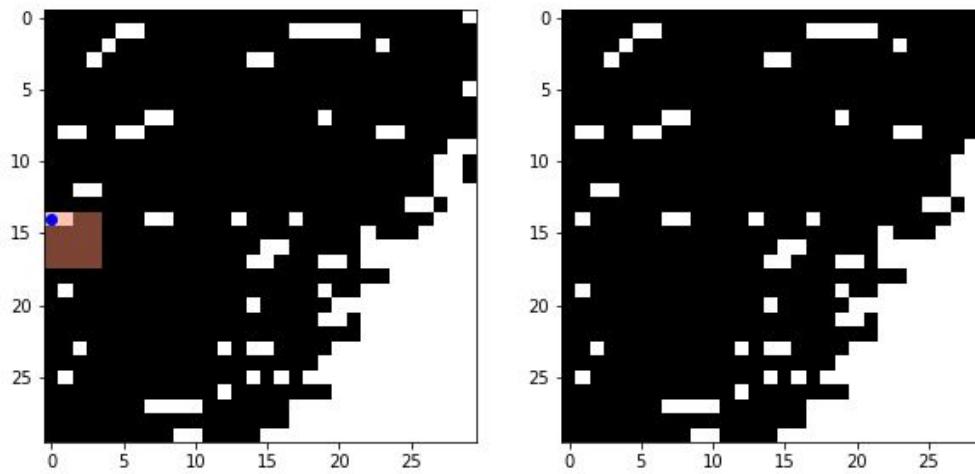
# Binary operations

To handle these cases, we use binary operators. These are filters that turn pixels ON/OFF depending on the local neighborhood. The simplest ones are the eroding and dilating filters.

## Erosion/dilation:

Whenever the subregion contains at least an OFF/ON pixel, turn that subregion location (here, upper left corner) OFF/ON

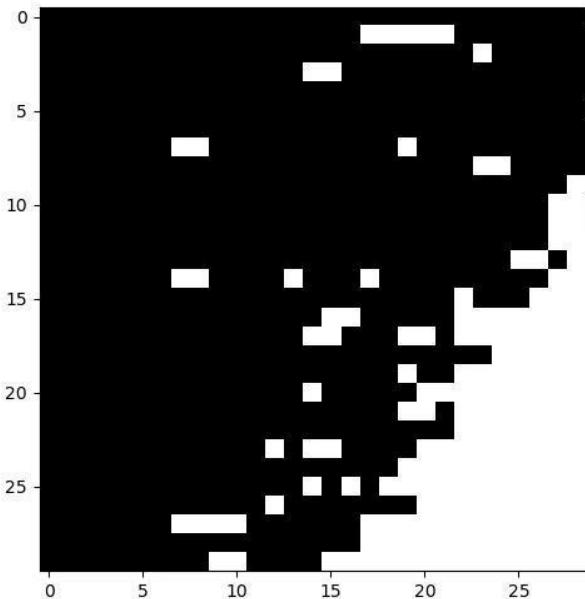
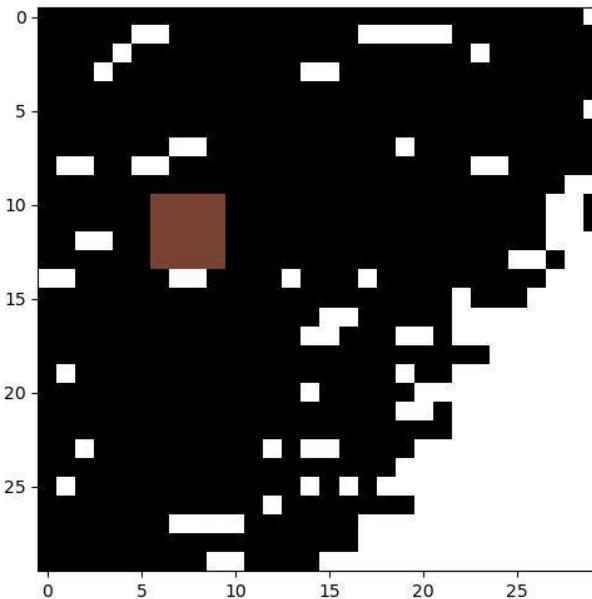
Erosion



You can play with the *binary* web-app

# Binary operations

To handle these cases, we use binary operators. These are filters that turn pixels ON/OFF depending on the local neighborhood. The simplest ones are the eroding and dilating filters.



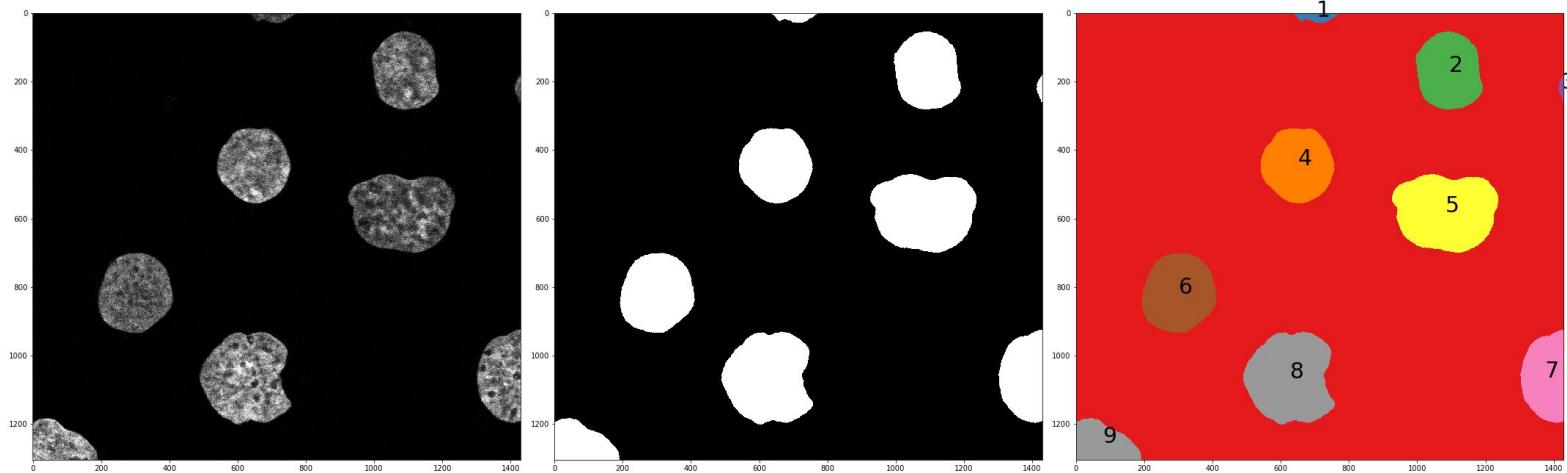
Operators can be combined (e.g. erosion + dilation)

Other more sophisticated operators exist (e.g. thinning)

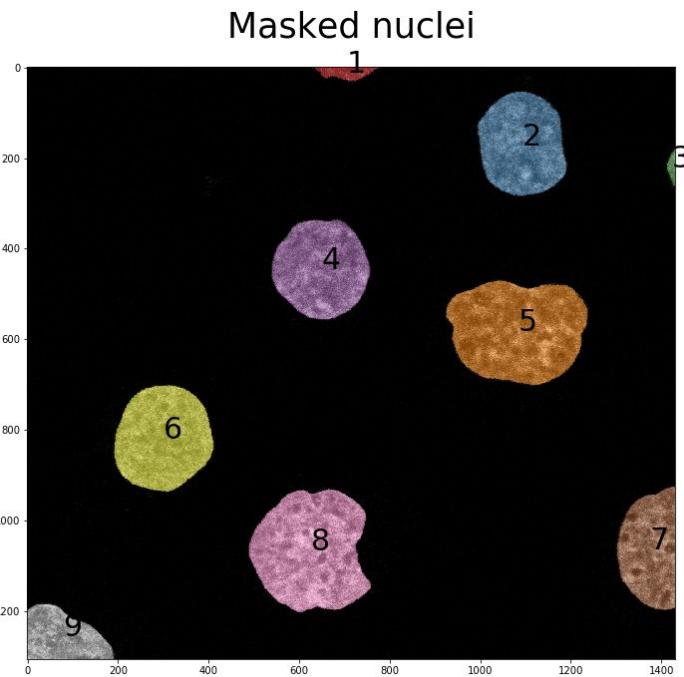
# Binary operations: regions

The cleaned binary image (middle), it can be used to recover information:

- About the structures themselves (size, shape etc.)
- About the “content” of these structures imaged in other channels by using the binary image as a **mask**



# Binary operations: region properties



Index	Area	Mean intensity	Eccentricity (shape)
1.0	2881.0	69.1	0.97
2.0	35179.0	105.3	0.56
3.0	877.0	88.4	0.97
4.0	36579.0	121.7	0.28
5.0	56095.0	86.1	0.68
6.0	39680.0	93.5	0.44
7.0	27606.0	121.2	0.87
8.0	55195.0	132.4	0.36
9.0	17178.0	146.0	0.83

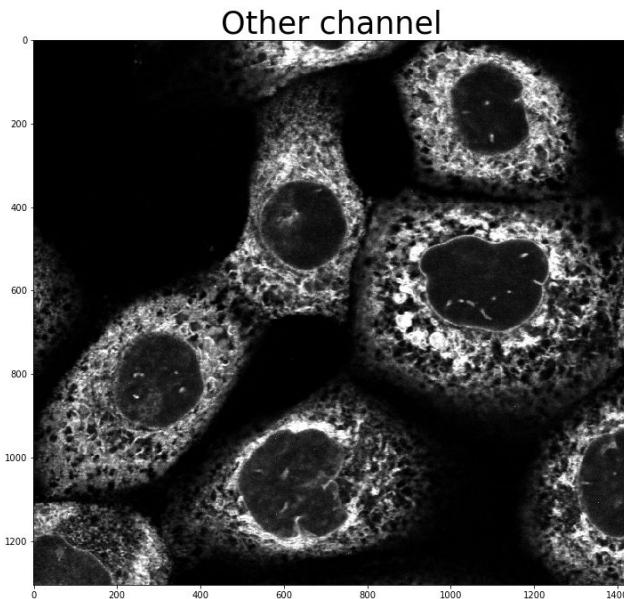
The mask can be used to recover information about the underlying image which can be any channel.

# Watershed segmentation

What can we do if we want to segment an image where boundaries are not as clean as e.g. for the nuclei?

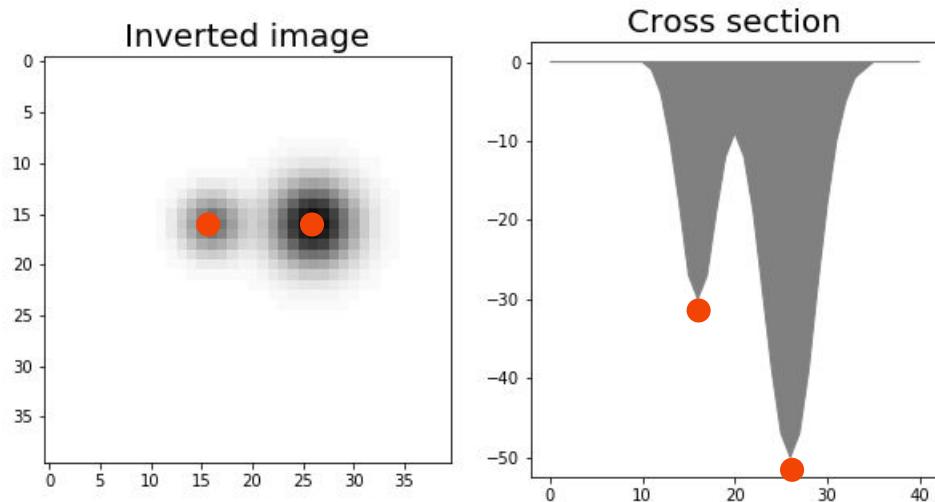
We can exploit our knowledge about nuclei (we know how many cells there are and where they are centered) and “propagate” it to a larger area.

For that we consider the nuclei as “seeds” from where we explore the area around.



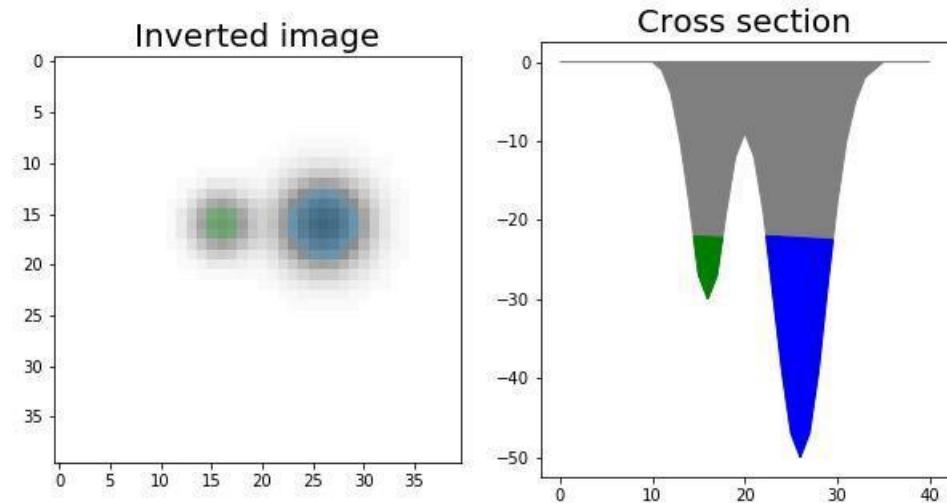
# Watershed segmentation

- We start filling from the bottom of the wells, in our case the nuclei **seeds**
- We fill up each well
- When two wells “meet” we create a boundary and extend it until the image is covered (various solutions)



# Watershed segmentation

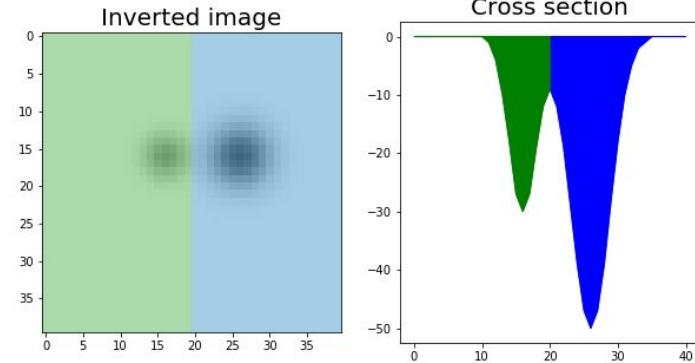
- We start filling from the bottom of the wells, in our case the nuclei **seeds**
- We fill up each well
- When two wells “meet” we create a boundary and extend it until the image is covered (various solutions)



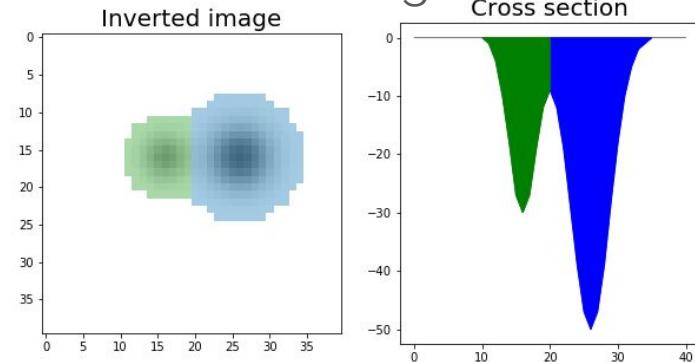
# Watershed segmentation

- We start filling from the bottom of the wells, in our case the nuclei **seeds**
- We fill up each well
- When two wells “meet” we create a boundary and extend it until the image is covered (various solutions)
- One can apply at the end a global mask to only retain interesting regions.

Final Segmentation



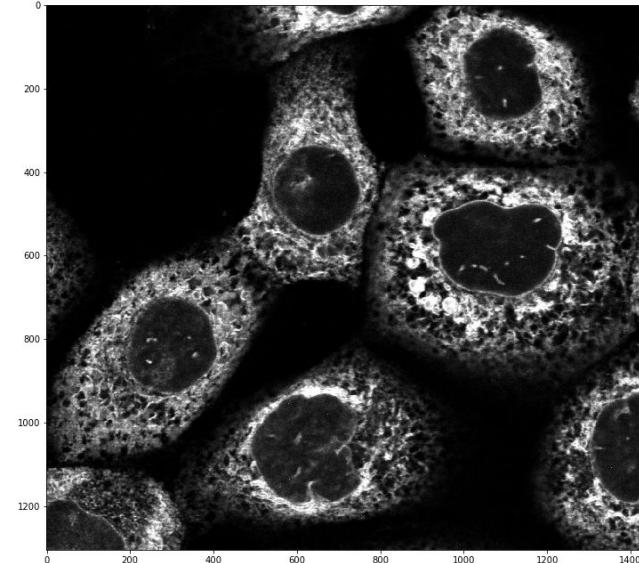
Masked Final Segmentation



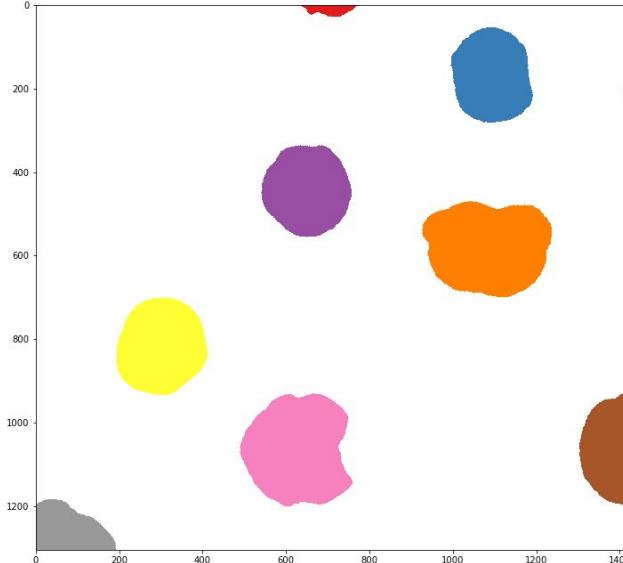
# Watershed segmentation

Applying this method on our difficult case, allows us to at least approximately segment the whole cells.

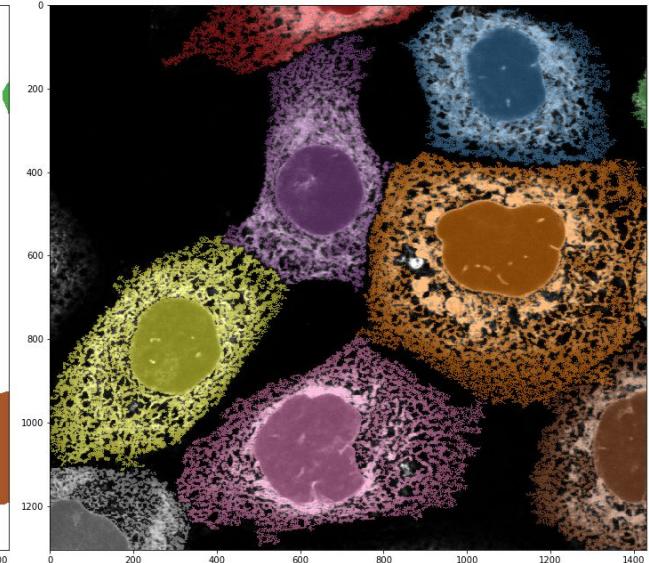
Other channel



Nuclei seeds

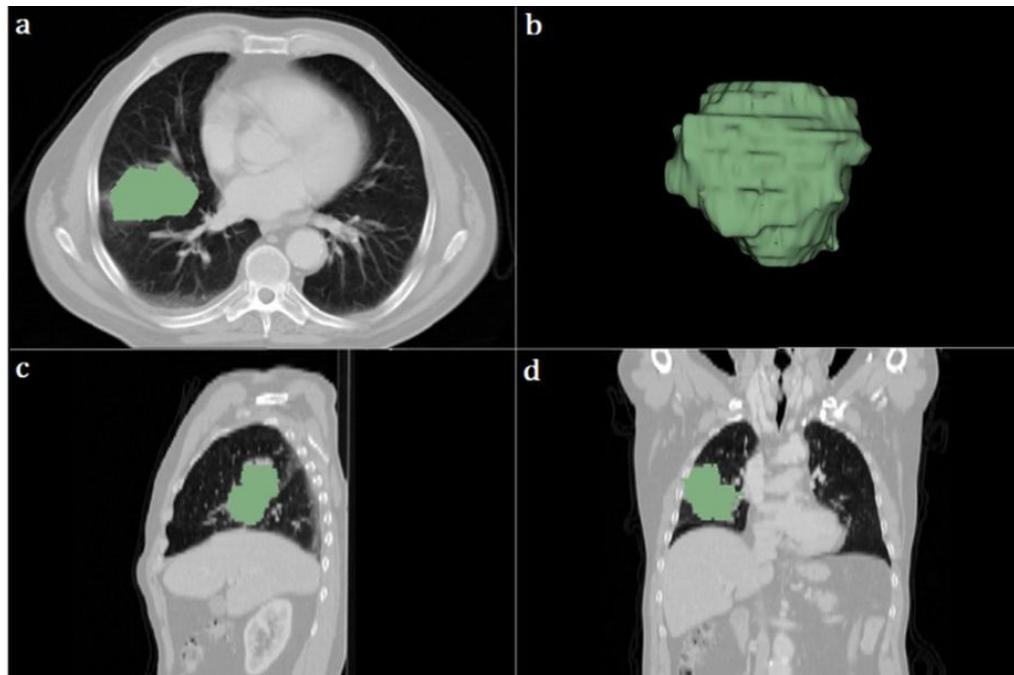


Labelled mask



# 3D segmentation

Most of the algorithms used in 2D can be extended to 3D at the price of computational time. Here for example, a tumor is segmented in 3D in a lung CT-scan.



Volumetric CT-based segmentation of NSCLC using 3D-Slicer, Velasquez et al. Scientific Rep. 2013

# VI. Advanced image processing

# Pixel classification

How do we classify things that we see ? Cat or not cat ?



Four legs ?	0.9
Horns ?	0.1
Pointy ears ?	0.7
Carapace ?	0.01
Fur ?	0.95
etc...	

From having seen many cats we know how these properties match with cats, i.e. our brains have been trained to associate a set of properties with cats.

→ Cat 80 %

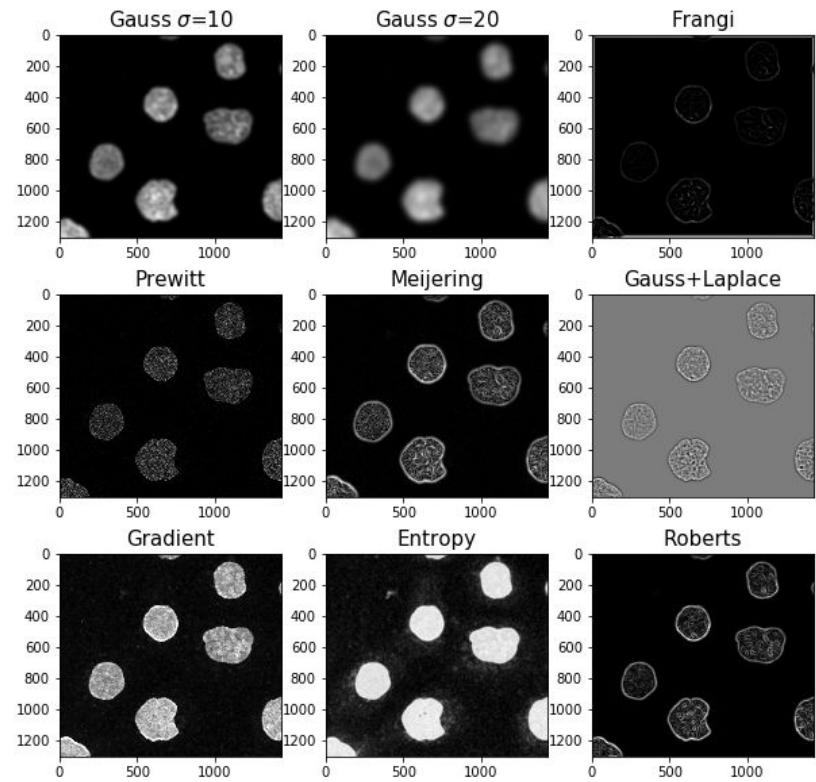
We can do the same for pixels !

# Pixel classification: calculating multiple properties

We can apply a series of filters to our image. For each pixel we have now a list of 9 values, one for each filter, that we can use as info (just like “four legs?”, “fur?”) to classify them.

Still missing: what information says “this is a nucleus”? We have to provide learning examples! In this artificial example, we already have the segmentation.

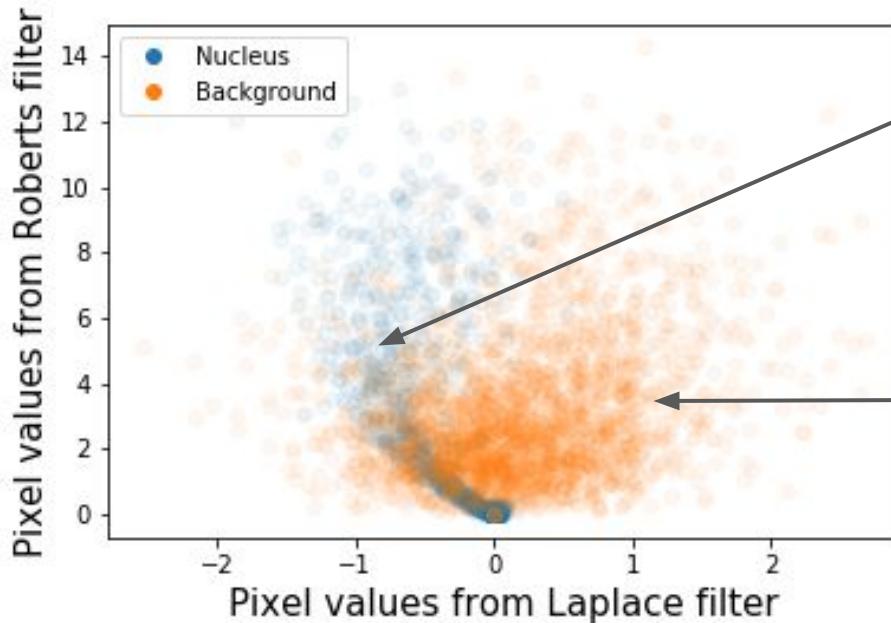
In a normal case we would have e.g. to manually draw regions to provide examples of “what is a nucleus”.



Various filters applied to nuclei image

# Pixel classification: with just two filters

Each colored dot represents a single pixel.

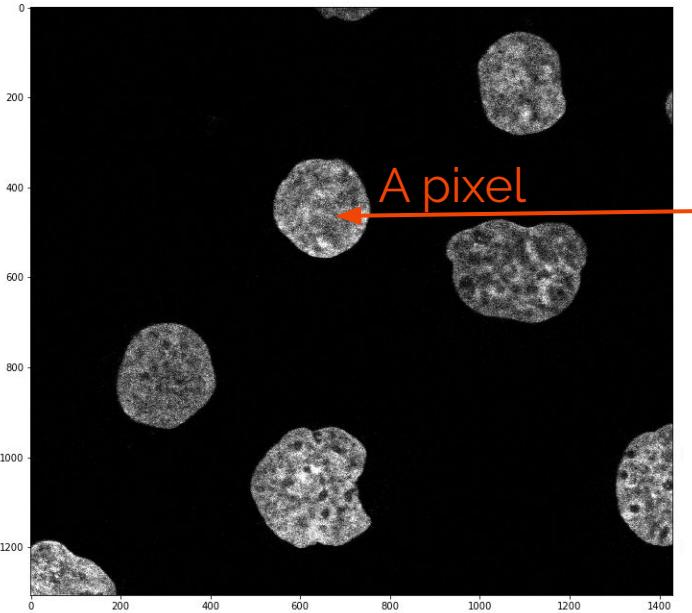


It's unclear what a pixel in this region is: nucleus or background ?

We need another variable!

A pixel with value 1 in Laplace filter and 4 in Roberts filter has a good chance of being Background

# Pixel classification



Gauss  
Frangi  
Gradient  
Laplace  
Entropy  
etc...

12  
1  
4  
Trained ML  
algorithm  
23  
0.1

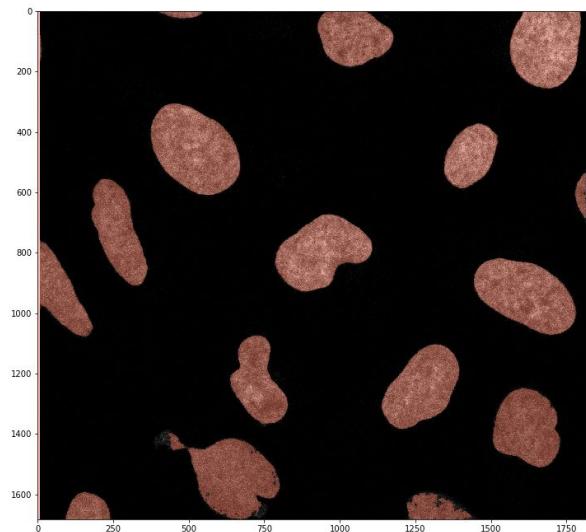
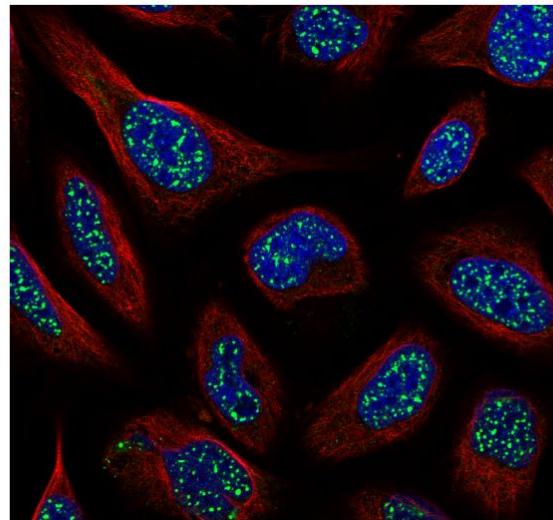
Nucleus  
85 %

A vertical brace groups the values 12, 1, 4, 23, and 0.1, with the text "Trained ML algorithm" positioned above the brace. An orange arrow points from the text "Nucleus 85 %" to the value 0.1.

# Pixel classification: unseen image

We can now use the trained classifier on an image it has never seen.

We don't have to worry about cleaning the image (binary operations). The classifier has learned that dark pixels surrounded by bright ones belong to nuclei.



<http://cellimagelibrary.org/images/41688>

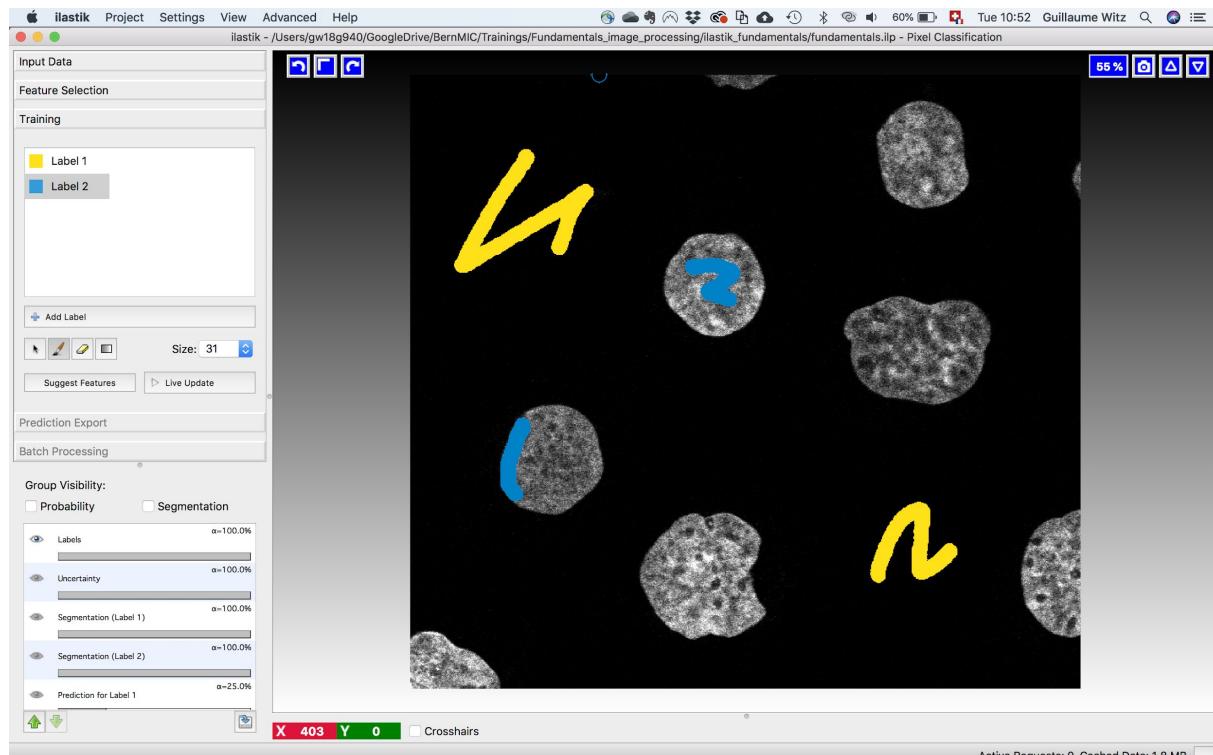
<https://www.proteinatlas.org/>

doi:10.7295/W9CIL41688

# Read-to-use solutions: Ilastik

Usually one does not have a segmentation available for training.

Ilastik is a software where one can manually indicate to what category regions belong.

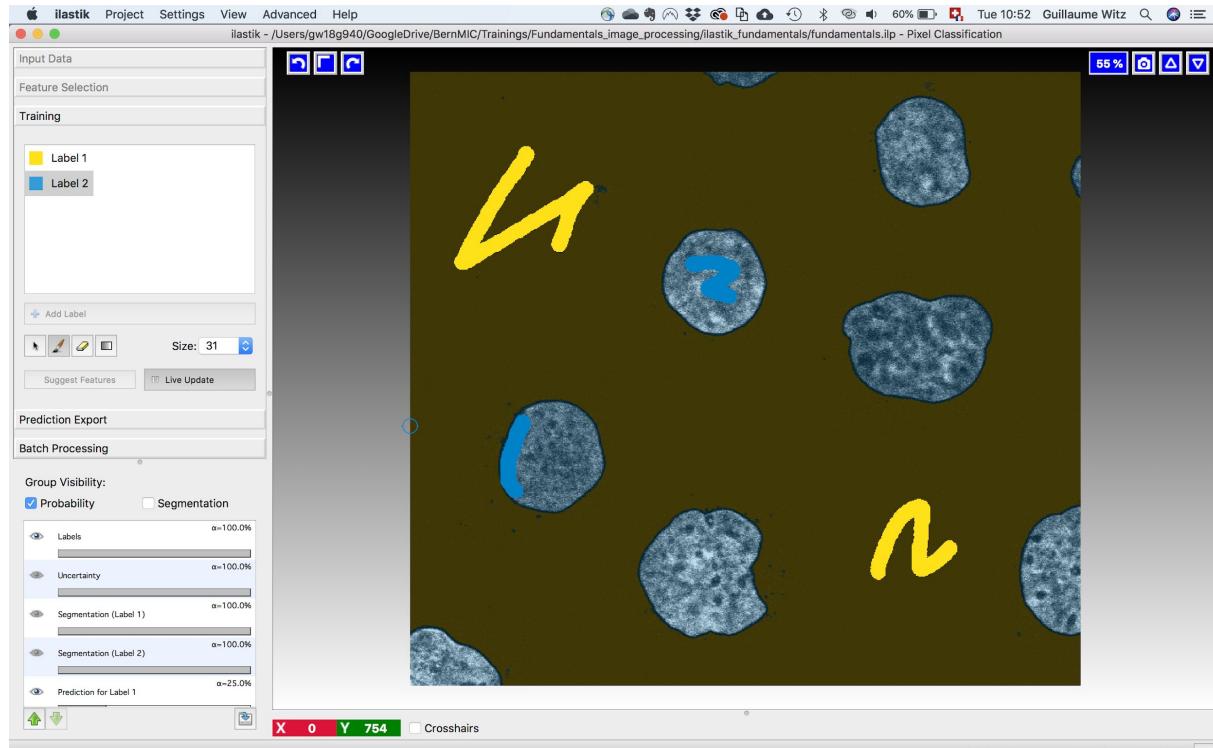


# Read-to-use solutions: Ilastik

Usually one does not have a segmentation available for training.

Ilastik is a software where one can manually indicate to what category regions belong.

Using that information the software then runs a classifier and provides a segmentation.

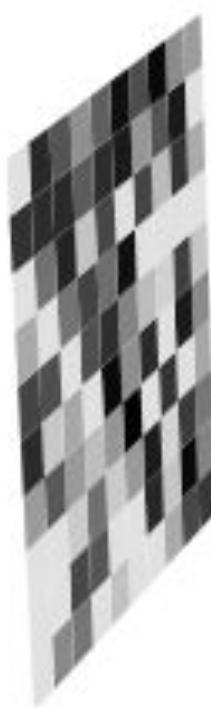


<https://www.ilastik.org>

# Deep learning

Deep learning uses the basic concepts covered before (filtering, downsampling etc.) and combines them in a powerful way.

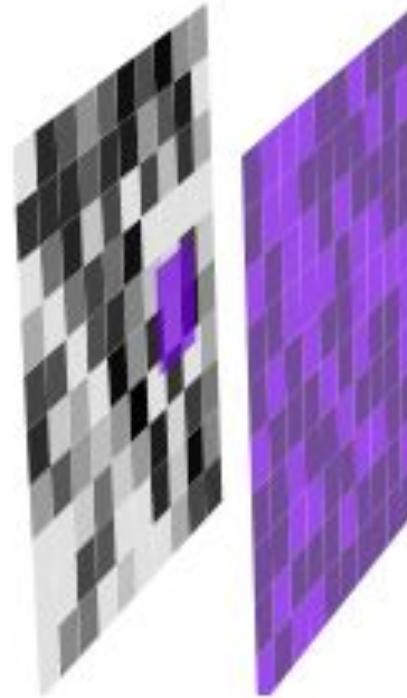
We start with an image we want to segment.



# Deep learning

## Convolution:

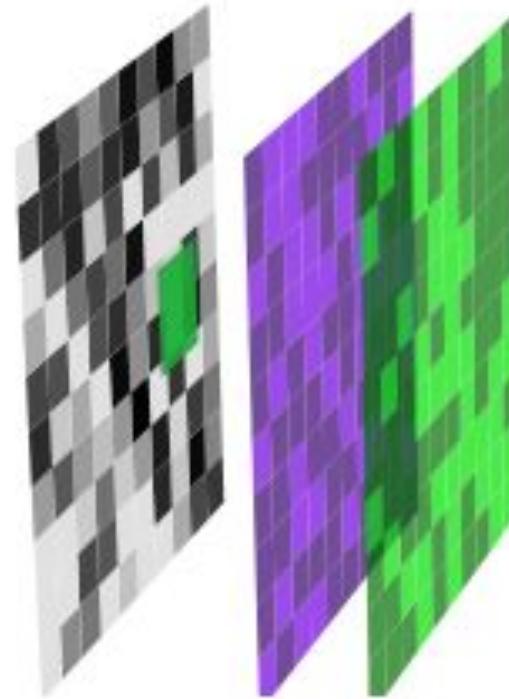
We filter the image with  
a series of filters to  
identify features



# Deep learning

## Convolution:

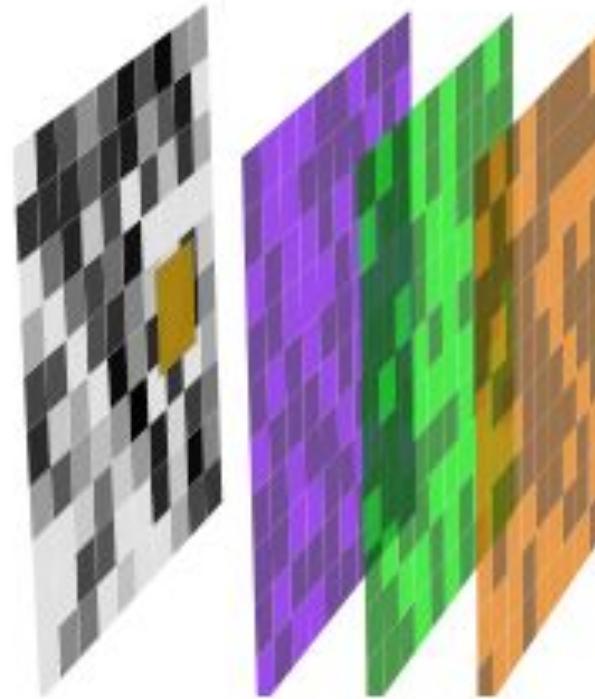
We filter the image with  
a series of filters to  
identify features



# Deep learning

## Convolution:

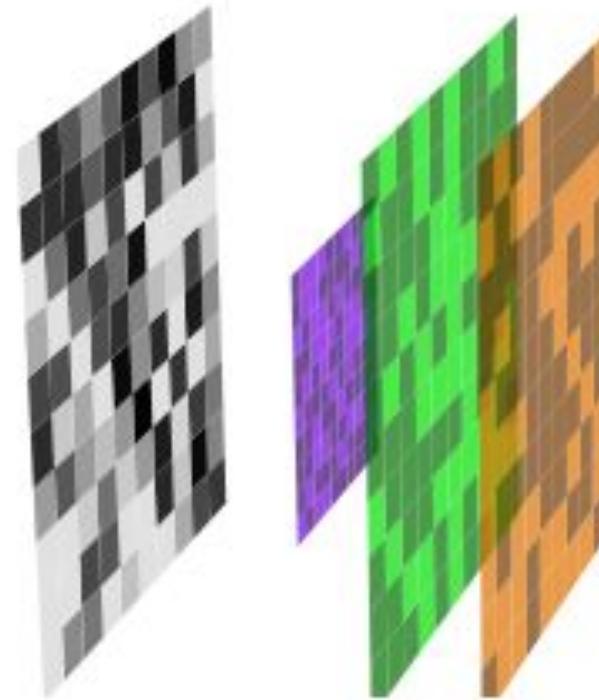
We filter the image with  
a series of filters to  
identify features



# Deep learning

## Pooling:

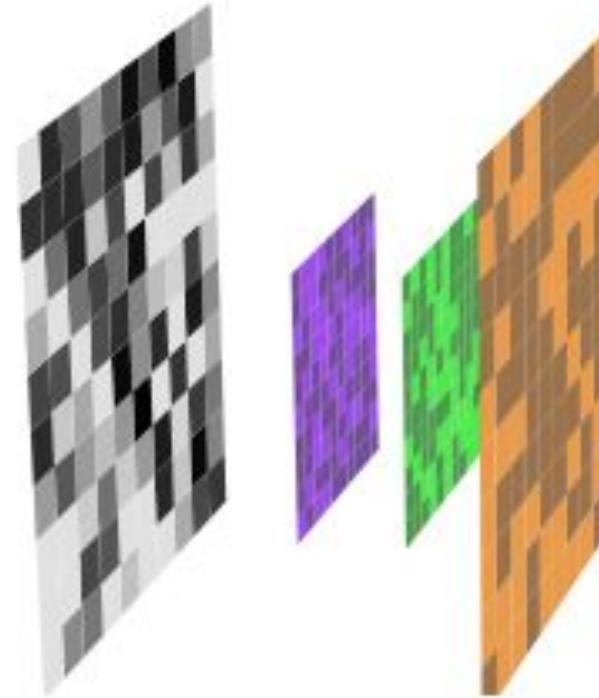
We bin the filtered images in order to look at another **scale**



# Deep learning

## Pooling:

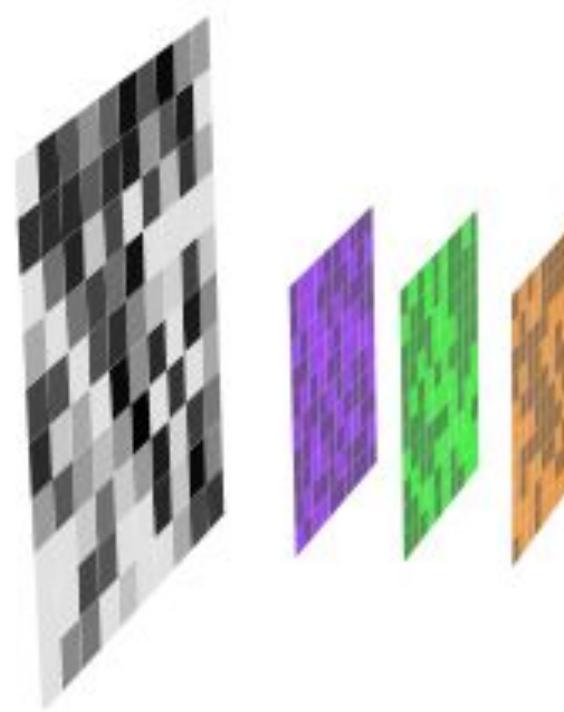
We bin the filtered images in order to look at another **scale**



# Deep learning

## Pooling:

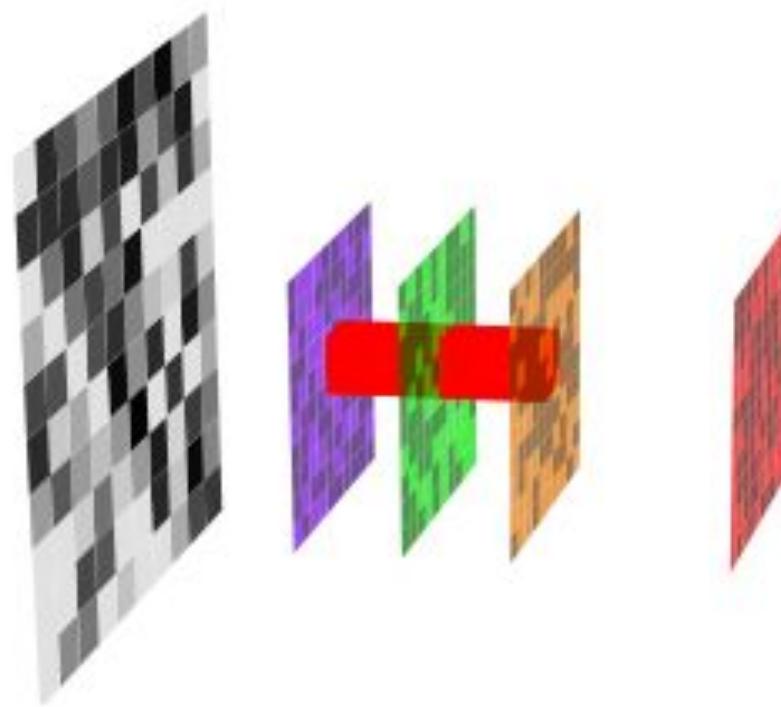
We bin the filtered images in order to look at another **scale**



# Deep learning

## Convolution:

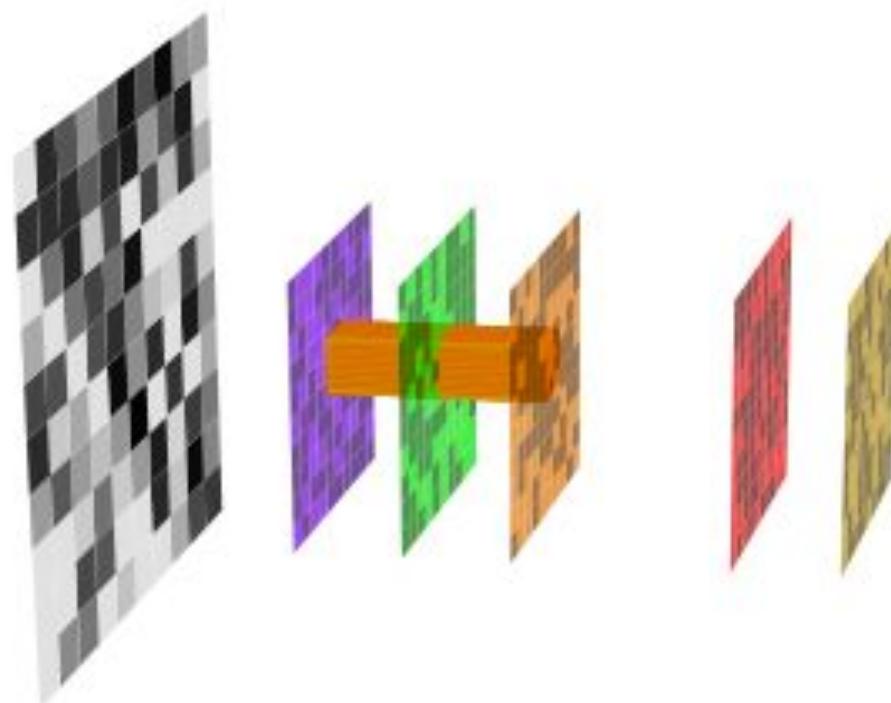
We make a new round of filtering, but now we combine the information of the different layers.



# Deep learning

## Convolution:

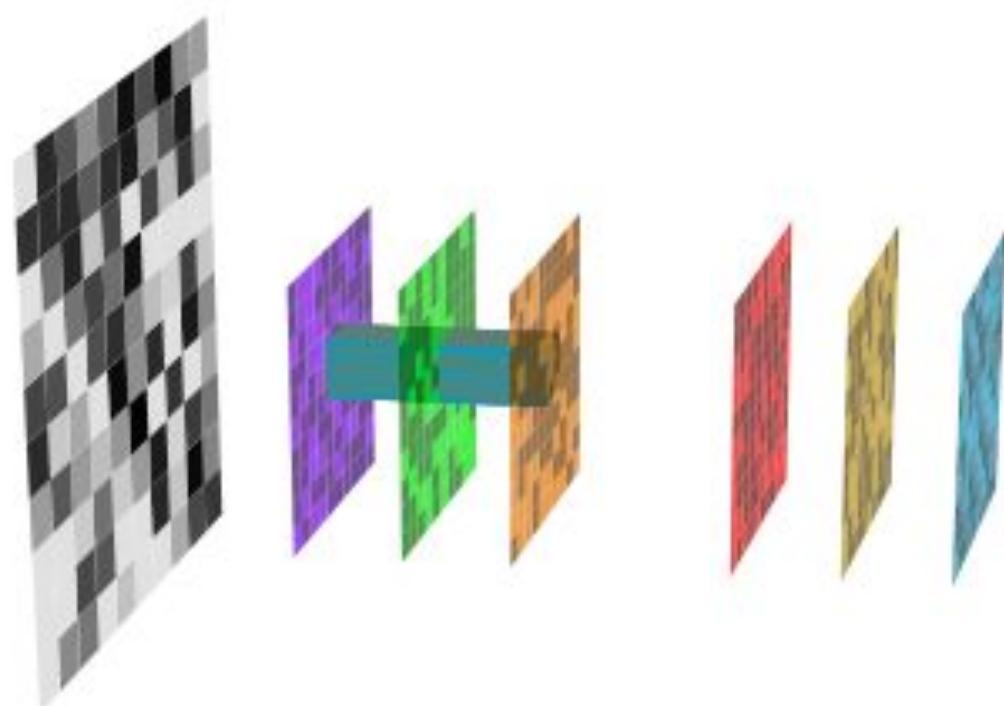
We make a new round of filtering, but now we combine the information of the different layers.



# Deep learning

## Convolution:

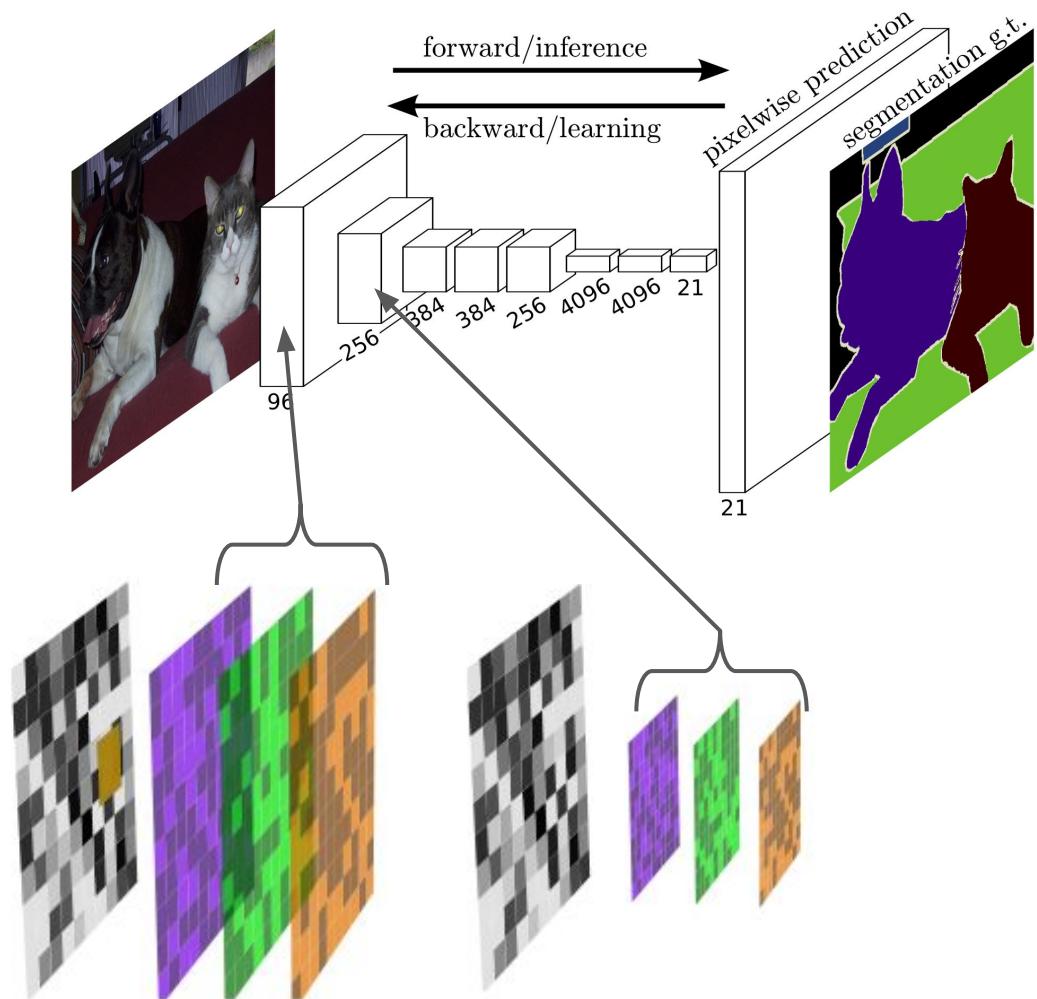
We make a new round of filtering, but now we combine the information of the different layers.



# Deep learning

## Segmentation:

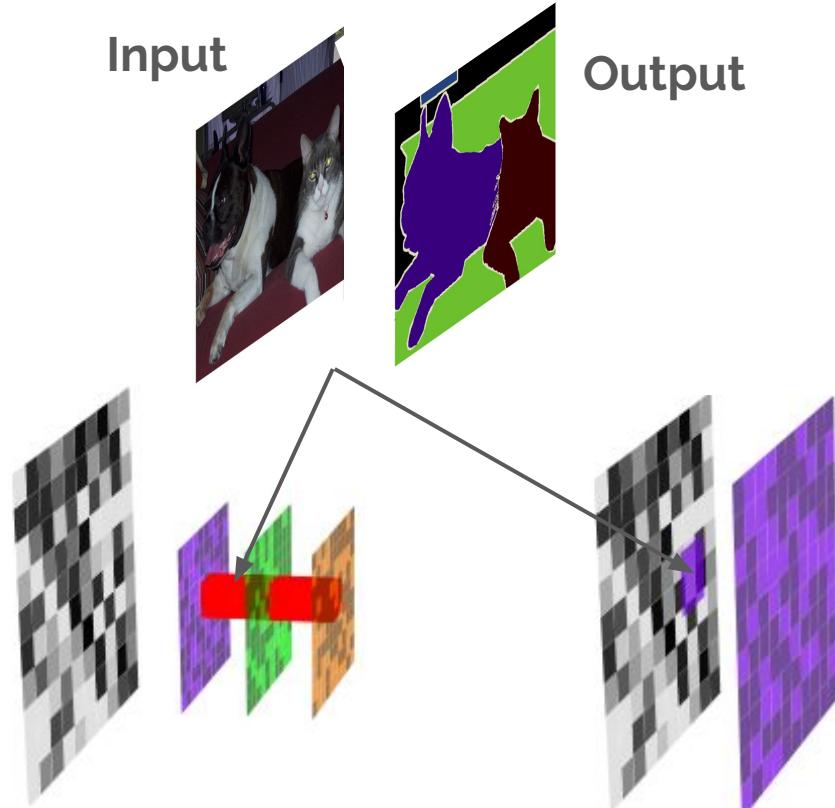
We can combine many of these layers (**deep** learning). The output of our network should be a map indicating the category of each pixel.



# Deep learning

## Training:

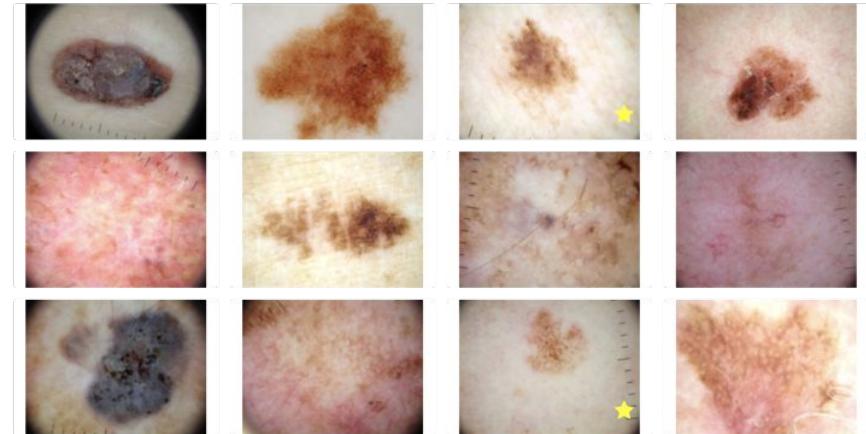
What gets trained here are all the filters. The filters are adjusted in many small steps (as done when fitting a function to data) so that the input of the network (cat picture) matches the output (segmentation) **for the training set**. Then **new images** can be passed through the network.



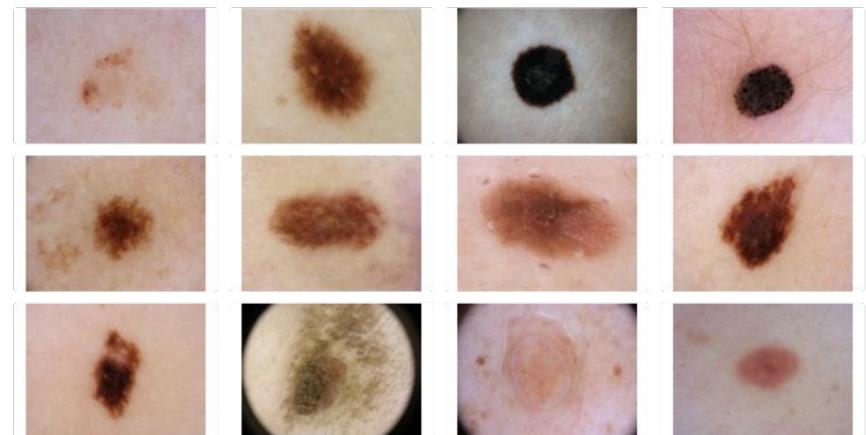
# Word of caution

In the pixel classifier seen previously, we designed the filters and could figure out which were important.

In the deep learning case, the inner workings of the network are (for the moment) very difficult to understand. One has therefore to be very careful during training, and thoroughly verify the quality of the results.



Malignant cancer (often with ruler)



Benign cancer (no ruler)

# Questions ?

If you want to learn more about image processing, the next date for my 1 day course is 21 January 2020. The material is available here:

[https://github.com/guiwitz/Python\\_image\\_processing](https://github.com/guiwitz/Python_image_processing)

You can find more infos about that course and the other courses of Science IT Support here:

[http://www.scits.unibe.ch/training/training\\_and\\_workshops/](http://www.scits.unibe.ch/training/training_and_workshops/)